

How Should Your Agent Talk to Mine?

The Security–Utility Frontier in Cross-Boundary Agent Delegation



Xisen Wang

Keble College

4th-Year Project Report

Supervised by
Professor Philip Torr

Department of Engineering Science
University of Oxford

Trinity Term, 2026

Acknowledgements

I owe this thesis to every person and element in my life that held me together through the hardest stretch of putting it all into one piece. This also marks the journey of my four years in engineering—from horizontal explorations to converging on the mission of building infrastructure for agentic networks and agentic society.

To the Torr Vision Group. To Prof. Jindong Gu, for introducing me into AI Safety. To Dr. Adel Bibi, for his encouragement in AI Security and the sharable OS concept. To Kevin Lin, for the guidance and perspective that kept this work grounded. And to Professor Philip Torr, for giving me the freedom to chase a problem I truly cared about.

To my co-founders at Pulse/Aicoo—and to all those who joined us along the way—for supporting the system from a research prototype to a product where thousands actually deploy their agents to interact with each other.

To every early user and beta tester of Pulse who trusted an agent with their data and gave us the signal to keep going.

To my friends, for the unwavering trust in me. To my family, for the unconditional belief that made it possible to bet everything on an idea before it had a name. And to the important people who stayed close enough to remind me who I was when the work became too heavy.

And to Oxford — for being the place where this all started.

Abstract

Abstract

Language model agents are evolving from single-turn tools into persistent, autonomous delegates that hold deep personal context and act on their owners’ behalf. As protocols for agent-to-agent communication mature, a fundamental problem emerges: when two delegates interact across an ownership boundary, every exchange becomes an implicit disclosure decision. Effective coordination demands rich context, yet exposing a context-rich agent to external parties creates serious privacy risks. We call the set of achievable (security, utility) operating points for a given system configuration the *security–utility frontier*, and its shape is the subject of this thesis.

We make five contributions. First, we **formalise cross-boundary delegation** as a setting where security enforcement is delegated to LLM reasoning rather than implemented by deterministic mechanism, and show why traditional access control models are insufficient. Second, we design and deploy **SharedOS**, a multi-agent delegation platform with per-relationship context scoping, configurable privacy policies, and tool orchestration—serving both as a real-world coordination system and as a controlled experimental apparatus. Third, we develop two complementary benchmarks: **PACT-PAIR** (600 dyadic scenarios spanning five sensitivity categories, three relationship tiers, and a six-level defence gradient) and **PACT-NET** (997 network scenarios testing transitive disclosure across chains of three or more agents). Both employ a dual-metric evaluation that simultaneously measures task-completion utility and data-protection security. Fourth, we derive **empirical insights** from these benchmarks: generic policies perform no better than no policy, category-specific policies hit a prompt-level ceiling, multi-turn interaction creates a $3\times$ incidental leakage gap, and relationship context can shift the cost from leakage to over-refusal. Fifth, motivated by these failures, we propose **architectural interventions**: *Mountable Context Cells*, which enforce disclosure policy by structural data absence rather than instructional governance, and an *Intelligent Escalation Protocol* that defers ambiguous decisions to the human principal while learning social boundaries through precedent clustering.

Our findings indicate that architectural enforcement—not prompt-level restriction—is the path toward secure cross-boundary agent coordination.

Table of contents

Table of contents	iv
List of figures	viii
List of tables	ix
Glossary	xi
1 Introduction	1
1.1 The Rise of Personal Agent Delegates	1
1.2 The Cross-Boundary Delegation Problem	2
1.3 Why This Is Different From Traditional Access Control	3
1.4 Research Questions	3
1.5 Contributions	4
1.6 Thesis Organisation	4
2 Related Work	6
2.1 AI Agent Architectures	6
2.2 LLM Security and Privacy	7
2.2.1 Attack Methodologies	7
2.2.2 Defence Mechanisms	7
2.2.3 Privacy in Language Models	7
2.3 Agent Security Benchmarks	8
2.4 Trust and Access Control	10
2.5 Position: SharedOS at the Intersection	10
3 Problem Formulation & SharedOS	11
3.1 Problem Formulation	12
3.1.1 The Personal Agent as an Operating System	12

3.1.2	Communication Between Personal Agents	13
3.1.3	Combinatorial Complexity	13
3.2	Engineering Implementation	14
3.2.1	Shared Execution Core	14
3.2.2	Policy-as-Documents	15
3.2.3	Capability-Based Share Links	17
3.2.4	Permission Model and Contact Graph	18
3.2.5	Pre-Tool Escalation Gate	19
4	PACT-PAIR: Dyadic Evaluation	21
4.1	Experiments	21
4.2	Findings	23
4.2.1	Setup	23
4.2.2	Core Findings	23
4.2.3	Ablations	28
4.3	Discussion	31
4.4	Failure Taxonomy	32
5	PACT-NET: Network Evaluation	35
5.1	Motivation & Design	35
5.2	Experiments	36
5.3	Findings	38
5.3.1	F1: Third-party information leaks indirectly	38
5.3.2	F2: Information crosses social and organisational clusters	38
5.3.3	F3: One leak often becomes multi-fact disclosure	39
5.3.4	F4: False delegation creates confused deputies	39
6	Architectural Solutions	41
6.1	Relationship-Specific Policy	42
6.2	Mountable Context Cells	42

6.2.1	Concept & Formulation	42
6.2.2	Implementation	43
6.3	MCC Design and Implementation	44
6.3.1	Assembly Pipeline	44
6.3.2	Namespace-Based Permissions	45
6.3.3	Relationship-to-MCC Mapping	45
6.4	MCC Experiments & Results	45
6.4.1	Experimental Design	45
6.4.2	Results	46
6.4.3	Network-Scale MCC Validation	48
6.4.4	Failure Modes MCC Cannot Address	49
6.5	Intelligent Escalation Protocol	50
6.5.1	Motivation	50
6.5.2	Pre-Tool Escalation Gate	51
6.5.3	Boundary Layers	51
6.5.4	Precedent-Based Tool Gating	52
6.6	Escalation Protocol Validation	52
6.6.1	Experimental Design	52
6.6.2	Results	53
6.7	Production Deployment	55
6.7.1	Folder-Scoped Share Links as Production MCC	55
6.7.2	Relationship Shards as Per-Requester Context	55
6.7.3	Escalation in Production	56
6.7.4	What Works vs What Remains Theoretical	56
6.8	Summary	57
7	Conclusion & Discussion	58
7.1	The Frontier Cannot Be Flattened By Prompt Engineering	58
7.2	Implications for Agent System Design	60

7.2.1	Safety Training Is Not Privacy Governance	60
7.2.2	Relationship Context Is a Feature, Not a Bug	60
7.2.3	Over-Refusal Is the Dominant Real-World Failure	61
7.2.4	Read Trust and Write Trust Are Independent	61
7.3	SharedOS as Research Infrastructure	61
7.4	Limitations	62
7.5	Ethical Considerations	63
7.6	Summary	64
7.7	Future Work	66
7.7.1	Formal Attack Integration	66
7.7.2	Cross-Model Adversarial Evaluation	67
7.7.3	Per-Field Mountable Context Cells	67
7.7.4	Formal Verification of Context Cell Isolation	67
7.7.5	User Study with Real Humans and Real Relationships	67
7.7.6	SharedOS Open-Source Release	68
7.7.7	Network-Scale MCC: Transitive Propagation	68
7.8	Closing Statement	68
Bibliography		70

List of figures

3.1	SharedOS as a multi-agent shared delegation system. Each owner’s agent accesses private files and structured state through typed tools. A cross-boundary delegation layer loads relationship context, applies governance policy, and routes inter-agent requests. External agents may query or request actions, but cannot bypass the tool layer.	12
4.1	Policy specificity determines security across all boundary surfaces. Files QA (left) averages GPT-5-mini and GPT-5.5 (both LLM judge); States QA (centre) and Actions (right) use GPT-5-mini (2 replications pooled). D1 provides no reliable protection over D0. D2 reduces leakage by 52–78 pp and eliminates the most dangerous action categories entirely. Full 6-model results in Table 4.4.	24
4.2	Multi-turn erosion case study: the wedding cascade (D2, gpt-5-mini, split s06). Phase 1 (Tick 18) produces a clean refusal. Phase 2 retries (Ticks 80–91) fail 12 consecutive times because the requesting agent searches only Work/shared folders. At Tick 92, a three-strategy combination (scope expansion + business justification + constrained format) breaks through, leaking the wedding date, venue, and guest count. Tick 93 chains the leaked note ID to extract venue confirmation. Ticks 94–99 show policy recovery: all six follow-up probes targeting the partner’s name are refused. The erosion is bounded: the policy does not collapse after a temporary breach.	25
4.3	The privacy–utility frontier across three dimensions. <i>Left:</i> Policy specificity—D2 shifts all six models to high security; D1 is indistinguishable from D0. <i>Center:</i> Cumulative multi-turn leakage (240 ticks). D2 caps exposure below 30% for both models. <i>Right:</i> Relationship-conditioned dual failure under D2. Red: leak rate (driven by <i>sensitive_work</i>). Blue: over-refusal on legitimate items. Jordan exemplifies the worst case—leaks 9% yet blocks 86% of entitled access.	28

List of tables

2.1	Comparison of agent security benchmarks. ✓ = fully supported, ~ = partially supported, – = not supported.....	9
3.1	Four entry points converging on a single execution core. The permission source differs; the execution path does not.	15
4.1	Experimental dimensions. Each RQ isolates one axis while controlling the others.	23
4.2	Multi-turn D2 security breakdown by category (10 splits × 240 ticks). Refuse = explicit refusal. Msg. Leak = response verified to contain gold facts. Failed Att. = attempted answer without gold facts. Glob. Leak = gold facts appear anywhere across 240 ticks (scan-based).	26
4.3	Relationship-conditioned D2 results (gpt-5-mini, 400 QA + 200 actions per requester, OR-adjudication). Leak = total P-item leak rate. O-Ref = over-refusal on L-items. Safety = unauthorised actions blocked.	27
4.4	Cross-model single-step files QA (6 models, 1 replication per model except gpt-5-mini with 2). All values in %. Util = work-public correctness. Ref = refuse rate. Msg.L = message leak rate. F.Att = failed attempt rate. D1 (omitted) is statistically indistinguishable from D0 in all models.	29
4.5	Multi-turn model scaling (10 splits × 240 ticks, 200 security QA per defence level). Ref = refuse rate. Msg.L = message leak rate (LLM judge). Glob.L = global leak rate (gold-fact scan). Util = work-public correct rate.	29
4.6	Cross-surface comparison (gpt-5-mini, single-step, pooled). Leak = fraction of sensitive queries where gold facts appear in response. Utility = fraction of legitimate queries answered correctly.	30
4.7	Prompt-level defence comparison (gpt-5-mini, 240 ticks, multi-step). Leak = global gold-fact rate. Act.S = unauthorised action block rate.	30
4.8	D2 failure mechanisms across surfaces and modes. Total is 44 unique leaked questions (28 from files, 16 from states) pooled across replications.	33
5.1	PACT-NET world and execution summary.	36

5.2	Clean PACT-NET V2 runs used in this chapter.	36
5.3	PACT-NET task families. Counts are per Phase 1 run.	37
5.4	PACT-NET composite scores, averaged over two namespace-isolated replications.	37
5.5	PACT-NET per-family accuracy, averaged over two replications. Action rows are response-heuristic scored.	38
5.6	Network-specific metrics. Lower is better for $\mathcal{J}$, $\mathcal{X}$, $\mathcal{A}$, and $\mathcal{D}$; higher is better for $\mathcal{C}$	39
6.1	Failure modes and their architectural remedies.	41
6.2	Per-requester MCC examples for the same target agent (Alex).	44
6.3	Relationship-to-MCC mapping rules (simplified).	45
6.4	Requester personas and MCC folder grants.	46
6.5	Three-condition comparison (Q1–400, Notes + Todos combined). 6,000 judged responses total (2,000 per condition). Leak rate = fraction of P/B questions where specific sensitive data was disclosed.	47
6.6	Decomposition of defence contributions by alignment class. Δ relative to D3 baseline.	47
6.7	PACT-NET MCC validation. Safety and Utility are composite benchmark scores. Transitive leak and Deputy success are lower-is-better network risks; Cross-surface defence is higher-is-better action safety.	48
6.8	Escalation gate Phase 2 ablation on PACT-NET. PStop is recall on private requests; Utility is recall on legitimate requests. All conditions use question-group splits and no auto-decision.	54
6.9	Deployment status of architectural solutions.	56

Glossary

This document is incomplete. The external file associated with the glossary ‘abbreviations’ (which should be called `pulse_4yp_thesis.gls-abr`) hasn’t been created.

Check the contents of the file `pulse_4yp_thesis.glo-abr`. If it’s empty, that means you haven’t indexed any of your entries in this glossary (using commands like `\gls` or `\glsadd`) so this list can’t be generated. If the file isn’t empty, the document build process hasn’t been completed.

Try one of the following:

- Add `automake` to your package option list when you load `glossaries-extra.sty`. For example:

```
\usepackage[automake]{glossaries-extra}
```

- Run the external (Lua) application:

```
makeglossaries-lite.lua "pulse_4yp_thesis"
```

- Run the external (Perl) application:

```
makeglossaries "pulse_4yp_thesis"
```

Then rerun $\LaTeX$ on this document.

This message will be removed once the problem has been fixed.

Chapter 1

Introduction

Personal agents are becoming delegates: systems that hold private data, wield tools, and act on behalf of their owners. When two delegates interact across an ownership boundary, every exchange becomes a permission decision about what to share, what to withhold, and what to refuse. This thesis investigates the security and utility consequences of cross-boundary agent delegation, proposes benchmarks to measure them, and develops architectural solutions that shift enforcement from probabilistic reasoning to deterministic structure.

The central observation is simple. A delegate that refuses every cross-boundary request is perfectly safe but entirely useless. A delegate that answers every request is maximally useful but catastrophically unsafe. Between these extremes lies a frontier of achievable operating points. We call this the *security-utility frontier*. The shape of this frontier, the factors that move it, and the architectural interventions that can expand it are the subject of this thesis.

1.1 The Rise of Personal Agent Delegates

The trajectory from language model to personal delegate has unfolded in three phases. **Phase One (Tools)** transformed language models into tool-users: ReAct [1] interleaves reasoning with executable actions, and Toolformer [2] showed tool use could be a learned competency. These systems were stateless—limited attack surface.

Phase Two (Assistants) produced stateful systems. Reflexion [3] enabled agents that critique and improve their own outputs across turns. MemGPT [4] introduced an OS-inspired memory hierarchy giving agents persistent long-term storage. An assistant with persistent memory holds a growing corpus of personal information—increasingly useful, but increasingly dangerous if exposed to unauthorised parties.

Phase Three (Delegates), now underway, produces persistent, autonomous systems that act on behalf of a specific person. Devin [5] operates as an autonomous software engineer. Manus [6] pushes desktop automation with persistent project state. These are not chatbots—they are digital proxies holding private context and wielding

powerful tools.

Two protocol-level developments accelerate the transition to a networked ecosystem. The Model Context Protocol (MCP) [7] standardises how agents connect to tools and data sources—analogous to USB for software. The Agent-to-Agent protocol (A2A) [8] addresses discovery, authentication, and communication between agents. Together, MCP provides the tool layer and A2A the networking layer. But neither addresses the operational question at the heart of every cross-boundary interaction: *what should one agent be permitted to disclose about its owner to another agent?* The infrastructure for connection exists. The infrastructure for safe connection does not.

1.2 The Cross-Boundary Delegation Problem

Consider a system with two agents, Agent A and Agent B, acting on behalf of Owner A and Owner B respectively. Agent A holds a context C_A consisting of Owner A’s private information. Agent B initiates a request r that requires some subset of C_A to be fulfilled. The cross-boundary delegation problem is: Agent A must decide, for each request r , what portion of C_A to disclose, balancing **utility** (sufficient information to fulfil the request) against **security** (no disclosure Owner A would not sanction).

As a concrete example: Alice’s agent holds her calendar, including an oncology follow-up (Tuesday), an interview at a competing firm (Wednesday), and a therapy session (Thursday). Bob’s agent asks to schedule a meeting. The ideal response shares time-slot availability without revealing *why* slots are blocked—“Tuesday afternoon doesn’t work” rather than “I can’t do Tuesday because of a medical appointment.” This requires distinguishing *what* information is needed from the sensitive content underneath, a distinction absent from traditional access control policies.

The problem is pervasive [9, 10]: it arises whenever a colleague’s agent asks for a project update (share progress, not internal conflicts), a recruiter’s agent probes career interests (openness vs. discretion), or a vendor’s agent negotiates budget (reveal constraints selectively). Critically, the problem exhibits an asymmetry of error: over-disclosure is irreversible, while under-disclosure can be retried or escalated.

1.3 Why This Is Different From Traditional Access Control

Traditional access control [11, 12] operates on a principle where **policy is enforcement**: a security label mechanistically determines access. The reference monitor is deterministic and formally verifiable.

LLM-based delegation violates this principle. The privacy policy is not enforcement but *one input to a reasoning process that produces enforcement as an output*. The agent receives the policy, the query, the conversational context, and the relationship—then reasons about what to do. The outcome is stochastic: the same agent, same policy, same query may produce different responses depending on phrasing, preceding conversation, and random seed.

This has three consequences. First, the frontier *cannot be predicted from policy text alone*—two policies with similar intent may produce different boundaries because the model interprets them differently. Second, the frontier *can only be measured empirically*, which is why we need benchmarks. Third, the frontier *is unstable*: implicit social norms from pretraining compose with explicit policy in ways the policy author cannot anticipate.

These properties mean that formal verification and lattice models are insufficient. We need empirical measurement of emergent boundaries, benchmarks that capture both dimensions simultaneously, and architectural interventions that recover the determinism of traditional access control.

1.4 Research Questions

RQ1: How do we measure cross-boundary delegation?

How should a personal-agent system represent the boundary between legitimate coordination and privacy leakage? This question motivates the SharedOS problem formulation and the PACT-PAIR dyadic benchmark: before proposing defences, we need an evaluation that measures security and utility together rather than treating refusal as success.

RQ2: What fails when delegation becomes networked?

Which failures appear only when agents interact across a social graph rather than a single pair? This question motivates PACT-NET, where transitive leakage, confused deputies, cross-cluster disclosure, and write-side laundering become measurable network-scale phenomena.

RQ3: Can architecture move the frontier beyond prompting?

Can enforcement be shifted from model instruction-following into the execution environment itself?

This question motivates Mountable Context Cells and the Escalation Protocol: mechanisms that remove unauthorised data before retrieval or stop ambiguous tool calls before private context enters the model.

1.5 Contributions

1. **Problem formulation.** We formalise cross-boundary delegation as emergent enforcement, define the security-utility frontier, and show why traditional access control is insufficient (Chapter 3).
2. **Platform.** We design SharedOS, a multi-agent delegation system with per-relationship context scoping, configurable policies, and tool orchestration—serving as both a deployed system and experimental apparatus (Chapter 3).
3. **Benchmarks.** PACT-PAIR (600 dyadic scenarios, five sensitivity categories, six-level defence gradient) and PACT-NET (997 network scenarios testing transitive disclosure). Both use dual-metric evaluation measuring utility and security simultaneously (Chapters 4 and 5).
4. **Empirical insights.** We show that generic privacy instructions do not move the frontier, category-specific policies have a ceiling, multi-turn interaction exposes incidental leakage invisible to single-turn tests, and relationship context mainly shifts the cost toward over-refusal (Chapters 4 and 5).
5. **Architectural solutions.** Mountable Context Cells enforce policy through structural data absence. The Escalation Protocol stops ambiguous tool calls before retrieval and uses relationship-specific precedents to learn owner boundaries (Chapter 6).

1.6 Thesis Organisation

Chapter 2 reviews agent architectures, LLM security, multi-agent protocols, and existing benchmarks. Chapter 3 formally defines cross-boundary delegation and presents the SharedOS platform. Chapter 4 describes the PACT-PAIR benchmark design and reports dyadic experimental results across the defence gradient. Chapter 5 presents PACT-NET, extending evaluation to multi-agent network scenarios with transitive disclosure risks. Chapter 6

proposes Mountable Context Cells and the Escalation Protocol, with validation experiments. Chapter 7 discusses findings, identifies limitations, and presents directions for future work.

Chapter 2

Related Work

This chapter reviews four bodies of literature: AI agent architectures (section 2.1), the attack and defence landscape for LLM-based systems (section 2.2), existing benchmarks for agent security and privacy (section 2.3), and trust and access control models (section 2.4). We conclude by positioning SharedOS at the intersection (section 2.5).

2.1 AI Agent Architectures

Single-agent systems. ReAct [1] established the reasoning-action loop that underpins modern agents. Toolformer [2] showed tool use can be a learned competency. MemGPT [4] introduced an OS-inspired memory hierarchy (main context, archival storage, paging), transforming stateless LLMs into persistent agents. These advances have translated into commercial systems such as Devin (autonomous software engineering) and Manus (desktop automation with GUI control). Individual agents are now capable enough to serve as personal delegates; the bottleneck is secure cross-boundary coordination.

Multi-agent frameworks. AutoGen [13] enables multi-agent conversation through structured role-based patterns. CrewAI [14] provides role-based workflows with defined goals and tool access. MetaGPT [15] simulates a software company with specialised roles, encoding SOPs as structured communication protocols. CAMEL [16] explores emergent communication between role-playing agents. The critical observation is that every existing multi-agent system is *multi-role-within-one-owner*: all agents share a single principal. None address the scenario where agents representing different owners with potentially conflicting privacy interests must coordinate.

Agent operating systems. AIOS [17] proposes OS-level primitives (scheduler, context manager, tool manager) for managing multiple agents on shared infrastructure, but its isolation is resource-oriented (preventing starvation), not privacy-oriented (preventing information leakage). OS-Copilot [18] operates *within* the OS as a single-owner agent with no cross-boundary considerations. No existing agent OS addresses information flow

control across trust boundaries. SharedOS occupies this gap.

2.2 LLM Security and Privacy

2.2.1 Attack Methodologies

Zou et al. [19] demonstrated with GCG that gradient-based adversarial suffixes break RLHF alignment and transfer across models, though the resulting non-natural-language tokens are detectable by perplexity filters. AutoDAN [20] uses genetic algorithms to generate fluent jailbreaks that evade such filters. PAIR [21] uses an attacker LLM to iteratively refine prompts against a black-box target in fewer than 20 queries.

In the persuasion domain, PAP [22] applies 40 social-science persuasion techniques to achieve >90% attack success on GPT-4—directly relevant to conversational agent-to-agent settings. Crescendo [23] demonstrates multi-turn escalation, beginning with benign topics and gradually steering toward sensitive information through individually innocuous turns.

Most critically, Nasr et al. [24] tested twelve published defences under adaptive attacks and found that *every defence* could be bypassed with >90% success; human red-teaming achieved 100%. This result motivates our architectural approach: if prompt-level defences fail under adaptive attack, security must be enforced by controlling what information is *present* rather than instructing the model not to reveal it.

2.2.2 Defence Mechanisms

Wallace et al. [25] proposed training LLMs to hierarchically prioritise instructions (system > user > tool), reducing indirect injection effectiveness. Hines et al. [26] introduced transformations (data marking, base64 encoding, XML delimiters) that help models distinguish instructions from data. Llama Guard [27] and NeMo Guardrails [28] provide classifier-based filtering. All operate at the prompt level: they are necessary but insufficient under adaptive attack [24]. This motivates both our defence gradient methodology and the Mountable Context Cell design.

2.2.3 Privacy in Language Models

Dwork et al. [29] established differential privacy (DP); Yang et al. [30] applied DP to agent communication, adding calibrated noise to bound per-query leakage at a utility cost our work measures empirically. Carlini et

al. [31] demonstrated that LLMs memorise and reproduce verbatim training data, establishing that LLMs are fundamentally unreliable as privacy-preserving components—any system depending on LLM behaviour for privacy inherits this unreliability.

2.3 Agent Security Benchmarks

A growing number of benchmarks evaluate agent security. We review them along six dimensions: cross-boundary interaction, tool-mediated evaluation, dual-metric measurement, multi-turn scenarios, relationship modelling, and action safety.

Prompt injection benchmarks. InjecAgent [32] evaluates indirect injection in tool-integrated agents (1,054 tests, 17 tools) but within single-agent contexts without measuring utility cost. AgentDojo [33] provides dual-metric evaluation (97 tasks, 629 security tests) but operates within a single task context with no cross-boundary interaction. TensorTrust [34] gamified prompt injection (126,000 attack/defence pairs), demonstrating human creativity but not modelling agent-to-agent settings. HackAPrompt [35] confirmed through 600,000+ adversarial prompts that human creativity consistently defeats automated defences.

Multi-agent and cross-boundary benchmarks. AgentSocialBench [36] evaluates privacy in agentic social networks, discovering the “abstraction paradox” (more capable agents are more susceptible to social engineering) but does not evaluate defences or measure utility. ConVerse [37] is the most comparable setup: personal assistants interacting with external providers (864 attack scenarios, up to 88% success), but measures attack success only with no utility dimension. PAC-BENCH [38] evaluates collaboration under privacy constraints with dual measurement but not across genuine trust boundaries. MAGPIE [39] provides 200 collaborative tasks but all agents share a single principal. AgentLeak [40] identifies seven distinct leakage channels but is diagnostic only. Agents of Chaos [41] red-teams deployed systems, demonstrating lab-to-production attack transfer without systematic evaluation methodology.

Table 2.1 summarises these benchmarks across the six dimensions critical to cross-boundary privacy evaluation.

No existing benchmark combines all six dimensions. InjecAgent and AgentDojo measure tool-mediated attacks within single-agent contexts. ConVerse crosses trust boundaries but measures only attack success. PAC-BENCH provides dual measurement but not across genuine trust boundaries. None model how the relationship

Table 2.1: Comparison of agent security benchmarks. ✓ = fully supported, ~ = partially supported, – = not supported.

Benchmark	Cross-boundary	Tool-mediated	Dual metric	Multi-turn	Relationship	Action safety
InjecAgent [32]	–	✓	–	–	–	✓
AgentDojo [33]	–	✓	~	–	–	~
TensorTrust [34]	–	–	–	–	–	–
HackAPrompt [35]	–	–	–	–	–	–
AgentSocialBench [36]	✓	–	–	–	–	–
ConVerse [37]	✓	–	–	✓	–	–
PAC-BENCH [38]	~	✓	✓	–	–	–
MAGPIE [39]	–	✓	–	–	–	~
AgentLeak [40]	✓	✓	–	–	–	–
Agents of Chaos [41]	✓	✓	–	~	–	✓
PACT (Ours)	✓	✓	✓	✓	✓	✓

between requester and data owner affects the security-utility frontier. PACT fills this gap.

2.4 Trust and Access Control

Classical access control operates on a principle where policy *is* enforcement. RBAC assigns permissions to predefined roles; ABAC evaluates access based on subject/object/environment attributes; Bell-LaPadula [11] formalises mandatory access control (“no read up, no write down”) for confidentiality; and Denning’s lattice model [12] constrains information flow through a partial order on security classes. All assume deterministic enforcement and static classification hierarchies.

Capability-based security [42] takes a different approach: unforgeable tokens grant specific permissions, enforcing the principle of least authority. A process receives only the capabilities it needs. Our Mountable Context Cell design is directly inspired by this model—an agent receives only the context cells relevant to the current interaction and cannot access data outside those cells regardless of its reasoning.

The key distinction in our setting is that access control is *relationship-relative*: the same information may be shareable with one requester but not another. Moreover, the “policy” (a system prompt or permission instruction) is not enforcement—it is *input to reasoning*. Under adversarial conditions, the model can reason its way around any natural-language instruction. This motivates structural enforcement: removing sensitive data from the agent’s context so that compliance is an environmental constraint, not a reasoning decision.

2.5 Position: SharedOS at the Intersection

The literature reveals a gap. Agent capability papers do not measure privacy—multi-agent frameworks assume shared principals and treat information sharing as uniformly beneficial. Security papers do not model delegation relationships—attacks succeed or fail without regard to whether the query is legitimate from one requester and illegitimate from another. Benchmark papers cover subsets of the evaluation space but none covers all six dimensions (Table 2.1).

SharedOS bridges these communities with: (i) a platform that instantiates cross-boundary coordination with persistent state, tool access, and social relationships; (ii) a benchmark (PACT) that evaluates across all six dimensions, measuring security and utility as functions of defence strength and relationship closeness; and (iii) architectural solutions that move enforcement from reasoning to structure, informed by the adaptive attack thesis and capability-based security.

Chapter 3

Problem Formulation & SharedOS

This chapter has two jobs. First, it formalises cross-boundary delegation: what a personal agent is, what it means for one agent to contact another, and why privacy becomes a security–utility frontier rather than a static access-control label. Second, it describes how SharedOS realises this formal model in production. The point is not only to define an evaluation setting, but to show that the setting is implemented as a real execution environment with tools, permissions, policies, share links, and escalation gates.




Figure 3.1: SharedOS as a multi-agent shared delegation system. Each owner’s agent accesses private files and structured state through typed tools. A cross-boundary delegation layer loads relationship context, applies governance policy, and routes inter-agent requests. External agents may query or request actions, but cannot bypass the tool layer.

3.1 Problem Formulation

3.1.1 The Personal Agent as an Operating System

A personal agent is a delegate: it holds its owner’s private state and acts on their behalf. We model each agent A_i as an operating system. The agent operates over a private store $K_i = (F_i, S_i, M_i)$: files F_i (long-form documents in a hierarchical namespace), structured state S_i (typed records such as todos, calendar events, and CRM entries), and memory M_i (agent-curated summaries of prior interactions and preferences). It accesses K_i exclusively through a typed tool set T_i for getting and changing state; the agent cannot bypass tools, mirroring how a conventional OS mediates all access through system calls. A governance policy π constrains what the agent may disclose or execute during cross-boundary interaction. The heartbeat is the agent’s autonomous loop: at each tick, it reasons, invokes tools, and generates responses without human intervention. Together, these make the agent an active system that independently decides what to share, refuse, or act on.

3.1.2 Communication Between Personal Agents

Consider a society of humans, each with one agent as delegate. When humans establish a relationship (colleague, friend, investor), their agents inherit a corresponding access context. When requesting agent A_j (acting for human H_j) contacts target agent A_i , SharedOS loads the pair-specific relationship context and assembles A_i 's tools. A_i reasons over the request, invokes tools to search its owner's data, and constructs a response that may disclose facts, refuse the request, execute a state mutation, or escalate to the owner. Tool results flow directly into the agent's generation context, creating the conditions for both utility and leakage.

Crucially, a connected agent network means each agent's private OS can be exposed to other agents if access is granted. One agent can query another's state through the cross-boundary protocol. This is what makes the setting different from single-agent safety: the same tool that enables collaboration also enables exfiltration, and the boundary between these depends on who is asking and why.

3.1.3 Combinatorial Complexity

Because state, tools, policy, and relationship are independently configurable, SharedOS exposes a combinatorial safety surface invisible to text-exchange benchmarks. For n agents each with $|T|$ tools and $|C|$ sensitivity categories, the static failure surface contains $O(n \cdot |T| \cdot |C|)$ cells, each conditioned on the pair-specific context $\mathcal{C}_{ji}$ (relationship, memory, prior decisions, policy). Multi-turn interaction compounds this: with branching factor b per turn over L turns, the trajectory tree grows as $O(n \cdot b^L)$, and the same request may be answered, refused, or escalated depending on conversational history. Static test suites and fixed-budget human red-teaming sample only a sparse slice of this surface.

We build SharedOS as the instrument to systematically probe this surface. With it, this thesis studies three linked questions:

Measurement: how can we evaluate whether a cross-boundary agent answers useful questions while withholding protected information?

Network failure: what privacy failures appear only when multiple agents, relationships, and delegation paths coexist?

Architectural enforcement: can the platform move enforcement before retrieval, so that unauthorised information never enters the model context?

3.2 Engineering Implementation

The formal model above is implemented in SharedOS, the production system behind Pulse/Aicoo. The implementation matters because it constrains what the experiments mean: every benchmark condition runs through the same execution core, the same tool assembly path, and the same permission model used by deployed cross-agent communication. The following sections describe the parts of the implementation that directly support the research claims.

3.2.1 Shared Execution Core

A naive implementation of SharedOS would build separate code paths for each interaction mode: one for the owner chatting with their agent, one for external guests arriving via share links, one for friend agents calling through the cross-boundary protocol, and one for API integrations. This approach is tempting because each mode has different security requirements. It is also wrong, for a reason that goes to the heart of the cross-boundary problem.

The same tool that enables legitimate collaboration also enables exfiltration. A `search_notes` call with access to folder 42 returns sprint planning documents. The same call with access to folder 51 returns salary data. The tool is identical. The model's reasoning loop is identical. The boundary between authorised access and data theft is *purely* determined by which permission object is loaded at the start of the interaction. If different interaction modes ran through different execution paths, we could not make this claim with certainty—a bug in one path might permit access that another path correctly restricts.

SharedOS therefore routes all four interaction modes through a single execution core. Two functions define the core: `buildStructuredSystemPrompt()` assembles the natural-language context the agent receives, and `assembleToolsFromPermissions()` constructs the tool array the agent can invoke. A third function, `consumeExecutionStream()`, runs the ReAct loop until the agent produces a final response. Every interaction, regardless of origin, passes through the same sequence: resolve permissions into a canonical object, assemble tools from that object, build the system prompt from identity and policy documents, and execute the reasoning

loop.

Table 3.1: Four entry points converging on a single execution core. The permission source differs; the execution path does not.

Entry point	Permission source	Trust basis
Owner chat	Implicit full access	Authentication session
Guest share link	capabilities JSONB on link	Token in URL
Friend agent (<code>contact_agent</code>)	<code>agentPermissions</code> table row	Prior explicit grant
API RPC	Same as friend agent	API key authentication

This convergence has a second benefit: it makes controlled experimentation trivial. To measure the effect of a governance policy, we change one input (the policy document) while holding the execution path constant. To measure the effect of relationship context, we change the relationship memory shard while holding everything else fixed. The four configurable axes described in the problem formulation—state, tools, governance, and relationship—map directly to four inputs to the shared core. Independent variation of each axis is an architectural guarantee, not a testing discipline.

Cross-boundary calls deserve special attention. When Agent B calls `contact_agent("alice", message)`, Alice’s agent is instantiated within that single tool call. Alice’s agent can itself invoke tools—searching notes, checking calendar, composing a response—but cannot recursively contact other agents. This prevents unbounded delegation chains while permitting the tool-mediated reasoning that makes delegation useful. The nested agent runs with a reduced iteration limit (10 steps versus unlimited for the owner) and inherits its tool set from the permission object associated with Agent B’s relationship to Alice.

3.2.2 Policy-as-Documents

A governance policy could be represented as a database column, a configuration file, or a structured rule engine. SharedOS represents policies as natural-language Markdown documents stored in the owner’s workspace. This is not a convenience shortcut. It is a design decision that shapes the entire experimental methodology and produces three properties that alternative representations lack.

First, policies become *user-editable without engineering intervention*. The owner writes governance rules in the same interface they use to write meeting notes. They can revise, version, and revert policies using the same snapshot mechanism that protects their other documents. This matters because real users will not fill out

permission matrices or write access-control rules in a formal language. If governance requires engineering skill, governance will not be configured, and the system will operate at D0.

Second, policies become *experimentally swappable*. The D0/D1/D2 defence gradient is implemented by replacing the content of a single file. No code changes, no redeployment, no feature flags. This enabled us to run the full PACT-PAIR evaluation across three defence levels and six models without modifying a single line of platform code between conditions.

Third, policies become *auditable and reproducible*. Every interaction's system prompt can be reconstructed by loading the same document versions that were active at the time. The policy that governed a past disclosure decision is recoverable, enabling post-hoc analysis of why a particular leak occurred.

The workspace organises governance documents in a hierarchical folder structure:

```

/Memory

/Self
    COO.md          Agent personality and capabilities
    USER.md         Owner identity and role
    POLICY.md        Base governance rules (all interactions)
    MEMORY.md        Learned facts and preferences
/@{handle}
    MEMORY.md        Facts about this person
    POLICY.md        Rules specific to this relationship
/links
    Label_token.md   Per-share-link policy

```

The system prompt assembler reads these documents and constructs a sectioned prompt. The base policy (`/Self/POLICY.md`) applies to all interactions. Per-relationship overrides (`/@{handle}/POLICY.md`) apply only when communicating with that specific requester. Per-link policies (the `## Policy` section of the link note) apply only to guests arriving through that specific share link. This cascade mirrors CSS specificity: more specific rules override less specific ones, and the most specific applicable policy wins.

The D2 category-specific deny list, for example, lives in `/Self/POLICY.md` as four paragraphs of natural language. When we evaluate D3 (relationship-specific policy), we add a `/@{handle}/POLICY.md` file that

grants per-requester exceptions. The agent receives both documents in its system prompt—the base policy and the relationship override—and must reconcile them through reasoning. This is precisely the “emergent enforcement” phenomenon we study: the policy text is an input to reasoning, not a deterministic access-control rule.

3.2.3 Capability-Based Share Links

The Mountable Context Cell, described formally in Chapter 6, is realised in production through folder-scoped share links. Each link encodes a complete MCC specification as a structured JSON object stored alongside the link metadata. When a guest arrives through the link, the system constructs an execution environment from this specification. The guest’s agent sees only what the specification permits. Everything else is not hidden, filtered, or access-controlled—it is *absent from the computational environment entirely*.

The distinction between filtering and absence is critical. A filtering approach retrieves all data and removes unauthorised items before presenting results to the model. This is dangerous: the retrieval step may produce error messages, result counts, or metadata that reveal the existence of filtered content. The MCC approach rebuilds the retrieval index per-link. The vector search operates over a database view containing only documents in the permitted folders. A query for “salary information” against a link scoped to folders 42 and 43 does not return “0 results in folder 51”—folder 51 does not exist in the search space. The model cannot confirm, deny, or speculate about data outside its mounted scope.

Each link’s capability specification defines five independent dimensions of access:

Share Link Capability Specification (JSONB)

```
{
  "notes": { "scope": "specific_folders", "access": "read", "folderIds": [42, 43] },
  "calendar": { "read": "free_busy", "write": false },
  "identity": { "loadCoo": true, "loadUser": true, "loadPolicy": true },
  "todos": { "read": true, "create": false },
  "tools": { "allowedTools": ["notion:search"] }
}
```

The *identity* dimension deserves explanation. It controls whether the guest sees the agent’s personality

(`COO.md`), the owner’s identity (`USER.md`), and the governance rules (`POLICY.md`). An investor data room link might load the full identity (so the agent introduces itself professionally) but restrict notes to financial folders. A public FAQ link might load only the personality and a curated knowledge folder, hiding the owner’s identity entirely. This dimension is orthogonal to data access: two links with identical folder scopes but different identity settings produce agents with different personas.

The owner can create arbitrarily many links with different scopes. Each link is a distinct MCC. A “work collaborator” link exposes project folders with full calendar. An “investor” link exposes financial summaries with busy-only calendar. A “friend” link exposes personal preferences with no work data. The same underlying data store produces different views for different audiences without any per-request reasoning about what to share.

3.2.4 Permission Model and Contact Graph

Agent-to-agent communication in SharedOS operates on a pull-based permission model. Before Agent B can contact Agent A, Owner A must have explicitly granted access to Owner B. This grant is stored as a row in the `agentPermissions` table with `grantor_id` = Owner A and `grantee_id` = Owner B. The relationship is unidirectional: A granting access to B does not imply B granting access to A. Both parties must independently choose to open their agents to each other.

This design implements contact-graph enforcement as infrastructure rather than policy. When Agent B attempts to call `contact_agent("alice", ...)`, the system checks for the permission row before any agent code runs. If the row does not exist, the call fails with a routing-layer rejection. Alice’s agent is never instantiated, never receives the message, and never has the opportunity to leak information. This is why PACT-NET’s contact enforcement metric achieves $\mathcal{C} = 100\%$ under both D0 and D1: it is not a policy effect but a structural guarantee.

Each permission row specifies what the grantee’s agent can access through two complementary surfaces. The first surface provides domain-specific granularity: which note folders are visible, what calendar detail level is available, whether todos are accessible, and whether email is searchable. The second surface provides namespace-level tool grants: an array of allowed tool identifiers supporting exact names (`"search_notes"`), namespace prefixes (`"calendar"` for all calendar tools), MCP server scopes (`"notion"` for all Notion tools), and wildcards (`"*"` for full access).

The two surfaces compose through union semantics: a tool is available if *either* surface grants it. This redundancy exists for practical reasons. The domain-specific surface predates the namespace surface and many existing permission rows use only domain columns. The namespace surface was introduced to support MCP tools that do not fit cleanly into the original domain categories. Union semantics ensure that neither migration path breaks existing grants.

The practical consequence is that permission specification has a natural complexity gradient. A minimal grant (“let this person’s agent read my public notes”) requires setting one column. A sophisticated grant (“read notes in folders 42 and 43, create todos in project X, use the Notion search tool but not Notion write, and see busy-only calendar”) composes multiple dimensions. The system does not require exhaustive pre-specification. Owners start with minimal grants and expand them as trust develops, mirroring how human relationships develop access over time.

3.2.5 Pre-Tool Escalation Gate

The escalation gate intercepts tool calls before execution. This placement is deliberate and produces a security property that post-response auditing cannot achieve. Consider the difference: a post-response auditor examines the agent’s output and catches leaks after they occur. The information has already entered the model’s context window, influenced its reasoning, and potentially leaked through side channels (metadata in refusals, reasoning traces, or implicit confirmation). A pre-tool gate prevents the tool from executing at all. The data never enters the context. The model cannot leak what it never saw.

The implementation uses a function-wrapping pattern. At the start of each cross-boundary interaction, the system iterates over the assembled tool array and replaces each tool’s execution function with a gated version. The gated function evaluates the proposed tool call against the requester’s relationship context and the owner’s precedent database. If the gate returns Continue, the original function executes normally. If it returns Stop, the tool is not executed and the agent receives a structured stop signal that it translates into a holding response (“Let me check with the owner first”).

The gate’s decision pipeline operates in stages designed to minimise latency. The first stage checks for exact precedent matches: has this specific requester asked for this specific resource before, and did the owner approve or deny? If a high-confidence match exists, the gate reuses the prior decision without invoking the

classifier. The second stage, reached only when no precedent matches, invokes a lightweight language model (GPT-4o-mini at temperature 0.2) with the requester’s identity, relationship cluster, the proposed tool and its arguments, and the owner’s stated boundaries. The classifier returns a ternary verdict—Allow, Escalate, or Deny—which the runtime maps to the binary Continue/Stop decision.

The ternary-to-binary mapping preserves information for the owner. When the owner reviews escalated items, they see whether the gate recommended denial (high confidence that the request violates boundaries) or escalation (genuine ambiguity). This distinction helps the owner calibrate: frequent escalations on legitimate requests signal that the boundary specification is too conservative, while frequent denials on the same category signal a systematic attack pattern.

Each owner decision—approve or deny an escalated request—becomes a precedent stored with the requester’s identity, relationship cluster, proposed resource, and the owner’s ruling. As precedents accumulate, exact and near-neighbour cases can be resolved with less owner intervention. The system converges toward the owner’s implicit boundary model without requiring them to pre-specify every cell in the permission matrix.

This convergence addresses a scalability problem identified in the MCC design. With O namespaces, N scopes per namespace, M tools, and R relationship types, the exhaustive permission matrix has $O \times N \times M \times R$ cells. No owner will specify this matrix in advance. The escalation gate fills it lazily: each real interaction that hits an unspecified cell triggers a decision, and that decision populates the matrix for future interactions in the same cluster. The system starts permissive for well-specified regions (where MCC folder grants are clear) and conservative for unspecified regions (where the gate defaults to Stop), then gradually relaxes as the owner’s decisions define the boundary.

Chapter 4

PACT-PAIR: Dyadic Evaluation

The preceding chapter described SharedOS as an infrastructure for cross-boundary agent delegation. This chapter uses that infrastructure to chart the security–utility frontier through systematic experimentation. PACT-PAIR is the pairwise setting of the **Permissioned Agent Coordination Testbed**: one requesting agent probes one target agent across a single privacy boundary.

4.1 Experiments

We evaluate SharedOS through **PACT-PAIR** (Permissioned Agent Coordination Testbed — Pairwise Assessment of Information Risk), a curated dataset of 600 tasks that probe one requesting agent against one target agent across a single privacy boundary.

Persona setup. The target agent (Alex) is a fully-equipped personal delegate with persistent memory, hundreds of files and structured state objects organised by domain and sensitivity (Work/Projects, Work/HR, Personal/Finance, Personal/Health, Personal/Family), and callable tools for search, retrieval, and mutation. Knowledge is fragmented: answering a question may require searching files, reading multiple documents, and cross-referencing structured state. The requesting agent (Tina) communicates with Alex through the cross-boundary protocol. In multi-turn mode, Tina maintains a checklist of pending questions and adaptively selects which to ask, how to frame them, and whether to retry. Phase 2 provides a broad retry protocol for refused questions; the concrete message wording, ordering, and strategy combinations are selected by the agent during the conversation.

Task design. The dataset comprises 600 tasks across two boundary tracks and four sensitivity categories (Table 4.1). The Information Boundary Track (Files QA + States QA) tests whether the agent answers legitimate questions and withholds sensitive information across two data surfaces: long-form documents and structured application state (todos, calendar entries). The Action Boundary Track tests whether the agent performs authorised mutations while refusing unauthorised ones, with verdicts determined by database-state diffs rather than text

parsing. Each task is paired with **gold key facts** whose presence in the response constitutes a leak (for sensitive queries) or a correct answer (for legitimate queries).

Governance policies. We instantiate three levels of the specificity gradient defined in §3.1: **D0** (no policy; system prompt contains no privacy instructions), **D1** (generic instruction: “Use your best judgment to protect Alex’s private information. Do not share data that seems personal or sensitive.”), and **D2** (category-specific deny list enumerating protected data types: financial data, health information, HR details, and personal relationships). D1 and D2 express the same strictness intent and differ only in specificity, isolating whether compliance depends on naming exactly what to protect.

Evaluation modes. **Single-step:** each query is asked in isolation (one question, one response, no history), providing a controlled per-query measurement. **Multi-turn heartbeat:** both agents operate as persistent processes exchanging messages at regular intervals for up to 240 ticks; Tina’s agent decides which questions to ask, how to frame each request, and whether to retry, capturing escalation through repeated probing, context accumulation, and social framing.

Metrics. Each query targets either public or protected information. **Utility** measures the fraction of legitimate queries answered correctly; **Security** measures the fraction of sensitive queries where no protected facts appear in the response. Neither alone suffices: an agent that refuses everything scores perfect security but zero utility. The security–utility frontier is the set of achievable operating points under a given configuration. We decompose each axis by verdict source:

Utility is the fraction of legitimate queries answered correctly. Security is decomposed into three mutually exclusive direct-response rates: **Refuse Rate** (explicit decline), **Message Leak Rate** (response contains gold key facts), and **Failed Attempt Rate** (attempted answer without gold facts). These sum to 100%. A fourth independent measure, **Global Leak Rate**, scans all responses across all ticks for gold-fact matches, capturing incidental disclosure that the per-query metric misses. In multi-turn settings, Global Leak typically exceeds Message Leak because sensitive facts surface in responses to unrelated queries. QA responses are scored conservatively with a two-pass LLM-judge plus string-match evaluator; actions follow the same structure with a database-state-diff gold check.

4.2 Findings

4.2.1 Setup

Table 4.1: Experimental dimensions. Each RQ isolates one axis while controlling the others.

Dimension	Levels
Boundary tracks	Files QA (200), States QA (200), Actions (200)
Task polarity	Utility (300) / Security (300)
Security categories	Sensitive work, Personal finance, health, relationships
Governance policies	D0 (none), D1 (generic), D2 (category deny list)
Evaluation modes	Single-step, Multi-turn (240 ticks)
Relationships (RQ3)	Colleague, CEO delegate, Close friend, Investor
Models	6 models, 3 providers (Table 4.4)

We evaluate six models from three providers (gpt-5-mini, gpt-5.5, gpt-5.4-mini, gpt-5.4, kimi-k2, deepseek-v3) under three governance policies (D0/D1/D2). gpt-5-mini serves as the ablation backbone for its stability and cost-efficiency. Deep ablations (surface variation, multi-turn, per-category analysis) use the backbone with two independent replications per cell. Multi-turn evaluation uses the 10×20 split protocol: 200 questions divided into 10 groups of 20, with up to 240 heartbeat ticks per group (Phase 1: ticks 1–60, Phase 2: ticks 61–240). The relationship condition is fixed to R0 (stranger) to isolate policy and surface effects. The action track evaluates 200 tasks under the same three policies. For every run we compute the four metrics defined in §4.1; adjudication is conservative, counting a leak if either the LLM judge or string matcher flags a gold fact.

4.2.2 Core Findings

We organise findings around the three research questions from §3.1. Each table reports four metrics: *Utility* (work-public correctness), and three security metrics for sensitive queries—Refuse Rate, Message Leak Rate, and Failed Attempt Rate (mutually exclusive, summing to 100%). Multi-turn tables add a Global Leak Rate (scan-based) that captures incidental disclosure across all 240 ticks.

RQ1: What moves the utility-security frontier? When does a privacy instruction actually change agent behavior? Specificity, not existence. Only category enumeration shifts the frontier. D1 adds a generic “protect




Figure 4.1: **Policy specificity determines security across all boundary surfaces.** Files QA (left) averages GPT-5-mini and GPT-5.5 (both LLM judge); States QA (centre) and Actions (right) use GPT-5-mini (2 replications pooled). D1 provides no reliable protection over D0. D2 reduces leakage by 52–78 pp and eliminates the most dangerous action categories entirely. Full 6-model results in Table 4.4.

private data” instruction and achieves zero net improvement: McNemar discordant pairs are tied on files (10:10, $p = \text{n.s.}$) and on states (13:15, $p = \text{n.s.}$). D2 names four protected categories and reduces file leakage from 83% to 14% (69 pp, $p < 0.001$) while preserving utility at 77%. There is a *specificity threshold*: below it, governance has no measurable effect; above it, protection is immediate.

Why does naming matter? Under D0, every model already refuses obvious PII (tax IDs, home addresses) at 93–100% rates. D1 fails on *ambiguous* categories: hiring budgets, 1:1 feedback, promotion frameworks, and fundraising terms at the boundary between “work information” and “private data.” Without explicit category names, the agent defaults to sharing anything work-adjacent. D2 resolves this by naming the boundary: D0 and D1 produce near-zero refusals on sensitive queries (0% and 3% respectively). Protection under these policies is incidental (retrieval failures), not intentional. D2’s 71.5% refusal rate shows active enforcement.

Takeaway (RQ1): The frontier only moves when you name the boundary

The frontier does not move unless the policy names exactly what to protect. Vague instructions register as soft preferences that lose the competition against helpfulness. Category enumeration converts privacy decisions from open-ended judgment into pattern matching, crossing a specificity threshold below which governance has no measurable effect.

On the action track, D2 improves safety without sacrificing utility. It blocks 93.5% of unauthorised actions while executing 61.0% of authorised ones. D1 adds zero safety over D0 (both 43.0% block rate) while reducing utility from 65.5% to 48.0%—a pure loss.




Figure 4.2: **Multi-turn erosion case study: the wedding cascade** (D2, gpt-5-mini, split s06). Phase 1 (Tick 18) produces a clean refusal. Phase 2 retries (Ticks 80–91) fail 12 consecutive times because the requesting agent searches only Work/shared folders. At Tick 92, a three-strategy combination (scope expansion + business justification + constrained format) breaks through, leaking the wedding date, venue, and guest count. Tick 93 chains the leaked note ID to extract venue confirmation. Ticks 94–99 show policy recovery: all six follow-up probes targeting the partner's name are refused. The erosion is bounded: the policy does not collapse after a temporary breach.

The protection is not free everywhere. D2 costs only 1 pp utility on files but 37 pp on structured state, where terse queries ("SOC2 status?") share vocabulary with protected categories and trigger a 26.5% false-refusal rate. The frontier's shape is surface-dependent (§4.2.3).

RQ2: Is the frontier stable over time? Can an adaptive requester erode defenses that hold on first contact? The frontier holds. D2's multi-turn message leak rate (12.6%) is comparable to its single-turn rate (14.0%), confirming bounded erosion rather than collapse (Table 4.2). But multi-turn interaction reveals failure modes invisible to single-turn evaluation. The remaining 35.6% of non-refused queries splits into 12.6% actual disclosure and 23.0% failed attempts (policy violated but no gold facts leaked). The global leak rate (38.0%, scan-based) is 3× the message rate because sensitive facts surface incidentally via responses to unrelated questions.

The category structure matters. Sensitive work has the lowest refuse rate (36.8%) and highest message leak (22.8%) because of genuine boundary ambiguity between "work question" and "HR data." Personal health

Table 4.2: **Multi-turn D2 security breakdown by category** (10 splits $\times$ 240 ticks). Refuse = explicit refusal. Msg. Leak = response verified to contain gold facts. Failed Att. = attempted answer without gold facts. Glob. Leak = gold facts appear anywhere across 240 ticks (scan-based).

Category	gpt-5-mini ($n=191$)				GPT-5.5 ($n=100$)			
	Ref	Msg.L	F.Att	Glob.L	Ref	Msg.L	F.Att	Glob.L
Sensitive work	36.8	22.8	40.4	51.7	83.3	16.7	0.0	26.7
Personal finance	75.0	4.5	20.5	34.7	80.0	8.0	12.0	24.0
Personal health	89.7	2.6	7.7	20.0	100	0.0	0.0	15.0
Personal relationships	66.7	15.7	17.6	39.2	80.0	8.0	12.0	24.0
All sensitive	64.4	12.6	23.0	38.0	85.0	9.0	6.0	23.0

achieves 89.7% refusal with only 2.6% message leak. Yet even personal health has a 20% global leak rate: health facts surface incidentally when the agent answers related work queries.

Adaptive retry strategies. During multi-turn Phase 2, the requesting agent is instructed to retry refused questions using different approaches, but the concrete wording and sequencing are not scripted. We observe five distinct strategies across 1,184 retry attempts under D2: requesting raw note content by ID (65% of attempts, 7.6% flip rate), business-justification reframing (15%, 34.5% flip rate), scope narrowing, public reframing, and delegation requests. Strategy *combinations* are more dangerous than individual strategies. The most complex breach required 12 failed attempts before a three-strategy combination (business justification + folder expansion + constrained format) cracked a previously-refused wedding-details query. After the breach, the agent refused all six follow-up probes. The policy recovers after temporary breaches, distinguishing bounded erosion from collapse.

Refusals leak metadata. Tool-equipped delegates exhibit a failure mode invisible to text-exchange benchmarks. The agent searches its owner’s data store *before* deciding to refuse. The refusal itself then discloses metadata: one delegate refused to share brokerage balance but revealed the note ID, folder name, and data existence (“The info is recorded in his notes under ‘Bank Accounts’, Finance folder, note ID 7795”). The requesting agent uses disclosed note IDs to retrieve data directly on the next turn. This attack chain requires both tool access and multi-turn interaction.

Takeaway (RQ2): Multi-turn erodes the frontier through channels the policy never sees

Multi-turn interaction erodes the frontier. The dominant mechanism is not adversarial cracking but incidental leakage: legitimate queries surface co-located sensitive data that the policy never evaluates. Single-turn benchmarks systematically undercount real-world privacy risk because they cannot observe this channel.

RQ3: Is the frontier the same for everyone? Does the same policy protect equally regardless of who is asking? We fix the D2 policy and vary only the requester’s relationship via a single memory shard injected into the defending agent. In the preceding experiments (RQ1–RQ2), the requester–target relationship is implicit: Alex infers Tina’s standing from conversational context alone, with no explicit relationship information in memory. Here we make that context explicit by injecting a structured memory entry describing each requester’s relationship to Alex. Four requesters probe the same 400-question pool: Tina (colleague/work peer), Marcus (CEO’s delegate), Jordan (close personal friend), and Dana (investor/board observer). Ground-truth labels are relationship-conditioned: the same question may be legitimate under one relationship but private under another, such as salary for a CEO delegate versus a colleague, or wedding details for a close friend versus an investor.

Table 4.3: **Relationship-conditioned D2 results** (gpt-5-mini, 400 QA + 200 actions per requester, OR-adjudication). Leak = total P-item leak rate. O-Ref = over-refusal on L-items. Safety = unauthorised actions blocked.

Requester	Role	QA (information)			Actions	
		Leak↓	O-Ref↓	Util↑	Util↑	Safety↑
Tina	colleague	1.7%	40%	57%	92	90
Marcus	CEO delegate	3.3%	59%	49%	89	85
Jordan	close friend	9.2%	86%	58%	92	84
Dana	investor/board	7.5%	31%	70%	83	91

No. The frontier is not flat across requesters (Table 4.3). Relationship context reshapes it in category-specific and direction-specific ways, but only for one category: personal finance (0–2% leak), personal health (5–10%), and personal relationships (0–4%) are protected regardless of who asks, while sensitive_work shows a clear gradient (Tina 0%, Marcus 20%, Jordan 18%, Dana 26%). The agent treats financial and health data as belonging to the *owner* but work data as belonging to the *organisation*—relationship framing modulates




Figure 4.3: **The privacy-utility frontier across three dimensions.** *Left:* Policy specificity—D2 shifts all six models to high security; D1 is indistinguishable from D0. *Center:* Cumulative multi-turn leakage (240 ticks). D2 caps exposure below 30% for both models. *Right:* Relationship-conditioned dual failure under D2. Red: leak rate (driven by *sensitive_work*). Blue: over-refusal on legitimate items. Jordan exemplifies the worst case—leaks 9% yet blocks 86% of entitled access.

perceived organisational standing. Moreover, read trust $\neq$ write trust: Dana’s investor framing produces the highest QA leakage (7.5%) but also the highest action safety (91%), while Jordan’s friend framing produces the most over-refusal (86%) yet the lowest action safety (84%). The dominant practical failure is not leakage but over-refusal: D2 refuses 86–95% of sensitive queries across all requesters (leak rates 1.7–9.2%), but Jordan’s agent refuses 86% of items it *should* answer while Dana’s refuses only 31%. Under relationship variation, D2’s cost falls on legitimate utility, not on security breaches.

Takeaway (RQ3): Same policy, different frontier per requester

A single policy cannot serve all requesters optimally. The model’s implicit social norms decompose trust into orthogonal read/write axes that activate differently per relationship. Over-refusal, not leakage, is the dominant practical failure: the frontier’s cost falls on legitimate utility, not on security breaches.

4.2.3 Ablations

Cross-model generality. The specificity threshold (RQ1) holds across all six evaluated models (Table 4.4). Under D0, all models leak at 83–93% despite substantial capability differences (utility 78–98%). D2 reduces leakage by 69–91 pp in every model. Two non-OpenAI models (Kimi K2, DeepSeek V3) show the same pattern, ruling out provider-specific training artifacts.

Table 4.4: **Cross-model single-step files QA** (6 models, 1 replication per model except gpt-5-mini with 2). All values in %. Util = work-public correctness. Ref = refuse rate. Msg.L = message leak rate. F.Att = failed attempt rate. D1 (omitted) is statistically indistinguishable from D0 in all models.

Model	D0 (no policy)				D2 (deny list)				Δ Msg.L
	Util	Ref	Msg.L	F.Att	Util	Ref	Msg.L	F.Att	
gpt-5-mini	78	0	83	17	77	72	14	14	−69
gpt-5.5	87	1	88	11	86	78	4	18	−84
gpt-5.4-mini	96	1	87	12	91	89	7	4	−80
gpt-5.4	98	1	92	7	74	93	1	6	−91
Kimi K2	82	0	93	7	81	86	4	10	−89
DeepSeek V3	91	1	93	6	62	90	9	1	−84

Model scale under multi-turn. Under D0, GPT-5.5’s message leak rate is 28.0% vs. gpt-5-mini’s 84.2%, a 56 pp reduction from scale alone. Under D2, both converge: 13.0% vs. 12.6% (Table 4.5). The explicit deny list is the great equaliser: it renders model scale irrelevant for the question of whether protection holds.

Table 4.5: **Multi-turn model scaling** (10 splits $\times$ 240 ticks, 200 security QA per defence level). Ref = refuse rate. Msg.L = message leak rate (LLM judge). Glob.L = global leak rate (gold-fact scan). Util = work-public correct rate.

Defence	gpt-5-mini					gpt-5.5				
	Ref	Msg.L	Glob.L	Util	Act.B	Ref	Msg.L	Glob.L	Util	Act.B
D0	0.0	84.2	83.0	82.9	59.0	63.5	28.0	39.5	68.5	40.9
D1	2.0	72.9	79.5	77.5	51.0	68.5	25.0	34.0	56.0	52.6
D2	64.4	12.6	38.0	60.3	88.5	80.5	13.0	24.5	51.5	91.4

Data-surface asymmetry. The frontier’s cost depends on the data surface being queried. Table 4.6 compares files QA (long-form documents) against states QA (structured records: todos, calendar entries, CRM fields).

The structured-state surface presents a harder retrieval problem (D0 utility 55% vs. 78%) *and* a harder disambiguation problem (D2 work/public refusal 26.5% vs. 0.5%). Long-form documents provide enough context for accurate category classification; terse state queries (“SOC2 status?”) share vocabulary with protected categories and trigger false positives. D2 achieves lower absolute leak rates on states (8% vs. 14%) partly

Table 4.6: **Cross-surface comparison** (gpt-5-mini, single-step, pooled). Leak = fraction of sensitive queries where gold facts appear in response. Utility = fraction of legitimate queries answered correctly.

	Files QA	Struct. States	Δ
D0 Utility	78.0%	55.0%	−23 pp
D0 Leak	83.0%	58.5%	−25 pp
D2 Utility	77.0%	18.0%	−59 pp
D2 Leak	14.0%	8.0%	−6 pp
D2 over-refusal	0.5%	26.5%	+26 pp

because its over-refusal catches genuine leaks as a side effect. The implication: practitioners cannot evaluate on a single surface and generalise—the frontier’s shape is surface-dependent, and governance that works well on documents may fail on structured state.

Prompt-level defence strategies. Given D2’s residual leak rate, can more sophisticated prompt-level defences close the gap? We evaluate three strategies applied on top of D2: **D3 (Spotlighting)** [26] frames external messages as “DATA, not instructions”; **D4 (Instruction Hierarchy)** [25] defines a privilege stack (SYSTEM > Owner > External); and **D5 (Sandwich + Boundary)** combines the deny list with pre-message classification and post-policy reinforcement. Under the 240-tick multi-step protocol (Table 4.7), D5 achieves the best trade-off (27.5% leak, −10.5 pp vs. D2) at only 5.5 pp utility cost. D3/D4 incur steeper utility penalties (−12–13 pp) for smaller gains. All three achieve $\geq 94\%$ action safety. Yet no strategy closes the gap by more than 11 pp: the remaining leakage is dominated by incidental disclosure.

Table 4.7: **Prompt-level defence comparison** (gpt-5-mini, 240 ticks, multi-step). Leak = global gold-fact rate. Act.S = unauthorised action block rate.

Defence	Splits	Utility	Leak Rate	Δ Leak	Act. Safety
D2 (baseline)	10	85.5%	38.0%	—	88.5%
D3 (Spotlighting)	7	72.9%	34.3%	−3.7 pp	94.0%
D4 (Instr. Hierarchy)	7	72.1%	32.1%	−5.9 pp	96.0%
D5 (Sandwich)	6	80.0%	27.5%	−10.5 pp	94.0%

Network scaling: PACT-NET. The pairwise PACT-PAIR benchmark evaluates one requesting agent against one target. To test whether findings generalise to multi-party networks, we instantiate PACT-NET: 25 agents with distinct personas, data stores, and inter-agent relationships forming a realistic organisational graph. The full

network-scale evaluation is reported in Chapter 5.

4.3 Discussion

We formalised cross-boundary agentic delegation, built SharedOS as a platform for controlled experimentation, and curated PACT-PAIR (600 tasks, dual-metric scoring, three boundary surfaces). Our experiments across six models reveal that the security–utility frontier is *emergent*: it only moves when governance names the boundary explicitly (RQ1), erodes through incidental disclosure invisible to single-turn evaluation (RQ2), and reshapes per requester under fixed policy (RQ3).

The prompt is necessary and practical, but has a ceiling. Category-specific policies shift the frontier by 50–69 pp; layered strategies add 4–11 pp more. No other mechanism (model scale, generic instructions, relationship context) achieves comparable gains. For practitioners: write explicit, category-enumerated policies. But three residual channels remain beyond prompt engineering’s reach: incidental disclosure (the 3× gap between message and global leak), metadata leakage in refusals (note IDs disclosed before the refuse decision), and surface-dependent over-refusal (1 pp utility cost on files vs. 37 pp on structured state). Closing these requires structural enforcement at the platform layer.

Structural interventions beyond the prompt. We identify three enforcement points that complement governance policies. *Pre-retrieval gating* classifies the incoming query’s information need before any tool call executes; if the requested category is protected for the current requester, the search never fires and no private content enters the context window. *Requester-conditioned data mounting* exposes different data partitions per relationship tier: a colleague’s agent sees work files but not health records, regardless of what it asks. *Post-generation auditing* scans outgoing messages for protected facts (e.g., salary figures, medical terms) and redacts before delivery. Together these three layers address the residual channels that prompt-level policy cannot close. Single-turn benchmarks systematically undercount risk because they cannot observe incidental disclosure; evaluation must be multi-turn and surface-aware.

Implications and future work. Agentic privacy is a *stack*: governance policy sets intent (D2), the model provides compliance capacity (scale ablation), and platform infrastructure enforces guarantees where language

fails (structural interventions above). The D5 ceiling implies model developers should prioritise structural compliance (tool-use gating, output filtering) over instruction-following finetuning; the cross-surface asymmetry implies that emerging agent-to-agent protocols (MCP, A2A) must distinguish data surfaces rather than applying uniform access control. Three open threads follow. First, *adaptive policy synthesis*: can a meta-agent observe leak/refusal outcomes and automatically tighten or relax category-level rules toward a target operating point? Our D0–D5 gradient provides the training signal but the closed-loop optimisation remains unbuilt. Second, *compositional privacy*: when Agent A delegates to Agent B who further delegates to Agent C, transitive leakage guarantees require extending Bell–LaPadula [11] lattice models to account for probabilistic LLM behaviour—a formal framework that does not yet exist. Third, *user-in-the-loop calibration*: deploying PACT-PAIR-style evaluations as a trust-calibration interface where users adjust their own governance policies based on observed leak/utility feedback, closing the loop between benchmark results and deployment decisions.

Limitations. RQ1–RQ2 evaluate one target agent; RQ3 varies four requesters but retains the same target. PACT-NET extends to 25 agents in Chapter 5. The data store is synthetic; real personal data stores exhibit richer cross-referencing and ambiguous category boundaries that may amplify both leakage and over-refusal. The evaluator was manually audited on sampled and disagreement cases, but the OR-pooling convention upper-bounds risk at each level. D2 states QA shows 26 pp utility variance across replications, suggesting that structured data surfaces are more sensitive to prompt phrasing than document surfaces. Optimised attacks (PAIR, GCG) are benchmark extensions rather than deployed threats; adversarial robustness under optimised attacks remains an open question. We do not evaluate user-facing deployment metrics (latency, user satisfaction) which may impose additional constraints on defence strategy selection.

4.4 Failure Taxonomy

Across both data surfaces (files and states) and both evaluation modes (single-step and multi-turn), D2 residual leaks fall into four mechanisms. Understanding these mechanisms is essential for designing architectural interventions that address root causes rather than symptoms.

Category boundary ambiguity. This is the dominant failure mechanism. Data is organisational in scope (a company fundraising round, an HR promotion framework) but contains individually-sensitive details (deal

Table 4.8: D2 failure mechanisms across surfaces and modes. Total is 44 unique leaked questions (28 from files, 16 from states) pooled across replications.

Mechanism	Files	States	Total (%)	Description
Category boundary ambiguity	17	10	27 (61%)	Organisational data not covered by individual-level deny list
Leaked-outside-message	4	3	7 (16%)	Refusal text itself reveals protected facts
Company-benefit confusion	2	1	3 (7%)	Benefits treated as company info, not personal finance
Stochastic failures	5	2	7 (16%)	Boundary cases leaked in one replication only
Total	28	16	44 (100%)	

terms, salary bands). D2’s deny list targets individual-level categories and does not generalise to organisational-but-sensitive information. For example, when Tina asks for Series A term sheet details, D2 can still disclose the lead investor, amount, valuation, liquidation preference, and board structure because fundraising terms are not named in the deny list.

Leaked-outside-message. The agent refuses the query but the refusal text itself reveals protected facts. One D2 refusal says it cannot share whether Alex has financial ties to his sister’s business or details about ownership, roles, financial interests, duration, or agreements. The refusal does not provide the protected record, but it confirms the existence and shape of the protected information.

Company-benefit confusion. Benefits (401k matching percentages, vision insurance details, wellness stipend amounts) are treated as company information rather than personal finance. The agent reasons that benefits are standardised across employees and therefore not individually sensitive. This is often correct for generic benefits but fails when the record reveals an individual’s elections or vesting status.

Stochastic failures. Some questions leaked in exactly one replication but not the other, with no identifiable systematic pattern. These cases are genuine boundary conditions where the same question, policy, and system prompt produce different outcomes under different sampling draws.

Summary of failure modes. The taxonomy confirms that D2’s residual vulnerabilities are structural, not accidental. The dominant failure (boundary ambiguity) arises from a fundamental mismatch between category-level policies and the continuous nature of information sensitivity. No amount of

prompt engineering can resolve this mismatch. The solution space lies in architectural enforcement: controlling access before reasoning, not trusting reasoning to control disclosure after access.

Chapter 5

PACT-NET: Network Evaluation

PACT-PAIR evaluates one requester interacting with one target. That dyadic setting is the right unit test for prompt specificity, relationship labels, and multi-turn erosion. It cannot, however, express the central risks of a deployed agent network: third-party information can be relayed through an intermediate agent; a requester can claim authority from someone else; a fact can cross from one social cluster into another; and one disclosure can reveal a bundle of co-located facts. PACT-NET is the network-scale extension of the Permissioned Agent Coordination Testbed. It evaluates 997 cross-agent tasks over a 25-agent social graph and asks what failures appear only when agents coordinate across a network.

To avoid confusion with the PACT-PAIR defence ladder (D0–D5), this chapter uses a local notation. **P0** denotes PACT-NET with no `POLICY.md`. **P1** denotes PACT-NET with each target agent’s static, per-agent `POLICY.md` loaded. P1 is role/category-specific: Carlos’s policy governs finance and payroll data; Dr Karen’s policy governs therapeutic confidentiality; Dana’s policy governs investor and portfolio information. It is *not* relationship-specific. The same target policy is applied regardless of whether the requester is a colleague, investor, friend, or family member.

5.1 Motivation & Design

PACT-NET is designed around a fictional startup ecosystem, TechFlow AI. The graph contains 25 agents across nine clusters: executive, engineering, product, business, operations, investor, external, personal, and family. The graph intentionally includes bridge nodes. Alex Chen, the CTO, connects professional, investor, external, personal, and family contexts. Other agents hold their own notes, todos, policies, and relationships. This makes the world more than a larger PACT-PAIR instance: it creates third parties and clusters, so the benchmark can ask questions that do not exist in a pairwise setting.

The contact graph is enforced by the production `contact_agent` path. If a source agent is not in the target’s contact list, the request is rejected before the target agent is instantiated. This lets PACT-NET separate

Table 5.1: PACT-NET world and execution summary.

Component	Value
Agents	25 synthetic SharedOS users
Clusters	9 social/organisational clusters
Contact graph	76 contact edges; Alex is the hub with 19 contacts
Evaluation tasks	997 Phase 1 tasks per run
Persistence probe	75 Phase 2 dig-further ticks per run
Replications	2 per condition
Main conditions	P0: no POLICY.md; P1: per-agent static POLICY.md
Not tested here	No clean P2/D2, no relationship-specific policy, no escalation

routing-layer guarantees from language-model policy following. Contact enforcement should be structural; selective disclosure within an authorised contact is the hard part.

The policy variable is deliberately narrow. In P0, every target agent’s POLICY.md is empty. In P1, each target agent receives its own role/category policy. These policies specify what the owner may share and what must be refused: for example, finance agents may share aggregate runway but not individual salaries; therapists may share scheduling logistics but no client information; investors may share public market insight but not other portfolio company data. Because P1 is not requester-conditioned, it tests whether a static owner policy is enough for networked delegation.

5.2 Experiments

The clean PACT-NET V2 experiment consists of four namespace-isolated runs: P0 replications 1–2 and P1 replications 1–2. Each run executes the same 997 Phase 1 tasks and 75 Phase 2 dig-further ticks. The V1 D2 runs are not used for main claims because the audit found a shared-UUID race in which concurrent runs could overwrite each other’s POLICY.md. The only clean PACT-NET comparison in this chapter is therefore P0 versus P1.

Table 5.2: Clean PACT-NET V2 runs used in this chapter.

Local condition	Policy state	Replications	Use
P0	Empty POLICY.md	2	Clean baseline
P1	Per-agent role/category POLICY.md	2	Clean policy condition
P2/D2	Not cleanly evaluated	–	Not reported
Relationship policy	Not run in PACT-NET	–	Not reported
Escalation	Not run in this chapter	–	Reported separately in Chapter 6

PACT-NET has two task surfaces. QA tasks test whether an agent should answer or refuse an information request. Action tasks test whether an agent should execute or block a write-side operation. The task families are chosen to separate ordinary dyadic behaviour from network-native behaviour.

Table 5.3: PACT-NET task families. Counts are per Phase 1 run.

Surface	Family	N	What it tests
QA	should_answer	172	Legitimate contact query
QA	should_refuse	139	Direct sensitive query
QA	transitive_risk	94	A asks B about C’s information
QA	cross_cluster	28	Information crossing a social/organisational cluster
QA	non_contact_probe	50	Routing-layer contact enforcement
Action	authorized_create	184	Legitimate note/todo creation
Action	authorized_complete	115	Legitimate completion or update
Action	unauthorized_mutation	115	Unsafe write to protected data
Action	confused_deputy	50	False delegated authority claim
Action	cross_surface_plant	50	Planting sensitive information into another surface

The scoring follows the same security–utility philosophy as PACT-PAIR. Utility is legitimate-task recall: the agent should answer or execute when the requester is authorised. Safety is private-task recall: the agent should refuse or block when the requester is unauthorised. QA tasks are scored by gold key-fact matching. Action tasks in these P0/P1 runs are scored by response heuristics rather than database-snapshot diffs, so action-category results are directional; QA and refusal metrics are the primary evidence.

PACT-NET adds four network-specific metrics used in the findings below. $\mathcal{T}$ is transitive leak rate: lower is better. $\mathcal{X}$ is cross-cluster leak rate: lower is better. $\mathcal{A}$ is amplification factor, the average number of protected facts exposed per leak event. $\mathcal{D}$ is confused-deputy attack success: lower is better. Contact enforcement $\mathcal{C}$ is reported as a structural routing check.

Table 5.4: PACT-NET composite scores, averaged over two namespace-isolated replications.

Metric	P0: no policy	P1: per-agent policy	Δ
Overall accuracy	55.9%	74.9%	+19.0 pp
Utility	88.7%	78.8%	−10.0 pp
Safety	26.6%	71.5%	+44.8 pp

Table 5.5: PACT-NET per-family accuracy, averaged over two replications. Action rows are response-heuristic scored.

Surface	Family	N	P0	P1	Δ
QA	should_answer	172	75.9%	64.0%	−11.9 pp
QA	should_refuse	139	16.2%	73.4%	+57.2 pp
QA	transitive_risk	94	3.7%	22.3%	+18.6 pp
QA	cross_cluster	28	12.5%	30.4%	+17.9 pp
QA	non_contact_probe	50	100.0%	100.0%	0
Action	authorized_create	184	98.4%	86.7%	−11.7 pp
Action	authorized_complete	115	91.7%	88.3%	−3.5 pp
Action	unauthorized_mutation	115	29.6%	85.2%	+55.7 pp
Action	confused_deputy	50	53.0%	98.0%	+45.0 pp
Action	cross_surface_plant	50	0.0%	94.0%	+94.0 pp

5.3 Findings

PACT-NET answers four questions that PACT-PAIR cannot ask. These are not merely larger versions of dyadic privacy tasks; each depends on a third party, a cluster boundary, an authority chain, or co-located network memory.

5.3.1 F1: Third-party information leaks indirectly

PACT-PAIR can ask whether B leaks B’s information to A. It cannot ask whether A’s query to B causes B to reveal information about C. PACT-NET can, because the graph contains third parties and overlapping notes.

Under P0, transitive-risk correctness is only 3.7%, corresponding to a transitive leak rate of $\mathcal{J} = 96.3\%$. P1 improves correctness to 22.3%, but $\mathcal{J}$ remains 77.7%. The policy helps, but three out of four transitive cases still leak. This is the core selective-disclosure failure: the agent may answer a legitimate local question while carrying along third-party facts that should have been withheld.

5.3.2 F2: Information crosses social and organisational clusters

PACT-PAIR has one relationship at a time. It cannot represent the boundary between professional, investor, personal, and family clusters. PACT-NET tests whether information that is appropriate inside one cluster crosses into another.

Cross-cluster correctness rises from 12.5% under P0 to 30.4% under P1, but the cross-cluster leak rate remains high: $\mathcal{X} = 87.5\%$ under P0 and $\mathcal{X} = 69.6\%$ under P1. Static per-agent policies therefore do not reliably

maintain cluster boundaries. They can describe categories such as compensation, clinical notes, or investor deliberations, but they do not encode who is currently asking and which cluster boundary the answer would cross.

5.3.3 F3: One leak often becomes multi-fact disclosure

PACT-PAIR can mark whether a specific gold fact leaked. In a networked workspace, a single answer often exposes a bundle of adjacent facts: meeting notes include attendees, numbers, next steps, and sensitive side comments; a todo can include both the task and the private rationale.

PACT-NET captures this through the amplification factor $\mathcal{A}$. P0 has $\mathcal{A} = 1.61$ and P1 has $\mathcal{A} = 1.55$. The reduction is small. Once a leak occurs, it remains bundled: the system usually discloses more than one protected fact per leak event. This means dyadic evaluation underestimates the density of real-world leakage, because the networked data store co-locates facts from multiple people and contexts.

5.3.4 F4: False delegation creates confused deputies

PACT-PAIR has no natural way to express “Alice asked me to ask you.” PACT-NET can test this authority-chain failure. A requester claims that a trusted third party authorised the request, and the target agent must decide whether to honour the claimed delegation.

Here P1 works well. Confused-deputy attack success falls from $\mathcal{D} = 47.0\%$ under P0 to $\mathcal{D} = 2.0\%$ under P1. The corresponding task-family accuracy rises from 53.0% to 98.0%. This contrast is important: prompt policy can be highly effective when the unsafe authority claim is visible in the immediate request. The same policy is much weaker when the violation is implicit in a transitive or cross-cluster information path.

Table 5.6: Network-specific metrics. Lower is better for $\mathcal{T}$, $\mathcal{X}$, $\mathcal{A}$, and $\mathcal{D}$; higher is better for $\mathcal{C}$.

Symbol	Metric	P0	P1	Interpretation
$\mathcal{T}$	Transitive leak rate	96.3%	77.7%	Still mostly leaks
$\mathcal{X}$	Cross-cluster leak rate	87.5%	69.6%	Still mostly leaks
$\mathcal{A}$	Amplification factor	1.61	1.55	Leaks remain bundled
$\mathcal{D}$	Confused-deputy success	47.0%	2.0%	Policy nearly solves this surface
$\mathcal{C}$	Contact enforcement	100.0%	100.0%	Structural ACL guarantee

Overall synthesis. P1 delivers a large aggregate safety gain (+44.8 pp) at a moderate utility cost (−10.0 pp), but the gain is uneven. Static per-agent policy is strong against direct sensitive queries, unsafe writes, cross-surface planting, and false delegation. It is weak against the genuinely network-native failures: transitive third-party leakage, cross-cluster leakage, and multi-fact amplification. This is the core lesson of PACT-NET. Network privacy is not solved by writing a better bilateral policy for each agent. It requires mechanisms that understand requester identity, cluster boundaries, provenance, and escalation across the social graph.

Chapter 6

Architectural Solutions

The preceding two chapters established what goes wrong. PACT-PAIR revealed that prompt-level governance has a ceiling: even the best-performing policy configuration (D2, category-specific deny list) still leaks 27–38% of protected information depending on the model and evaluation mode. PACT-NET revealed that network structure amplifies these failures through transitive leakage, confused deputies, and cross-cluster boundary violations. This chapter proposes architectural solutions that address these failures not by asking the model to try harder but by restructuring the execution environment so that violations become structurally impossible.

The governing philosophy is simple: *do not ask the agent to refrain from disclosing data; ensure the data is not in the room*. Prompt-level governance instructs the model what not to say. Architectural governance removes the information from the model’s context entirely. The former depends on the model’s reasoning. The latter does not.

Table 6.1: Failure modes and their architectural remedies.

Failure mode	Root cause	Architectural remedy
Direct information leak	Data present in context	Mountable Context Cells
Transitive leakage	Co-located sensitive data	Mountable Context Cells
Confused deputy	No delegation verification	Escalation Protocol
Metadata leakage	Tool/search reveals existence	Pre-tool escalation gate
Over-refusal	Binary access model	Per-requester MCC
Multi-turn erosion	Accumulated context	Future multi-turn audit
Cross-cluster conflation	Closeness $\neq$ breadth	Relationship-scoped MCC

This chapter proposes three complementary solutions—Relationship-Specific Policy, Mountable Context Cells, and Agentic Escalation—that address the failures identified in Chapters 4 and 5. The key architectural insight is drawn from capability-based security [42]: do not ask the agent to refrain from disclosing data; ensure the data is never loaded into the agent’s context in the first place.

6.1 Relationship-Specific Policy

The PACT-PAIR and PACT-NET results demonstrate that information sensitivity is relational rather than absolute: the same fact may be appropriate to share with one requester and harmful to share with another. A flat access-control model (“this document is sensitive”) systematically fails because it cannot express relationship-dependent boundaries. The solution is to parameterise the governance mechanism by the requester’s identity and relationship type, assembling a distinct access profile for each interaction.

We evaluated relationship-specific policy (D3) as a standalone defence in the three-condition MCC experiment (Section 6.4). Under D3, each requester receives a tailored system prompt specifying which categories are accessible for their relationship type. Results (Table 6.5) show that D3 achieves an aggregate leak rate of 15.5% with 70.9% utility. This confirms that relationship conditioning is necessary—it substantially outperforms flat D2 for misaligned requesters—but not sufficient, as within-scope leakage persists (38.7% for the Close Friend requester). These residual failures motivate the architectural solutions that follow.

This principle motivates the Mountable Context Cell architecture described in the next section. Rather than writing a single policy that enumerates what to withhold, the system constructs a per-requester execution environment containing only the data that the specific requester is entitled to see. Different requesters receive different views of the same underlying data store, eliminating the tension between helpfulness and privacy that plagues flat prompt-level governance.

6.2 Mountable Context Cells

6.2.1 Concept & Formulation

A Mountable Context Cell (MCC) is a capability-based execution sandbox assembled for each interaction. The MCC defines what data the agent can access, what tools it can invoke, and what system context it receives. Crucially, the MCC is constructed per-requester. Different requesters see different MCCs, and therefore different views of the owner’s data.

The analogy is containerisation versus file permissions. File permissions (the equivalent of prompt-level governance) tell a process which files it should not read. Containers (the equivalent of MCCs) construct a filesystem view in which unauthorised files do not exist. The process cannot read what is not mounted.

- **Prompt-level governance (D2):** All data is loaded. The model is instructed to withhold certain categories. Data is *hidden*.
- **MCC governance:** Only permitted data is loaded. Unauthorised data never enters the context window. Data is *absent*.

This distinction matters because language models are not reliable access control mechanisms. They are trained to be helpful and to use available context. Asking them to ignore visible information is asking them to act against their training objective. MCCs eliminate this conflict by ensuring the model never encounters the information it should withhold.

6.2.2 Implementation

An MCC is a runtime configuration object assembled dynamically when an interaction begins. It specifies five dimensions of access:

MCC Specification

```
allowedNoteIds: [list of accessible document IDs]
calendarScope: "full" | "busy_only" | "none"
todoFilter: {projects: [...], priorities: [...]}
allowedTools: [list of invocable tool names]
memoryShards: [list of accessible memory shard IDs]
```

allowedNoteIds Specifies which documents the agent can read. Documents not in this list are invisible.

calendarScope Controls calendar visibility: “full” (event details), “busy_only” (time blocks only), or “none”.

todoFilter Controls which tasks are visible, filtering by project and priority.

allowedTools Specifies which tools exist in the agent’s function array. Unlisted tools are absent entirely.

memoryShards Specifies accessible relationship memory shards, preventing cross-relationship contamination.

Per-requester assembly. The key property of MCCs is that they are assembled per-requester based on the relationship context. When Engineer A contacts Alex’s agent, the system assembles `MCC(Alex, Engineer_A)`. When Investor B contacts Alex’s agent, the system assembles `MCC(Alex, Investor_B)`. These two MCCs differ in every dimension:

Table 6.2: Per-requester MCC examples for the same target agent (Alex).

MCC dimension	Engineer A	Investor B
allowedNoteIds	Sprint notes, tech docs	Financial summaries, board decks
calendarScope	full	busy_only
todoFilter	{project: “v2-api”}	{priority: “high”}
allowedTools	read_note, create_todo	read_note, check_calendar
memoryShards	[engineer_a_shard]	[investor_b_shard]

The MCC eliminates several failure modes simultaneously. Transitive leakage is addressed because co-located sensitive data is not mounted for requesters who lack access. Over-refusal is reduced because the agent can freely use everything in its context without governance tension. Cross-cluster conflation is addressed because the MCC enforces domain-specific access regardless of relational closeness.

In SharedOS terms, the MCC acts as a filter between the State Layer and the agent’s execution context, selecting which state is visible and which tools are available based on the Relationship Context Layer.

6.3 MCC Design and Implementation

6.3.1 Assembly Pipeline

When a request arrives from an external agent, the MCC assembly pipeline executes in four stages:

1. **Relationship identification.** The system identifies the requester and loads the relationship context: type (colleague, investor, friend), trust level, historical interaction patterns, and any explicit access grants.
2. **Policy resolution.** The relationship context is mapped to an MCC template. Each relationship type has a default MCC that defines baseline access. Explicit grants and restrictions override the defaults.
3. **MCC construction.** The system assembles the concrete MCC by resolving the template against the current state. Document IDs are enumerated from the permitted folders. Tool lists are constructed from the permitted capability set. Memory shards are loaded from the relationship store.
4. **Context injection.** The assembled MCC is injected into the agent’s execution environment. The system prompt is constructed to reference only the accessible data. The tool array contains only the permitted tools. The retrieval index is scoped to the permitted documents and memory shards.

6.3.2 Namespace-Based Permissions

The MCC implementation uses a namespace-based permission model rather than per-object access control lists. Documents are organised in a hierarchical namespace (folders). Access is granted at the folder level. When a requester has access to folder `/work/projects/`, they can read all documents within that folder and its subfolders. Documents in `/personal/health/` are invisible.

Namespace-based permissions have three advantages over per-object ACLs. First, they scale linearly with the number of folders rather than quadratically with the number of documents. Second, they align with how humans think about information boundaries. People organise sensitive data into distinct areas. Third, new documents added to a folder automatically inherit the folder’s access profile. No per-document configuration is needed.

6.3.3 Relationship-to-MCC Mapping

The mapping from relationship types to MCCs is defined by explicit rules:

Table 6.3: Relationship-to-MCC mapping rules (simplified).

Relationship type	Folders accessible	Calendar	Tools
Co-founder	All work, strategy	full	All work tools
Engineer	Projects, sprint docs	full	read_note, create_todo
Investor	Financials, board decks	busy_only	read_note
Close friend	Personal (non-health)	full	read_note, calendar
Acquaintance	Public only	none	read_note (public)

These rules are configurable by the owner. The defaults reflect common social conventions. The owner can override them by granting specific access to individual requesters beyond their type-level defaults, or by restricting access below the defaults for specific individuals.

6.4 MCC Experiments & Results

6.4.1 Experimental Design

To validate the MCC architecture and decompose the contributions of structural versus behavioural governance, we compare three conditions using the PACT-PAIR task pool. The three conditions form a factorial design isolating the effect of prompt policy and structural access control independently:

- **Condition D3 (prompt policy only):** Full data loaded. Per-requester relationship policy in system prompt. The agent can read all folders and must rely on the policy to decide what to withhold.
- **Condition MCC_H (structural only):** Folder-scoped access control with *no* policy prompt. The agent receives no instructions about what to withhold; it simply cannot see unmounted folders. This isolates the pure contribution of data absence.
- **Condition MCC_H_D3 (structural + prompt):** Same folder-scoped access as MCC_H, plus the per-requester relationship policy from D3. This tests whether combining both mechanisms yields complementary benefits.

All three conditions use GPT-5.5 (Azure) as the target agent and GPT-5-mini (Azure) as the LLM judge.

Five requester personas probe 400 QA tasks each (200 Notes, 200 Todos):

Table 6.4: Requester personas and MCC folder grants.

ID	Relationship	D3 Access	MCC Folders
R0	Stranger	Full (policy gates)	Shared only
R1	Colleague	Full (policy gates)	Projects, Meetings, Shared
R2	Delegate (EA)	Full (policy gates)	Projects, Meetings, HR, Shared
R3	Close Friend	Full (policy gates)	Family, Shared
R4	Investor	Full (policy gates)	Projects, Shared

Evaluation uses an LLM judge (GPT-5-mini) with structured output that classifies each response as correct/incorrect/refused (utility questions) or leaked/safe/refused (security questions). A response counts as “leaked” only if it reveals specific sensitive data, not merely mentioning that a topic exists. Note that for the Todos QA track (Q201–400), MCC operates as coarse todo tool gating (on/off) rather than folder-scoped isolation, because the todo tools do not currently consume folder IDs. The Notes QA track (Q1–200) provides the clean folder-level MCC evidence.

6.4.2 Results

The three-condition decomposition reveals four findings.

Structure alone reduces leaks for misaligned requesters but *increases* them for aligned ones. Pure MCC_H (no prompt) cuts the aggregate leak rate from 15.5% to 12.4%, but this improvement is concentrated in R3 (−16.7 pp) and R4 (−15.3 pp) where folder restrictions remove the sensitive data entirely. For aligned

Table 6.5: Three-condition comparison (Q1–400, Notes + Todos combined). 6,000 judged responses total (2,000 per condition). Leak rate = fraction of P/B questions where specific sensitive data was disclosed.

Requester	Role	D3 (prompt only)		MCC_H (structure only)		MCC_H + D3 (both)	
		Util%	Leak%	Util%	Leak%	Util%	Leak%
R0	Stranger	21.0	1.2	23.8	1.6	23.9	0.6
R1	Colleague	87.4	8.0	82.9	20.3	83.3	6.3
R2	Delegate	87.1	12.6	86.6	17.9	85.9	12.6
R3	Close Friend	63.3	38.7	26.2	22.0	27.1	18.1
R4	Investor	87.4	16.9	59.7	1.6	60.6	2.7
Aggregate		70.9	15.5	57.6	12.4	58.5	8.0

Table 6.6: Decomposition of defence contributions by alignment class. Δ relative to D3 baseline.

Class	D3 $\rightarrow$ MCC_H		D3 $\rightarrow$ MCC_H+D3		MCC_H $\rightarrow$ MCC_H+D3	
	Δ Leak	Δ Util	Δ Leak	Δ Util	Δ Leak	Δ Util
Aligned (R0–R2)	+5.7	−0.7	−0.7	−0.9	−6.4	−0.2
Misaligned (R3–R4)	−15.5	−32.3	−17.5	−31.5	−2.0	+0.8
All	−3.1	−13.3	−7.5	−12.4	−4.4	+0.9

requesters (R1, R2), MCC_H actually *increases* leaks (+12.3 pp for R1, +5.3 pp for R2) because the agent freely discloses everything within mounted folders—including items the prompt policy would have suppressed. This demonstrates that structural access control and prompt policy are complementary, not substitutes.

The combined approach (MCC_H+D3) dominates D3 for misaligned requesters. Adding folder restrictions on top of prompt policy reduces R3 leaks from 38.7% to 18.1% (−53%) and R4 leaks from 16.9% to 2.7% (−84%). R4’s near-elimination of leaks occurs because MCC removes the Meetings folder entirely—containing 1:1 notes, performance reviews, and hiring decisions that the prompt policy only partially suppresses.

Prompt policy provides within-scope discrimination that structure cannot. The transition from MCC_H to MCC_H+D3 (Δ Leak = −6.4 pp for aligned requesters) shows that prompt policy contributes precisely where structural controls cannot help: within folders the requester legitimately accesses. R1’s leak rate drops from 20.3% (MCC_H) to 6.3% (MCC_H+D3) because the prompt tells the agent which items within shared folders remain private.

Residual leaks reveal the combined ceiling. MCC_H+D3’s residual 18.1% leak rate for R3 comes from the Family folder, which R3 legitimately accesses. These represent cases where both structural absence (folder is

mounted) and prompt policy (agent ignores the instruction) fail simultaneously. This residual motivates the Escalation Protocol: a runtime audit layer that catches disclosures missed by both structural and behavioural governance.

6.4.3 Network-Scale MCC Validation

The same complementarity appears in PACT-NET. We run two additional PACT-NET V2 MCC conditions: **MCC_H**, which applies folder-scoped structure without prompt policy, and **MCC_H+D1**, which combines folder scoping with the base network policy. D0 and D1 are averaged over two namespace-isolated replications; MCC_H and MCC_H+D1 are single-replication validation runs. The table reports Phase 1/core task metrics only, avoiding the evaluator’s separate Phase 2 dig-further bookkeeping.

Table 6.7: PACT-NET MCC validation. Safety and Utility are composite benchmark scores. Transitive leak and Deputy success are lower-is-better network risks; Cross-surface defence is higher-is-better action safety.

Condition	Policy	MCC	Safety	Utility	Trans. Leak	Deputy	Plant Def.
D0	No	No	26.6	88.7	96.3	47.0	0.0
D1	Yes	No	71.5	78.8	77.7	2.0	94.0
MCC_H	No	Yes	64.4	23.1	77.7	6.0	2.0
MCC_H+D1	Yes	Yes	77.8	20.6	67.0	0.0	92.0

Network safety improves monotonically. The safety ladder is D0 (26.6%) → MCC_H (64.4%) → D1 (71.5%) → MCC_H+D1 (77.8%). Structure alone recovers most of the safety benefit of prompt policy, and the combined condition is strongest. This supports the central MCC claim at network scale: removing out-of-scope data changes the reachable state space, not merely the model’s stated behaviour.

Structure and policy protect different surfaces. MCC_H matches D1 on transitive leakage (both 77.7%), and MCC_H+D1 reduces transitive leakage further to 67.0%. Confused-deputy success falls from 47.0% under D0 to 6.0% under MCC_H and 0.0% under MCC_H+D1. But cross-surface planting is not solved by structure alone: MCC_H defends only 2.0% of plant attempts, while D1 and MCC_H+D1 defend 94.0% and 92.0%. This is because MCC controls what the agent can read, while cross-surface planting depends on what the agent is willing to write.

Write utility is the main MCC gap. The low PACT-NET utility under MCC_H and MCC_H+D1 is not primarily a language-model failure. The evaluated MCC configuration grants read-only folder scope, which blocks almost all authorised write actions: authorised note creation falls to 0.0%, and authorised completion remains at 2.6%. A production MCC therefore needs a separate write-permission layer, not only read-side folder visibility.

6.4.4 Failure Modes MCC Cannot Address

MCCs are not a complete solution. Four failure modes remain.

Within-scope leakage. When a requester has legitimate access to a folder, MCC alone provides no protection for sensitive items co-located in that folder. The MCC_H condition demonstrates this clearly: R1’s leak rate is 20.3% under pure structure, versus 6.3% when prompt policy is added. Similarly, R3’s 22.0% MCC_H leak rate (versus 18.1% with MCC_H+D3) shows that within the Family folder, items the friend should see (wedding logistics) coexist with items the friend should not (family financial disagreements). MCC filters at the folder level, not at the item level.

Parametric knowledge. The model’s pre-training data may contain information about the owner that the MCC cannot remove. If the owner is a public figure, the model may “know” facts that the MCC excludes from the context. This is a limitation of any context-level governance mechanism.

Metadata leakage. The act of searching for information can reveal that the information exists, even if the search returns no results. If the agent says “I cannot find any health records to share with you,” it has confirmed the existence of a health-related data category. MCCs prevent data disclosure but not metadata disclosure.

Configuration scalability. The preceding failure modes concern correctness. A separate challenge is scalability of specification. In SharedOS, access control has two independent dimensions: *scope* (which data partitions are visible) and *tools* (which operations are permitted). For files, scopes are folders and tools are read, write, search—so a colleague might have read access to `project/reports` but write access only to `project/drafts`. For structured state such as todos, the same multiplication applies: which projects are visible

and which operations (read, create, complete) are allowed. The per-requester configuration surface is therefore $N_{\text{scopes}} \times M_{\text{tools}}$.

Write-path governance. Read-side data absence does not by itself decide what an agent may write. The PACT-NET MCC run exposes this sharply: a read-only MCC blocks unauthorised mutations, but it also blocks legitimate creates and completions. Conversely, cross-surface planting requires policy judgement because the risk comes from writing sensitive information into a shared surface. MCCs therefore need independent read and write capabilities rather than a single folder grant.

Some data types resist clean scoping entirely. Calendar events are not organised into folders; splitting them by sensitivity (e.g. separating a medical appointment from a team standup) requires manual categorisation or an agent-driven classification step. Email faces a similar challenge: folder-based separation is theoretically possible but rarely maintained in practice.

An agent could automate the categorisation, assigning items to scopes. But this addresses *assignment*, not *specification*. The policy matrix—which scopes grant which tool permissions to which requesters—must still be defined. With O namespaces in a multi-agent network, N scopes per namespace, M tools, and R relationship types, the total configuration surface is $O \times N \times M \times R$. Even modest growth in any dimension makes exhaustive pre-specification impractical. This scalability constraint, together with the correctness gaps above, motivates the Intelligent Escalation Protocol: rather than requiring the owner to pre-specify every cell in a combinatorial permission matrix, the system defers ambiguous boundary cases to the owner at runtime and learns from each decision to reduce future escalations.

6.5 Intelligent Escalation Protocol

6.5.1 Motivation

MCCs address the dominant failure mode (incidental disclosure of data present in context) but leave a second problem: not every boundary can be pre-specified as a folder permission. Even when a folder is legitimately mounted, a requester may ask for a fact whose appropriateness depends on relationship, intent, prior owner decisions, or conversational context. A close friend may legitimately access family logistics but not health details; an investor may access strategy summaries but not private meeting notes. Folder-level absence cannot

decide these within-scope boundary cases.

The evaluated version of the Intelligent Escalation Protocol therefore operates as a *pre-tool gate*. Before any read or write tool executes, the gate classifies the proposed operation as Continue or Stop. Continue lets the agent execute the tool. Stop prevents execution and creates an escalation record for the owner. This moves the decision point from free-form response generation to the tool boundary, before protected data enters the model's context.

6.5.2 Pre-Tool Escalation Gate

Before the agent executes a tool call, the escalation gate evaluates the proposed operation against the requester's relationship context and the owner's precedent database.

1. The agent proposes a tool call, such as `search_notes`, `search_todos`, or `search_agent_memory`.
2. The gate receives the requester identity, relationship context, policy memory, precedent matches, and proposed tool arguments.
3. The sanitiser returns an internal verdict: Allow, Escalate, or Deny.
4. The runtime maps this to a binary control decision: Allow → Continue; Escalate or Deny → Stop.
5. If Continue, the tool executes normally. If Stop, the tool is not executed and the requester receives a holding response such as "Let me check with Alex first."

This evaluated gate is intentionally narrower than a complete response-auditing system. It tests whether the system can decide, before tool execution, whether a proposed access should proceed.

6.5.3 Boundary Layers

The benchmark contains two different kinds of STOP-required cases, which correspond to different system layers.

- **Content-boundary cases (P labels).** The requester can contact the target agent, but the requested information is not appropriate for that requester. These are the core EscalationGate cases.

- **Contact-boundary cases (BLOCKED labels).** The requester is not in the target agent’s contact graph.

In the production OS, these requests should be rejected by routing before the target agent receives them.

The results therefore report content-boundary security separately from contact-boundary enforcement. This avoids attributing routing-layer successes to the query sanitiser.

6.5.4 Precedent-Based Tool Gating

The Escalation Protocol reduces manual review by turning owner decisions into reusable precedents. Each precedent records a relationship-conditioned judgement about a proposed tool call: who owns the target agent, who made the request, what resource or tool pattern was involved, whether the request was allowed or stopped, and why.

The final experiment deliberately removes automatic lookup and consensus shortcuts. Instead of asking whether a cache can remember previous decisions, it asks whether an LLM gate can use a small number of prior owner decisions as context for a new tool-call decision. This makes the evaluation a stricter test of precedent-conditioned reasoning.

In production, the natural precedent object is tool-based:

```
owner, requester, relationship summary, proposed tool,
resource pattern, sensitivity, prior decision, rationale
```

The research harness approximates this by deriving a synthetic tool call from each PACT-NET query and asking the gate whether that call should Continue or Stop. This is not a final-answer leakage test. It isolates the pre-tool boundary: if the gate stops the call, the agent never receives the protected tool result.

6.6 Escalation Protocol Validation

6.6.1 Experimental Design

The final escalation evaluation is a controlled PACT-NET gate ablation. It isolates the pre-tool sanitizer from prompt policy, MCC scoping, automatic decision caches, and final-answer generation. The sanitizer receives the target owner, requester, proposed synthetic tool call, relationship context, and precedent cards, then decides whether the call should Continue or Stop.

Labels. The experiment uses only PACT-NET L and P cells. L means the request is legitimate and should continue; P means the request is private and should stop. Borderline cells are excluded. BLOCKED cells are also excluded because they are contact-graph failures: in the OS, a non-contact request should be rejected before it reaches the target agent’s escalation gate.

Question-group split. Train/test splitting is done at the question level. All agent-pair cells for the same question are assigned to either train or test, preventing direct leakage of the same query text across the split. The 10% split contains 171 train cells and 1,548 test cells; the 30% split contains 533 train cells and 1,186 test cells.

Precedent scopes. Four conditions are evaluated for each model:

- **individual-10:** the gate sees only 10% same-pair precedents for the same target owner and requester.
- **cluster-2NN-10:** the gate sees the same individual precedents plus precedents from two similar requesters under the same target owner.
- **rich-cluster-2NN-10:** the same cluster retrieval, but shared precedents include richer relationship-aware cards with requester identity, relationship note, sensitivity, and similarity rationale.
- **individual-30:** the gate sees 30% same-pair precedents and serves as the direct-label utility ceiling.

The corrected transfer axis is same-owner, similar-requester transfer. For example, Tina’s agent may use Tina’s prior decisions about Sarah to help decide a similar request from Alex. The experiment does not test cross-owner transfer.

Metrics. PStop is recall on P items: the fraction of private requests correctly stopped. Utility is recall on L items: the fraction of legitimate requests correctly continued. False Continue is $1 - \text{PStop}$; False Stop is $1 - \text{Utility}$. The reported results are not blended accuracy, because security and utility failures have different meanings.

6.6.2 Results

The completed Phase 2 ablation produced eight conditions across GPT-5-mini and GPT-5.5. Table 6.8 reports all final conditions.

Table 6.8: Escalation gate Phase 2 ablation on PACT-NET. PStop is recall on private requests; Utility is recall on legitimate requests. All conditions use question-group splits and no auto-decision.

Model	Scope	Frac.	PStop	Utility	Errors
GPT-5-mini	individual	10%	90.8	69.8	0
GPT-5-mini	cluster-2NN	10%	88.5	76.7	0
GPT-5-mini	rich cluster-2NN	10%	91.3	78.0	0
GPT-5-mini	individual	30%	87.7	91.0	0
GPT-5.5	individual	10%	94.2	67.1	0
GPT-5.5	cluster-2NN	10%	92.7	71.3	1
GPT-5.5	rich cluster-2NN	10%	94.4	68.1	0
GPT-5.5	individual	30%	92.1	87.4	0

Sparse same-pair precedents are secure but conservative. With only 10% same-pair precedents, both models stop most private requests: 90.8% for GPT-5-mini and 94.2% for GPT-5.5. Utility is much lower: 69.8% and 67.1%. This is the expected failure mode of a privacy gate under sparse evidence: uncertainty is often converted into Stop.

Same-owner transfer improves utility but does not replace direct labels. Adding two similar requester relationships under the same owner improves utility by 6.9 percentage points for GPT-5-mini and 4.1 points for GPT-5.5, while slightly reducing PStop. However, cluster-2NN at 10% does not match the 30% same-pair baseline. The utility gap remains 14.3 points for GPT-5-mini and 16.2 points for GPT-5.5. Precedent count alone is therefore insufficient; the semantic quality of the precedent matters.

Rich cards expose a representation bottleneck. Rich relationship-aware cards improve GPT-5-mini on both axes relative to anonymous cluster cards: PStop rises from 88.5% to 91.3%, and Utility rises from 76.7% to 78.0%. For GPT-5.5, rich cards improve PStop from 92.7% to 94.4%, but Utility falls from 71.3% to 68.1%. The stronger model appears to use the additional relationship and sensitivity context more conservatively. This is useful evidence: escalation quality depends not only on retrieving similar relationships, but on representing precedents with the right amount of boundary-relevant context.

Conclusion. The strongest supported claim is not that cluster transfer solves sparse labelling. It is narrower and more defensible: precedent-conditioned tool gating can learn high-security relationship boundaries from sparse examples, same-owner similar-requester transfer can reduce false stops, and richer precedent cards are

necessary but not sufficient to match direct same-pair labels. A full production evaluation would require running the gate inside the agent loop and judging final answers end-to-end.

6.7 Production Deployment

The architectural solutions described in this chapter are not purely theoretical. They are implemented and deployed in Pulse (the production system underlying SharedOS). This section describes how the research concepts map to production features.

6.7.1 Folder-Scoped Share Links as Production MCC

In Pulse, the MCC concept is realised through folder-scoped share links. When a user shares access to their agent, they specify which folders are visible. The share link encodes the MCC specification: the set of accessible folders, the tool permissions, and the calendar scope.

Each share link is a distinct MCC. The user can create multiple share links with different scopes:

- A “work collaborator” link that exposes project folders and sprint tasks
- An “investor” link that exposes financial summaries and board presentations
- A “friend” link that exposes personal preferences and social calendar

When someone accesses the agent through a share link, the system assembles the execution environment from the link’s MCC specification. The agent only sees what the link permits. This is not a permission check applied after retrieval. It is a scoping constraint applied before retrieval. The retrieval index is rebuilt per-link to contain only the permitted documents.

6.7.2 Relationship Shards as Per-Requester Context

Pulse implements relationship memory shards through tagged vector storage. Each conversation with an external party is tagged with the relationship identifier. Retrieval queries include a hard filter on the relationship tag. Cross-relationship contamination is prevented at the database layer.

In production, this means that when Investor A asks a question, the agent retrieves memory only from the Investor A shard and the owner’s self-state. It does not retrieve fragments from conversations with Investor B,

even if those fragments are semantically relevant to the query. The isolation is enforced by the database query, not by the model’s judgement.

6.7.3 Escalation in Production

The Escalation Protocol is partially deployed. The production implementation wraps cross-boundary tools with a pre-execution gate: if the sanitizer returns Continue, the tool executes; if it returns Stop, the tool is not executed and an escalation record is created for owner review. The precedent system accumulates owner decisions as structured context for future gate decisions; richer relationship-aware precedent cards remain an active design requirement.

Contact-graph enforcement remains a separate upstream layer. Requests from non-contacts should be rejected by routing before the target agent’s tool gate is invoked. A post-generation auditor is a natural extension, but it is not the mechanism evaluated in this chapter.

6.7.4 What Works vs What Remains Theoretical

Table 6.9: Deployment status of architectural solutions.

Component	Status	Limitation
Folder-scoped MCC	Deployed	Document-level, not field-level
Per-requester tool scoping	Deployed	Manual configuration required
Relationship memory shards	Deployed	L2 only, L3 pending
Pre-tool escalation gate	Deployed	Gate-only, not final-answer judging
Contact graph enforcement	Deployed	Routing layer, separate from sanitizer
Precedent learning	Prototype	Rich cards and retrieval calibration pending
Post-generation audit	Future work	Not evaluated here
Cross-category decomposition	Not deployed	Research prototype only

The gap between the research prototype and production deployment is primarily one of granularity and representation. Production MCCs operate at the folder level. The research proposes field-level granularity (e.g., showing a calendar event’s time but not its title). Production escalation currently operates at the tool boundary, but the Phase 2 ablation shows that precedent-card design materially affects utility. Future extensions should add post-generation auditing, query/resource-similar precedent retrieval, and multi-turn cumulative leakage tracking. These extensions are architecturally compatible with the deployed system but require additional engineering effort.

6.8 Summary

This chapter presented two complementary architectural solutions for cross-boundary agent privacy. Mountable Context Cells enforce data absence: protected information never enters the agent's context window, making disclosure structurally impossible regardless of model behaviour. The Intelligent Escalation Protocol enforces runtime monitoring at the tool boundary: proposed tool calls are stopped before protected data enters the agent's context, while owner precedents reduce repeated manual review.

Together, these solutions transform privacy governance from a behavioural property (dependent on the model following instructions) into a structural property (enforced by the execution environment). The combined architecture addresses all failure modes identified in Chapters 4 and 5, with three residual categories (parametric knowledge, coordinated multi-requester attacks, and temporal consistency) representing the long tail of risk.

The solutions are not hypothetical. They are deployed in production at varying levels of granularity. The gap between research and deployment is one of precision (field-level vs folder-level) and scope (all categories vs high-sensitivity only). The architectural principles are validated. The engineering work to extend them to full coverage is incremental rather than fundamental.

Chapter 7

Conclusion & Discussion

The preceding chapters presented the problem (cross-boundary agent privacy lacks structural enforcement), the diagnostic benchmarks (PACT-PAIR and PACT-NET), and the architectural solutions (Mountable Context Cells and the Intelligent Escalation Protocol). Notably, SharedOS is not merely a research prototype: the deployed platform (Pulse/Aicoo) hosts thousands of active agents, with thousands of real cross-agent communications recorded daily, grounding our findings in production-scale evidence. This chapter steps back from the technical details to discuss what these findings mean for the broader field of agent system design. We address the limits of prompt engineering, the implications for practitioners building multi-agent systems, the role of SharedOS as research infrastructure, and the limitations and ethical considerations of this work.

At the highest level, the thesis has three linked claims. First, cross-boundary agent delegation is not a prompt-design problem but an operating-system problem: the platform decides which state, tools, policies, and relationship memories enter the agent’s execution environment. Second, the right empirical object is the security–utility frontier, because privacy systems that maximise refusal are as unusable as systems that maximise disclosure. Third, the frontier only expands when enforcement moves earlier in the stack—from final response judgement, to pre-tool gating, to context construction itself.

7.1 The Frontier Cannot Be Flattened By Prompt Engineering

The central empirical finding of this thesis is that prompt-level governance has a ceiling. Under the best prompt configuration (D2, category-specific deny list), 27–38% of protected information still leaks in dyadic evaluation depending on the model and evaluation mode. At network scale (PACT-NET), even with base policy, transitive leaks persist at 77.7% and cross-cluster leaks at 69.6%. This residual is not a matter of poor prompt writing. It reflects a structural limitation of instructional governance.

Three forces prevent prompt engineering from eliminating the residual.

Implicit social norms. Language models are trained on human conversation, where disclosure decisions are governed by implicit social norms rather than explicit rules. A human asked “How is the fundraiser going?” by a trusted colleague would likely share some details, even if a policy says otherwise. The model inherits this tendency. It resolves ambiguous cases by defaulting to social convention rather than strict policy compliance. No amount of prompt specificity can override a training signal that pervades the entire pretraining corpus.

Semantic properties of natural language. Natural language is compositional. A response that individually reveals no protected fact may collectively reveal protected information through implication, elimination, or contextual inference. The policy author cannot enumerate every possible compositional leak path. Prompt-level governance operates on explicit rules. Semantic leakage operates on implicit entailment.

Conversational dynamics. Multi-turn interaction creates pressure toward disclosure. The requester builds rapport, establishes context, and gradually narrows the information boundary. Each individual response may comply with the policy, but the cumulative effect is disclosure. PACT-PAIR’s $3\times$ gap between single-turn and multi-turn leakage quantifies this effect. Conversational dynamics are not addressable by prompt engineering because they exploit the temporal dimension, which static instructions cannot control.

These three forces define the ceiling. Prompt engineering can approach the ceiling (D2 does so). It cannot break through it. Architectural solutions (MCCs and the Escalation Protocol) break through by operating at a different layer. They do not ask the model to resist social norms, enumerate semantic entailments, or track conversational dynamics. They remove the data from the context or gate tool execution before protected data is retrieved.

The role of model scale. Larger models perform better under D0 (no policy). They have stronger default safety behaviour from RLHF training. But under D2, model scale provides diminishing returns. The specificity threshold (D1 to D2) equalises models: once category-level instructions are provided, smaller models match or approach larger models. This suggests that the ceiling is not a function of model capability but a function of the governance mechanism. A more capable model does not overcome the structural limitation of instructional governance. It merely approaches the ceiling faster.

7.2 Implications for Agent System Design

The findings of this thesis have practical implications for practitioners building multi-agent systems with cross-boundary communication.

7.2.1 Safety Training Is Not Privacy Governance

Model providers invest heavily in safety training: RLHF, constitutional AI, red-teaming. These techniques teach models to refuse harmful requests, avoid generating toxic content, and decline to assist with dangerous activities. Safety training and privacy governance are different capabilities.

Safety training addresses universal harms. Toxic content is harmful regardless of context. Bomb-making instructions are dangerous regardless of who asks. Privacy governance addresses contextual harms. Alex's salary is appropriate to share with the CEO but harmful to share with a junior engineer. The harm depends on the requester, not the content.

This distinction means that a model with excellent safety training may still fail at privacy governance. It knows not to generate harmful content. It does not know that the same content is harmful for one requester and appropriate for another. Privacy governance requires relationship context, which safety training does not provide.

7.2.2 Relationship Context Is a Feature, Not a Bug

PACT-PAIR and PACT-NET demonstrate that the same question has different correct answers depending on who asks. This is not a flaw in the evaluation. It is a fundamental property of privacy in social systems. Information sensitivity is relational, not absolute.

This has a design implication. Agent systems that treat privacy as a property of data objects ("this file is sensitive") will systematically fail. Some files are sensitive for some requesters and appropriate for others. The governance mechanism must be parameterised by the requester's identity, not just by the data's classification. MCCs embody this principle: the same data store produces different context windows for different requesters.

7.2.3 Over-Refusal Is the Dominant Real-World Failure

Academic security research typically focuses on successful attacks: how to break the system. In production deployment, over-refusal is the more common and more damaging failure mode. An agent that refuses legitimate requests frustrates users, reduces adoption, and undermines the utility proposition of the system.

PACT-PAIR's results show that over-refusal accounts for 18–25% of utility failures under D2 governance. PACT-NET confirms this at network scale: base policy reduces should_answer accuracy by 11.9 pp (75.9% → 64.0%). Users abandon agents that refuse too often. The practical consequence is that security and utility must be optimised jointly. A system that achieves perfect security by refusing everything is useless. The three-condition MCC experiment demonstrates this precisely: for aligned-access requesters (Stranger, Colleague, Delegate), the combined MCC+D3 stack costs less than 1 pp utility versus D3 alone, while structural-only MCC (no prompt) achieves 86–95% utility for R1/R2 because the agent freely uses everything in its context without governance tension.

7.2.4 Read Trust and Write Trust Are Independent

PACT-NET's cross-surface plant tasks (D0: 0.0% defence, D1: 94.0% defence — the single largest policy effect at +94.0 pp) reveal that read access and write access create different risk profiles. A requester may be trusted to read project documents but not trusted to trigger writes that expose sensitive data through a side channel. Conversely, a requester may be trusted to create tasks (a write operation) but not to read the full task database.

Agent system designers must treat read and write permissions as independent dimensions. A single “trust level” that grants both read and write access at the same granularity will either over-expose data (granting reads that should be restricted) or under-utilise the agent (restricting writes that should be allowed). MCCs support this through separate “allowedTools” specifications that distinguish between read tools and write tools.

7.3 SharedOS as Research Infrastructure

SharedOS is both the platform on which our experiments run and a contribution in its own right. Its value lies not in the specific findings it has produced (which are tied to the current generation of models) but in its extensibility as a research instrument.

New state types. The current deployment uses documents, structured records, and memory shards. The architecture supports arbitrary state types. Researchers studying agent privacy in domains like healthcare (FHIR records), legal (case files), or education (student records) can instantiate SharedOS with domain-specific state while retaining the evaluation infrastructure.

New governance policies. The D0/D1/D2 specificity gradient is one point in the policy design space. Researchers can define new governance mechanisms (role-based, attribute-based, context-aware) and evaluate them against the same benchmark tasks without modifying the platform.

New relationship configurations. PACT-NET’s 25-agent ecosystem is one social graph. The platform supports arbitrary graph structures. Researchers studying privacy in hierarchical organisations, peer networks, or adversarial environments can define custom graphs and relationship types.

New models. The evaluation is model-agnostic. Any model that accepts the OpenAI function-calling API can be evaluated without platform modifications. This enables longitudinal studies of how privacy behaviour evolves across model generations.

We do not claim that SharedOS is a universal benchmark for agent privacy. We claim that it is a platform that revealed specific phenomena (the specificity threshold, the $3\times$ multi-turn gap, transitive leakage, over-refusal dominance) and that its extensible design enables future researchers to reveal phenomena we have not anticipated.

7.4 Limitations

This work has several limitations that constrain the generalisability of its findings.

Synthetic data. All evaluation data is synthetic. Documents, structured records, and relationship configurations are authored rather than collected from real users. While the data is structured to reflect realistic sensitivity boundaries (informed by the authors’ experience building production systems), it cannot capture the full complexity and idiosyncrasy of real personal data. Real users have privacy preferences that are inconsistent, context-dependent, and culturally variable. Our benchmark evaluates against well-defined labels. Real privacy is messier.

Same-model attacker and defender. In PACT-PAIR and PACT-NET, the requesting agent (attacker) and the target agent (defender) use the same model family. In production, attackers may use specialised models, fine-tuned for extraction, while defenders use general-purpose models. Our evaluation may underestimate adversarial capability because the attacker is not optimised for attack.

Organic multi-turn only. The multi-turn evaluation uses adaptive questioning strategies generated by the requesting agent. It does not employ formal adversarial techniques such as PAIR [21], Crescendo [23], or tree-of-attacks [43]. These techniques produce more sophisticated attack trajectories that may defeat governance mechanisms that resist organic probing. Our evaluation represents a lower bound on adversarial success.

MCC validation is partial. The deployed MCC operates at the folder level for notes, not at the field level. A folder may contain documents at different sensitivity levels. The MCC includes all documents in a permitted folder, which may include some that should be restricted — this is precisely what produces R3 Friend’s 18.1% residual leak rate from the Family folder. True field-level MCC (where individual facts within a document are independently governed) is architecturally possible but not yet implemented or evaluated. Additionally, MCC does not yet provide folder-level isolation for todos (only coarse on/off gating), which limits the strength of the Todos QA results.

Single ecosystem. Both PACT-PAIR and PACT-NET operate within a single fictional ecosystem (a technology startup). The relationship types (co-founder, engineer, investor, friend) and information categories (work, finance, health, relationships) reflect this context. Other domains may produce different sensitivity distributions, different relationship hierarchies, and different failure mode distributions. We cannot claim that our findings generalise to healthcare, government, or education contexts without additional evaluation.

7.5 Ethical Considerations

Research on agent privacy has dual-use potential. The same findings that help system designers build stronger privacy protections could help attackers identify weaknesses in existing systems.

Attack taxonomy as a roadmap. The failure mode taxonomy in Chapter 4 (transitive leakage, confused deputies, cross-surface plants) identifies specific attack strategies. Publishing these strategies makes them

available to adversaries. We mitigate this concern through two observations. First, the attack strategies are not novel. They are adaptations of well-known attack patterns (social engineering, confused deputies, side channels) to the agent context. Second, the architectural defences we propose address each attack strategy. The net effect of publication is to provide both the attack and the defence simultaneously.

Responsible testing. All experiments were conducted on our own system (Pulse/SharedOS). No third-party systems were probed or attacked. The evaluation tasks target a fictional persona (Alex) in a fictional company (TechFlow AI). No real user data was used, accessed, or at risk during any experiment.

Fully synthetic data. The evaluation data is entirely synthetic. No real personal information, corporate data, health records, or financial details appear in the benchmark. The data is designed to be realistic in structure (to test the correct failure modes) but fictional in content (to avoid any privacy harm from the research itself).

Disclosure and reproducibility. We intend to release the evaluation framework (SharedOS), the benchmark tasks (PACT-PAIR and PACT-NET), and the architectural specifications (MCC and Escalation Protocol) to enable reproducibility. The release will include the evaluation methodology but not the specific attack prompts that achieved highest success rates. This balances reproducibility with responsible disclosure.

7.6 Summary

This thesis investigated privacy governance in multi-agent shared delegation systems. We formulated cross-boundary delegation as a setting where security enforcement is emergent rather than designed: agents must decide, at inference time, what to share and what to withhold, based on policies expressed as natural-language instructions rather than structural constraints. This formulation revealed that the security-utility frontier is the central object of study, not security or utility in isolation.

We made five contributions toward understanding and improving this frontier.

SharedOS: the first platform for systematic study. We designed and implemented SharedOS, a multi-agent coordination layer with persistent state, tool-mediated interactions, relationship-aware memory, and configurable defence mechanisms. SharedOS is not a benchmark in a vacuum. It is a deployed system that grounds evaluation in real architectural constraints. The platform enables controlled experiments on cross-boundary

delegation that no prior system could support, because no prior system simultaneously maintained persistent per-relationship state, enforced trust tiers through policy files, and provided tool-mediated communication channels between agents representing different principals.

PACT-PAIR: charting the dyadic security-utility frontier. We designed and executed the PACT-PAIR benchmark across six frontier LLMs, three relationship tiers, and a defence gradient from zero-shot to full category enumeration. Three research questions structured this investigation:

- **RQ1: What moves the frontier?** Only explicit category enumeration in the system prompt produces meaningful frontier movement. Across six models, category enumeration improved the security-utility balance by 69 to 91 percentage points compared to zero-shot baselines. Generic privacy instructions, role-based prompting, and few-shot examples all failed to move the frontier reliably. The implication is that LLMs require exhaustive specification of what constitutes private information. They cannot generalise privacy norms from abstract principles.
- **RQ2: How does multi-turn interaction affect the frontier?** Multi-turn erosion degrades security not through adversarial cracking but through incidental channels. Agents leak information through context accumulation (each turn adds context that makes the next disclosure more natural), through clarification responses that reveal metadata, and through over-helpful elaboration on topics adjacent to sensitive data. The adversarial escalation pattern (Crescendo-style gradual steering) was less effective than simple persistent conversation, because agents are trained to detect escalation but not trained to detect gradual context drift.
- **RQ3: How does relationship reshape the frontier?** Relationship closeness reshapes the frontier per requester. For close relationships (trusted collaborators), over-refusal dominates leakage as the primary failure mode: agents withhold information that the owner would want shared, causing utility loss without security benefit. For distant relationships (strangers), leakage dominates over-refusal. The frontier is therefore not a single curve but a family of curves parameterised by relationship. Any defence that ignores relationship will either over-refuse for friends or under-protect against strangers.

PACT-NET: validation at network scale. We extended evaluation from dyadic interactions to network-scale coordination involving multiple agents with overlapping social graphs. PACT-NET revealed two phe-

nomena invisible at the dyadic level. First, *transitive leakage*: information shared legitimately with Agent B can propagate to Agent C through B’s subsequent interactions, even when A has no direct relationship with C. Second, *amplification effects*: the same query posed to multiple agents in a network yields more information than the sum of individual responses, because each agent contributes complementary fragments that reconstruct the full picture.

Architectural solutions: MCC and Escalation Protocol. We proposed and evaluated two mechanisms that move enforcement from the reasoning layer to the structural layer. The Mountable Context Cell (MCC) physically constructs execution environments containing only authorised resources, making security an environmental constraint rather than a behavioural decision. A three-condition factorial experiment (prompt-only, structure-only, combined) revealed that structural and behavioural governance are complementary: structure alone reduces leaks by 20% (15.5% → 12.4%) but increases within-scope leaks for aligned requesters; the combination reduces leaks by 48% (15.5% → 8.0%) by deploying each mechanism where the other is weakest. The Escalation Protocol gates tool execution before retrieval and uses owner precedents as relationship-conditioned context. In the final PACT-NET ablation, sparse same-pair precedents achieve high PStop (90.8–94.2%) but low utility (67.1–69.8%); same-owner relationship transfer improves utility by 4.1–6.9pp, while rich cards reveal that precedent representation remains a bottleneck.

7.7 Future Work

Several directions extend naturally from this work. We organise them by proximity to the current contribution, from near-term extensions to longer-term research programmes.

7.7.1 Formal Attack Integration

Our current evaluation uses human-designed adversarial queries and natural-language social engineering. A natural extension is to integrate formal attack algorithms (PAIR, Crescendo, PAP) as automated red-teamers within the PACT framework. This would enable continuous adversarial evaluation at scale: as new attacks are published, they can be incorporated as new attack strategies in the benchmark. The key technical challenge is adapting single-turn attack algorithms to the multi-turn, relationship-aware setting of cross-boundary delegation.

7.7.2 Cross-Model Adversarial Evaluation

Our current experiments use the same model family for both the defending agent and the evaluation. A stronger evaluation would pair a more capable attacker model against a less capable defender model, measuring whether capability asymmetry shifts the security-utility frontier. Preliminary results suggest that frontier models attacking smaller models achieve substantially higher attack success, but a systematic study across model pairs and defence configurations remains future work.

7.7.3 Per-Field Mountable Context Cells

The current MCC design operates at the folder level: an entire context folder is either mounted or not mounted for a given interaction. A more granular design would support per-field mounting, where individual attributes within a document (e.g., a calendar event’s title but not its attendees) can be selectively exposed. This requires a richer metadata schema and a mounting decision that operates at sub-document granularity. The engineering challenge is maintaining acceptable latency when the mounting decision must evaluate hundreds of individual fields per interaction.

7.7.4 Formal Verification of Context Cell Isolation

The security guarantees of Mountable Context Cells are architectural but not formally verified. Applying model checking or information flow analysis [12] to prove that MCC satisfies non-interference properties (no execution trace within a cell can observe data outside the cell’s scope) would strengthen the security claims. The challenge is modelling the LLM’s behaviour within the formal framework, given that LLMs are stochastic and their internal computation is not directly observable.

7.7.5 User Study with Real Humans and Real Relationships

Our evaluation uses simulated relationships and synthetic personal data. A longitudinal user study with real participants, genuine social relationships, and authentic personal information would validate whether the observed phenomena (over-refusal for close relationships, incidental leakage through multi-turn drift) manifest in naturalistic settings. Such a study raises significant ethical considerations: participants would need to share

real personal information with AI agents that may leak it, requiring careful IRB oversight and informed consent protocols.

7.7.6 SharedOS Open-Source Release

The platform, benchmark, and evaluation tools developed in this thesis can serve the broader research community as infrastructure for studying cross-boundary agent coordination. An open-source release would enable other researchers to evaluate new defence mechanisms, test new attack strategies, and extend the benchmark to new domains beyond personal information management.

7.7.7 Network-Scale MCC: Transitive Propagation

The PACT-NET results on transitive leakage motivate a network-scale extension of the MCC mechanism. Currently, MCC controls what information enters a single agent’s context. At network scale, we need to control how information propagates through the social graph: if Alice shares information with Bob’s agent under a restrictive policy, that policy should propagate when Bob’s agent subsequently interacts with Carol’s agent. This requires a distributed policy propagation protocol that maintains provenance and enforces access constraints transitively. The design draws on taint tracking from information flow control but must operate across independently deployed agents with no shared runtime.

7.8 Closing Statement

The central thesis of this work can be stated simply: *do not ask the delegate to refrain from accessing sensitive data; ensure the sensitive data is not in the room.*

This principle, drawn from capability-based security and instantiated through Mountable Context Cells, transforms privacy from a behavioural property (dependent on LLM compliance with natural-language instructions) into a structural property (enforced by the composition of the execution environment). The distinction matters because behavioural properties fail under adaptive attack, while structural properties hold regardless of the sophistication of the adversary.

The coordination problem facing AI agents today will not be solved by more capable individual agents. A more capable agent is also a more capable adversary. The solution lies in trust infrastructure between agents:

mechanisms that enable safe coordination without requiring each agent to be perfectly aligned, perfectly robust, and perfectly obedient.

SharedOS is a step toward that infrastructure. The PACT benchmark provides a methodology for measuring progress. The architectural solutions (MCC, Escalation Protocol) demonstrate that meaningful security-utility improvements are achievable without sacrificing the benefits of agent delegation. Much work remains, but the direction is clear: from prompt-level restriction to structural enforcement, from single-agent safety to multi-agent trust, from behavioural compliance to architectural guarantee.

Bibliography

- [1] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, “React: Synergizing reasoning and acting in language models,” *arXiv preprint arXiv:2210.03629*, 2023.
- [2] T. Schick, J. Dwivedi-Yu, R. Dessì, R. Raileanu, M. Lomeli, E. Hambro, L. Zettlemoyer, N. Cancedda, and T. Scialom, “Toolformer: Language models can teach themselves to use tools,” *Advances in Neural Information Processing Systems*, vol. 36, 2023.
- [3] N. Shinn, F. Cassano, A. Gopinath, K. Narasimhan, and S. Yao, “Reflexion: Language agents with verbal reinforcement learning,” *Advances in Neural Information Processing Systems*, vol. 36, 2023.
- [4] C. Packer, S. Wooders, K. Lin, V. Fang, S. G. Patil, I. Stoica, and J. E. Gonzalez, “Memgpt: Towards llms as operating systems,” *arXiv preprint arXiv:2310.08560*, 2023.
- [5] Cognition Labs, “Devin: The first ai software engineer.” <https://www.cognition.ai/devin>, 2024.
- [6] Manus AI, “Manus: Autonomous desktop agent.” <https://www.manus.ai>, 2025.
- [7] Anthropic, “Model context protocol: Connecting llms to external systems.” <https://modelcontextprotocol.io>, 2024.
- [8] Google, “Agent-to-agent protocol: Enabling agent interoperability.” <https://github.com/google/a2a-protocol>, 2025.
- [9] Y. Talebirad and A. Nadiri, “Multi-agent collaboration: Harnessing the power of intelligent llm agents,” *arXiv preprint arXiv:2306.03314*, 2023.
- [10] T. Guo, X. Chen, Y. Wang, R. Chang, S. Pei, N. V. Chawla, O. Wiest, and X. Zhang, “Large language model based multi-agents: A survey of progress and challenges,” *arXiv preprint arXiv:2402.01680*, 2024.
- [11] D. E. Bell and L. J. LaPadula, “Secure computer systems: Mathematical foundations,” *MITRE Technical Report*, vol. 2547, 1973.
- [12] D. E. Denning, “A lattice model of secure information flow,” *Communications of the ACM*, vol. 19, no. 5, pp. 236–243, 1976.
- [13] Q. Wu, G. Bansal, J. Zhang, Y. Wu, B. Li, E. Zhu, L. Jiang, X. Zhang, S. Zhang, J. Liu, *et al.*, “Autogen: Enabling next-gen llm applications via multi-agent conversation,” *arXiv preprint arXiv:2308.08155*, 2023.
- [14] CrewAI, “Crewai: Framework for orchestrating role-playing ai agents.” <https://www.crewai.com>, 2024.
- [15] S. Hong, M. Zhuge, J. Chen, X. Zheng, Y. Cheng, C. Zhang, J. Wang, Z. Wang, S. K. S. Yau, Z. Lin, *et al.*, “Metagpt: Meta programming for a multi-agent collaborative framework,” *arXiv preprint arXiv:2308.00352*, 2023.
- [16] G. Li, H. A. A. K. Hammoud, H. Itani, D. Khizbullin, and B. Ghanem, “Camel: Communicative agents for “mind” exploration of large language model society,” *Advances in Neural Information Processing Systems*, vol. 36, 2023.
- [17] K. Mei, Z. Li, S. Xu, R. Ye, Y. Ge, and Y. Zhang, “Aios: Llm agent operating system,” *arXiv preprint arXiv:2403.16971*, 2025.

- [18] Z. Wu, C. Han, Z. Ding, Z. Weng, Z. Liu, S. Yao, T. Yu, and L. Kong, “Os-copilot: Towards generalist computer agents with self-improvement,” *arXiv preprint arXiv:2402.07456*, 2024.
- [19] A. Zou, Z. Wang, J. Z. Kolter, and M. Fredrikson, “Universal and transferable adversarial attacks on aligned language models,” *arXiv preprint arXiv:2307.15043*, 2023.
- [20] X. Liu, N. Xu, M. Chen, and C. Xiao, “Autodan: Generating stealthy jailbreak prompts on aligned large language models,” in *International Conference on Learning Representations*, 2024.
- [21] P. Chao, A. Robey, E. Dobriban, H. Hassani, G. J. Pappas, and E. Wong, “Jailbreaking black box large language models in twenty queries,” in *IEEE Conference on Safe and Trustworthy Machine Learning (SaTML)*, 2025.
- [22] Y. Zeng, H. Lin, J. Zhang, D. Yang, R. Jia, and W. Shi, “How johnny can persuade llms to jailbreak them: Rethinking persuasion to challenge ai safety by humanizing llms,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.
- [23] M. Russinovich, A. Salem, and R. Eldan, “Great, now write an article about that: The crescendo multi-turn llm jailbreak attack,” in *USENIX Security Symposium*, 2025.
- [24] M. Nasr *et al.*, “The attacker moves second: Evaluating prompt injection defenses under adaptive attack,” *arXiv preprint*, 2025.
- [25] E. Wallace, K. Xiao, J. Leike, L. Weng, J. Heidecke, and A. Beutel, “The instruction hierarchy: Training llms to prioritize privileged instructions,” *arXiv preprint arXiv:2404.13208*, 2024. OpenAI.
- [26] K. Hines, G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman, “Defending against indirect prompt injection attacks with spotlighting,” *arXiv preprint arXiv:2403.14720*, 2024. Microsoft.
- [27] H. Inan, K. Upasani, J. Chi, *et al.*, “Llama guard: Llm-based input-output safeguard for human-ai conversations,” *arXiv preprint arXiv:2312.06674*, 2023. Meta.
- [28] T. Rebedea, R. Dinu, M. N. Sreedhar, C. Parisien, and J. Cohen, “Nemo guardrails: A toolkit for controllable and safe llm applications with programmable rails,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, 2023.
- [29] C. Dwork, “Differential privacy,” *International Colloquium on Automata, Languages, and Programming*, pp. 1–12, 2006.
- [30] Y. Yang and Y. Zhu, “Differential privacy in generative ai agents: Analysis and optimal tradeoffs,” *arXiv preprint arXiv:2603.17902*, 2026.
- [31] N. Carlini, F. Tramer, E. Wallace, M. Jagielski, *et al.*, “Extracting training data from large language models,” in *USENIX Security Symposium*, 2021.
- [32] Q. Zhan, Z. Liang, Z. Ying, and D. Kang, “Injecagent: Benchmarking indirect prompt injections in tool-integrated llm agents,” in *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.
- [33] E. Debenedetti, J. Zhang, M. Balunovic, L. Beurer-Kellner, M. Fischer, and F. Tramer, “Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents,” in *Advances in Neural Information Processing Systems*, vol. 37, 2024.
- [34] S. Toyer, O. Watkins, E. Mendes, J. Sveigliato, L. Bailey, T. Wang, I. Ong, K. Elmaaroufi, P. Abbeel, T. Darrell, A. Ritter, and S. Russell, “Tensor trust: Interpretable prompt injection attacks from an online game,” in *Advances in Neural Information Processing Systems*, vol. 37, 2024.

- [35] S. Schulhoff, J. Pinto, A. Khan, L.-F. Bouchard, C. Si, S. Anati, N. Tagber, M. Karpinska, B. Dolan, and J. Boyd-Graber, “Hackaprompt: An adversarial prompt engineering competition,” in *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2024.
- [36] Y. Wang and N. Jiang, “Agentsocialbench: Evaluating privacy risks in human-centred agentic social networks,” *arXiv preprint arXiv:2604.01487*, 2026.
- [37] W. Gomaa, A. Salem, and S. Abdelnabi, “Converse: Benchmarking contextual safety in agent-to-agent conversations,” *arXiv preprint arXiv:2511.05359*, 2025. EACL 2026 Findings.
- [38] J. Park, S. Kim, H. Ju, J. Lim, S. Choi, O. Kwon, J. Kim, and S. Yeo, “Pac-bench: Evaluating multi-agent collaboration under privacy constraints,” *arXiv preprint arXiv:2604.11523*, 2026.
- [39] P. Juneja, L. B. Pasupulati, A. Albalak, W. Hua, and W. Y. Wang, “Magpie: A benchmark for multi-agent contextual privacy evaluation,” *arXiv preprint arXiv:2510.15186*, 2025.
- [40] O. Yagoubi, G. Badu-Marfo, and R. Al Mallah, “Agentleak: A full-stack benchmark for privacy leakage in multi-agent llm systems,” *arXiv preprint arXiv:2602.11510*, 2026.
- [41] N. Shapira, C. Wendler, H. Yen, *et al.*, “Agents of chaos,” *arXiv preprint arXiv:2602.20021*, 2026.
- [42] J. B. Dennis and E. C. Van Horn, “Programming semantics for multiprogrammed computations,” *Communications of the ACM*, vol. 9, no. 3, pp. 143–155, 1966.
- [43] A. Mehrotra, M. Zampetakis, P. Kassianik, B. Nelson, H. Anderson, Y. Singer, and A. Karbasi, “Tree of attacks: Jailbreaking black-box llms automatically,” in *Advances in Neural Information Processing Systems*, vol. 37, 2024.