

How Should Your Agent Talk to Mine?

The Security–Utility Frontier in Cross-Boundary Agent Delegation



Xisen Wang

Keble College

4th-Year Project Report

Supervised by

Professor Philip Torr

Professor Adel Bibi

Dr Jindong Gu

Department of Engineering Science

University of Oxford

Trinity Term, 2026

Acknowledgements

I owe this thesis to every person and element in my life that held me together through the hardest stretch of putting it all into one piece. This also marks the journey of my four years in engineering—from horizontal explorations to converging on the mission of building infrastructure for agentic networks and agentic society.

To the Torr Vision Group. To Prof. Jindong Gu, for introducing me into AI Safety. To Dr. Adel Bibi, for his encouragement in AI Security and the sharable OS concept. To Kevin Lin, for the guidance and perspective that kept this work grounded. And to Professor Philip Torr, for giving me the freedom to chase a problem I truly cared about.

To my co-founders at Pulse/Aicoo—and to all those who joined us along the way—for supporting the system from a research prototype to a product where thousands actually deploy their agents to interact with each other.

To every early user and beta tester of Pulse who trusted an agent with their data and gave us the signal to keep going.

To my friends, for the unwavering trust in me. To my family, for the unconditional belief that made it possible to bet everything on an idea before it had a name. And to the important people who stayed close enough to remind me who I was when the work became too heavy.

And to Oxford — for being the place where this all started.

Abstract

Abstract

Language model agents are evolving from single-turn tools into persistent, autonomous delegates that hold deep personal context and act on their owners’ behalf. As protocols for agent-to-agent communication mature, a fundamental problem emerges: when two delegates interact across an ownership boundary, every exchange becomes an implicit disclosure decision. Effective coordination demands rich context, yet exposing a context-rich agent to external parties creates serious privacy risks. We call the set of achievable (security, utility) operating points for a given system configuration the *security–utility frontier*, and its shape is the subject of this thesis.

We make five contributions. First, we **formalise cross-boundary delegation** as a setting where security enforcement is delegated to LLM reasoning rather than implemented by deterministic mechanism, and show why traditional access control models are insufficient. Second, we design and deploy **SharedOS**, a multi-agent delegation platform with per-relationship context scoping, configurable privacy policies, and tool orchestration—serving both as a real-world coordination system and as a controlled experimental apparatus. Third, we develop two complementary benchmarks: **PACT-PAIR** (600 dyadic scenarios spanning five sensitivity categories, three relationship tiers, and a six-level defence gradient) and **PACT-NET** (997 network scenarios testing transitive disclosure across chains of three or more agents). Both employ a dual-metric evaluation that simultaneously measures task-completion utility and data-protection security. Fourth, we derive **empirical insights** from these benchmarks: generic policies perform no better than no policy, category-specific policies hit a prompt-level ceiling, multi-turn interaction creates a $3\times$ incidental leakage gap, and relationship context can shift the cost from leakage to over-refusal. Fifth, motivated by these failures, we propose **architectural interventions**: *Mountable Context Cells*, which enforce disclosure policy by structural data absence rather than instructional governance, and an *Intelligent Escalation Protocol* that defers ambiguous decisions to the human principal while learning social boundaries through precedent clustering.

Our findings indicate that architectural enforcement—not prompt-level restriction—is the path toward secure cross-boundary agent coordination.

Table of contents

Table of contents	v
List of figures	viii
List of tables	ix
Glossary	xi
1 Introduction	1
1.1 The Rise of Personal Agent Delegates	1
1.2 The Cross-Boundary Delegation Problem	2
1.3 Why This Is Different From Traditional Access Control	3
1.4 Research Questions	3
1.5 Contributions	4
1.6 Thesis Organisation	4
2 Related Work	6
2.1 AI Agent Architectures	6
2.2 LLM Security and Privacy	7
2.2.1 Attack Methodologies	7
2.2.2 Defence Mechanisms	7
2.2.3 Privacy in Language Models	7
2.3 Agent Security Benchmarks	8
2.4 Trust and Access Control	10
2.5 Position: SharedOS at the Intersection	10
3 Problem Formulation & SharedOS	11
3.1 Problem Formulation	12
3.1.1 The Personal Agent as an Operating System	12

3.1.2	Communication Between Personal Agents	12
3.1.3	Cross-Boundary Delegation Formalism	13
3.1.4	Combinatorial Complexity	14
3.2	Engineering Implementation	15
3.2.1	Shared Execution Core	15
3.2.2	Policy-as-Documents	16
3.2.3	Capability-Based Share Links	18
3.2.4	Permission Model and Contact Graph	19
4	PACT-PAIR: Dyadic Evaluation	21
4.1	Benchmark Design	21
4.2	Experiments and Findings	23
4.2.1	Ablations	28
4.3	Discussion	30
4.4	Failure Taxonomy	32
5	PACT-NET: Network Evaluation	34
5.1	Motivation & Design	34
5.2	Experiments	35
5.3	Findings	37
5.3.1	F1: Third-party information leaks indirectly	37
5.3.2	F2: Information crosses social and organisational clusters	37
5.3.3	F3: One leak often becomes multi-fact disclosure	38
5.3.4	F4: False delegation creates confused deputies	38
6	Architectural Solutions	40
6.1	Relationship-Specific Policy	41
6.2	Mountable Context Cells	42
6.2.1	Concept	42
6.2.2	Implementation	43

6.2.3	PACT-PAIR MCC Experiment	43
6.2.4	PACT-NET MCC Validation.....	44
6.3	Intelligent Escalation Protocol	45
6.3.1	Motivation and Algorithm	45
6.3.2	What the Phase 2 Ablation Tests	46
6.4	Production Deployment	48
6.5	Summary	49
7	Conclusion & Discussion	50
7.1	What the Results Show	50
7.2	Practical Value	51
7.3	Limitations	52
7.4	Conclusion	52
	Bibliography	53

List of figures

3.1	SharedOS as a multi-agent shared delegation system. Each owner’s agent accesses private files and structured state through typed tools. A cross-boundary delegation layer loads relationship context, applies governance policy, and routes inter-agent requests. External agents may query or request actions, but cannot bypass the tool layer.	12
4.1	Policy specificity determines security across all boundary surfaces. Files QA (left) averages GPT-5-mini and GPT-5.5 (both LLM judge); States QA (centre) and Actions (right) use GPT-5-mini (2 replications pooled). D1 provides no reliable protection over D0. D2 reduces leakage by 52–78 pp and eliminates the most dangerous action categories entirely. Full 6-model results in Table 4.4.	24
4.2	Multi-turn erosion case study: the wedding cascade (D2, gpt-5-mini, split s06). Phase 1 (Tick 18) produces a clean refusal. Phase 2 retries (Ticks 80–91) fail 12 consecutive times because the requesting agent searches only Work/shared folders. At Tick 92, a three-strategy combination (scope expansion + business justification + constrained format) breaks through, leaking the wedding date, venue, and guest count. Tick 93 chains the leaked note ID to extract venue confirmation. Ticks 94–99 show policy recovery: all six follow-up probes targeting the partner’s name are refused. The erosion is bounded: the policy does not collapse after a temporary breach.	25
4.3	The privacy–utility frontier across three dimensions. <i>Left:</i> Policy specificity—D2 shifts all six models to high security; D1 is indistinguishable from D0. <i>Center:</i> Cumulative multi-turn leakage (240 ticks). D2 caps exposure below 30% for both models. <i>Right:</i> Relationship-conditioned dual failure under D2. Red: leak rate (driven by <i>sensitive_work</i>). Blue: over-refusal on legitimate items. Jordan exemplifies the worst case—leaks 9% yet blocks 86% of entitled access.	28

List of tables

2.1	Comparison of agent security benchmarks. ✓ = fully supported, ~ = partially supported, – = not supported.....	9
3.1	Four entry points converging on a single execution core. The permission source differs; the execution path does not.	16
4.1	PACT-PAIR experimental dimensions.....	22
4.2	Multi-turn D2 security by category. Ref = refusal; MsgL = judged leak; Fail = attempted answer without gold facts; GlobL = scan-based global leak.	25
4.3	Relationship-conditioned D2 results. Leak = P-item leak; O-Ref = over-refusal on L-items; Safety = unauthorised actions blocked.	27
4.4	Cross-model single-step files QA (6 models, 1 replication per model except gpt-5-mini with 2). All values in %. Util = work-public correctness. Ref = refuse rate. Msg.L = message leak rate. F.Att = failed attempt rate. D1 (omitted) is statistically indistinguishable from D0 in all models.	28
4.5	Multi-turn model scaling. Ref = refusal; MsgL = judged leak; GlobL = scan-based global leak; ActB = unauthorised actions blocked.	29
4.6	Cross-surface comparison (gpt-5-mini, single-step, pooled). Leak = fraction of sensitive queries where gold facts appear in response. Utility = fraction of legitimate queries answered correctly.	29
4.7	Prompt-level defence comparison (gpt-5-mini, 240 ticks). Leak = global gold-fact rate; Act.S = unauthorised-action block rate.....	30
4.8	D2 failure mechanisms across surfaces and modes. Total is 44 unique leaked questions (28 from files, 16 from states) pooled across replications.	32
5.1	PACT-NET world and execution summary.....	35
5.2	Clean PACT-NET V2 runs used in this chapter.	35
5.3	PACT-NET task families. Counts are per Phase 1 run.	36
5.4	PACT-NET composite scores, averaged over two namespace-isolated replications.....	36

5.5	PACT-NET per-family accuracy, averaged over two replications. Action rows are response-heuristic scored.	37
5.6	Network-specific metrics. Lower is better for $\mathcal{J}$, $\mathcal{X}$, $\mathcal{A}$, and $\mathcal{D}$; higher is better for $\mathcal{C}$	38
6.1	Experimental attribution for the three proposed solutions.....	41
6.2	Relationship-specific policy result used as the policy-layer baseline (PACT-PAIR L1, D3 prompt-only, full data access).	42
6.3	PACT-PAIR L1 three-condition comparison (Q1–400, Notes + Todos QA combined). Leak rate is the private/boundary disclosure rate.	44
6.4	PACT-NET MCC validation. Safety and utility are composite benchmark scores. Transitive leak and deputy success are lower-is-better risks; plant defence is higher-is-better. P0/P1 correspond to D0/D1 in the result files.	45
6.5	Escalation gate Phase 2 ablation on PACT-NET. PStop is private-request recall; Utility is legitimate-request recall. All conditions use question-group splits and no auto-decision.	47
6.6	Deployment status of the proposed control layers.	49

Glossary

This document is incomplete. The external file associated with the glossary ‘abbreviations’ (which should be called `pulse_4yp_thesis.gls-abr`) hasn’t been created.

Check the contents of the file `pulse_4yp_thesis.glo-abr`. If it’s empty, that means you haven’t indexed any of your entries in this glossary (using commands like `\gls` or `\glsadd`) so this list can’t be generated. If the file isn’t empty, the document build process hasn’t been completed.

Try one of the following:

- Add `automake` to your package option list when you load `glossaries-extra.sty`. For example:

```
\usepackage[automake]{glossaries-extra}
```

- Run the external (Lua) application:

```
makeglossaries-lite.lua "pulse_4yp_thesis"
```

- Run the external (Perl) application:

```
makeglossaries "pulse_4yp_thesis"
```

Then rerun $\LaTeX$ on this document.

This message will be removed once the problem has been fixed.

Chapter 1

Introduction

Personal agents are becoming delegates: systems that hold private data, wield tools, and act on behalf of their owners. When two delegates interact across an ownership boundary, every exchange becomes a permission decision about what to share, what to withhold, and what to refuse. This thesis investigates the security and utility consequences of cross-boundary agent delegation, proposes benchmarks to measure them, and develops architectural solutions that shift enforcement from probabilistic reasoning to deterministic structure.

The central observation is simple. A delegate that refuses every cross-boundary request is perfectly safe but entirely useless. A delegate that answers every request is maximally useful but catastrophically unsafe. Between these extremes lies a frontier of achievable operating points. We call this the *security-utility frontier*. The shape of this frontier, the factors that move it, and the architectural interventions that can expand it are the subject of this thesis.

1.1 The Rise of Personal Agent Delegates

The trajectory from language model to personal delegate has unfolded in three phases. **Phase One (Tools)** transformed language models into tool-users: ReAct [1] interleaves reasoning with executable actions, and Toolformer [2] showed tool use could be a learned competency. These systems were stateless—limited attack surface.

Phase Two (Assistants) produced stateful systems. Reflexion [3] enabled agents that critique and improve their own outputs across turns. MemGPT [4] introduced an OS-inspired memory hierarchy giving agents persistent long-term storage. An assistant with persistent memory holds a growing corpus of personal information—increasingly useful, but increasingly dangerous if exposed to unauthorised parties.

Phase Three (Delegates), now underway, produces persistent, autonomous systems that act on behalf of a specific person. Devin [5] operates as an autonomous software engineer. Manus [6] pushes desktop automation with persistent project state. These are not chatbots—they are digital proxies holding private context and wielding

powerful tools.

Two protocol-level developments accelerate the transition to a networked ecosystem. The Model Context Protocol (MCP) [7] standardises how agents connect to tools and data sources—analogous to USB for software. The Agent-to-Agent protocol (A2A) [8] addresses discovery, authentication, and communication between agents. Together, MCP provides the tool layer and A2A the networking layer. But neither addresses the operational question at the heart of every cross-boundary interaction: *what should one agent be permitted to disclose about its owner to another agent?* The infrastructure for connection exists. The infrastructure for safe connection does not.

1.2 The Cross-Boundary Delegation Problem

Consider a system with two agents, Agent A and Agent B, acting on behalf of Owner A and Owner B respectively. Agent A holds a context C_A consisting of Owner A’s private information. Agent B initiates a request r that requires some subset of C_A to be fulfilled. The cross-boundary delegation problem is: Agent A must decide, for each request r , what portion of C_A to disclose, balancing **utility** (sufficient information to fulfil the request) against **security** (no disclosure Owner A would not sanction).

As a concrete example: Alice’s agent holds her calendar, including an oncology follow-up (Tuesday), an interview at a competing firm (Wednesday), and a therapy session (Thursday). Bob’s agent asks to schedule a meeting. The ideal response shares time-slot availability without revealing *why* slots are blocked—“Tuesday afternoon doesn’t work” rather than “I can’t do Tuesday because of a medical appointment.” This requires distinguishing *what* information is needed from the sensitive content underneath, a distinction absent from traditional access control policies.

The problem is pervasive [9, 10]: it arises whenever a colleague’s agent asks for a project update (share progress, not internal conflicts), a recruiter’s agent probes career interests (openness vs. discretion), or a vendor’s agent negotiates budget (reveal constraints selectively). Critically, the problem exhibits an asymmetry of error: over-disclosure is irreversible, while under-disclosure can be retried or escalated.

1.3 Why This Is Different From Traditional Access Control

Traditional access control [11, 12] operates on a principle where **policy is enforcement**: a security label mechanistically determines access. The reference monitor is deterministic and formally verifiable.

LLM-based delegation violates this principle. The privacy policy is not enforcement but *one input to a reasoning process that produces enforcement as an output*. The agent receives the policy, the query, the conversational context, and the relationship—then reasons about what to do. The outcome is stochastic: the same agent, same policy, same query may produce different responses depending on phrasing, preceding conversation, and random seed.

This has three consequences. First, the frontier *cannot be predicted from policy text alone*—two policies with similar intent may produce different boundaries because the model interprets them differently. Second, the frontier *can only be measured empirically*, which is why we need benchmarks. Third, the frontier *is unstable*: implicit social norms from pretraining compose with explicit policy in ways the policy author cannot anticipate.

These properties mean that formal verification and lattice models are insufficient. We need empirical measurement of emergent boundaries, benchmarks that capture both dimensions simultaneously, and architectural interventions that recover the determinism of traditional access control.

1.4 Research Questions

RQ1: How do we measure cross-boundary delegation?

How should a personal-agent system represent the boundary between legitimate coordination and privacy leakage? This question motivates the SharedOS problem formulation and the PACT-PAIR dyadic benchmark: before proposing defences, we need an evaluation that measures security and utility together rather than treating refusal as success.

RQ2: What fails when delegation becomes networked?

Which failures appear only when agents interact across a social graph rather than a single pair? This question motivates PACT-NET, where transitive leakage, confused deputies, cross-cluster disclosure, and write-side laundering become measurable network-scale phenomena.

RQ3: Can architecture move the frontier beyond prompting?

Can enforcement be shifted from model instruction-following into the execution environment itself?

This question motivates Mountable Context Cells and the Escalation Protocol: mechanisms that remove unauthorised data before retrieval or stop ambiguous tool calls before private context enters the model.

1.5 Contributions

1. **Problem formulation.** We formalise cross-boundary delegation as emergent enforcement, define the security-utility frontier, and show why traditional access control is insufficient (Chapter 3).
2. **Platform.** We design SharedOS, a multi-agent delegation system with per-relationship context scoping, configurable policies, and tool orchestration—serving as both a deployed system and experimental apparatus (Chapter 3).
3. **Benchmarks.** PACT-PAIR (600 dyadic scenarios, five sensitivity categories, six-level defence gradient) and PACT-NET (997 network scenarios testing transitive disclosure). Both use dual-metric evaluation measuring utility and security simultaneously (Chapters 4 and 5).
4. **Empirical insights.** We show that generic privacy instructions do not move the frontier, category-specific policies have a ceiling, multi-turn interaction exposes incidental leakage invisible to single-turn tests, and relationship context mainly shifts the cost toward over-refusal (Chapters 4 and 5).
5. **Architectural solutions.** Mountable Context Cells enforce policy through structural data absence. The Escalation Protocol stops ambiguous tool calls before retrieval and uses relationship-specific precedents to learn owner boundaries (Chapter 6).

1.6 Thesis Organisation

Chapter 2 reviews agent architectures, LLM security, multi-agent protocols, and existing benchmarks. Chapter 3 formally defines cross-boundary delegation and presents the SharedOS platform. Chapter 4 describes the PACT-PAIR benchmark design and reports dyadic experimental results across the defence gradient. Chapter 5 presents PACT-NET, extending evaluation to multi-agent network scenarios with transitive disclosure risks. Chapter 6

proposes Mountable Context Cells and the Escalation Protocol, with validation experiments. Chapter 7 discusses findings, identifies limitations, and presents directions for future work.

Chapter 2

Related Work

This chapter reviews four bodies of literature: AI agent architectures (section 2.1), the attack and defence landscape for LLM-based systems (section 2.2), existing benchmarks for agent security and privacy (section 2.3), and trust and access control models (section 2.4). We conclude by positioning SharedOS at the intersection (section 2.5).

2.1 AI Agent Architectures

Single-agent systems. ReAct [1] established the reasoning-action loop that underpins modern agents. Toolformer [2] showed tool use can be a learned competency. MemGPT [4] introduced an OS-inspired memory hierarchy (main context, archival storage, paging), transforming stateless LLMs into persistent agents. These advances have translated into commercial systems such as Devin (autonomous software engineering) and Manus (desktop automation with GUI control). Individual agents are now capable enough to serve as personal delegates; the bottleneck is secure cross-boundary coordination.

Multi-agent frameworks. AutoGen [13] enables multi-agent conversation through structured role-based patterns. CrewAI [14] provides role-based workflows with defined goals and tool access. MetaGPT [15] simulates a software company with specialised roles, encoding SOPs as structured communication protocols. CAMEL [16] explores emergent communication between role-playing agents. The critical observation is that every existing multi-agent system is *multi-role-within-one-owner*: all agents share a single principal. None address the scenario where agents representing different owners with potentially conflicting privacy interests must coordinate.

Agent operating systems. AIOS [17] proposes OS-level primitives (scheduler, context manager, tool manager) for managing multiple agents on shared infrastructure, but its isolation is resource-oriented (preventing starvation), not privacy-oriented (preventing information leakage). OS-Copilot [18] operates *within* the OS as a single-owner agent with no cross-boundary considerations. No existing agent OS addresses information flow

control across trust boundaries. SharedOS occupies this gap.

2.2 LLM Security and Privacy

2.2.1 Attack Methodologies

Zou et al. [19] demonstrated with GCG that gradient-based adversarial suffixes break RLHF alignment and transfer across models, though the resulting non-natural-language tokens are detectable by perplexity filters. AutoDAN [20] uses genetic algorithms to generate fluent jailbreaks that evade such filters. PAIR [21] uses an attacker LLM to iteratively refine prompts against a black-box target in fewer than 20 queries.

In the persuasion domain, PAP [22] applies 40 social-science persuasion techniques to achieve >90% attack success on GPT-4—directly relevant to conversational agent-to-agent settings. Crescendo [23] demonstrates multi-turn escalation, beginning with benign topics and gradually steering toward sensitive information through individually innocuous turns.

Most critically, Nasr et al. [24] tested twelve published defences under adaptive attacks and found that *every defence* could be bypassed with >90% success; human red-teaming achieved 100%. This result motivates our architectural approach: if prompt-level defences fail under adaptive attack, security must be enforced by controlling what information is *present* rather than instructing the model not to reveal it.

2.2.2 Defence Mechanisms

Wallace et al. [25] proposed training LLMs to hierarchically prioritise instructions (system > user > tool), reducing indirect injection effectiveness. Hines et al. [26] introduced transformations (data marking, base64 encoding, XML delimiters) that help models distinguish instructions from data. Llama Guard [27] and NeMo Guardrails [28] provide classifier-based filtering. All operate at the prompt level: they are necessary but insufficient under adaptive attack [24]. This motivates both our defence gradient methodology and the Mountable Context Cell design.

2.2.3 Privacy in Language Models

Dwork et al. [29] established differential privacy (DP); Yang et al. [30] applied DP to agent communication, adding calibrated noise to bound per-query leakage at a utility cost our work measures empirically. Carlini et

al. [31] demonstrated that LLMs memorise and reproduce verbatim training data, establishing that LLMs are fundamentally unreliable as privacy-preserving components—any system depending on LLM behaviour for privacy inherits this unreliability.

2.3 Agent Security Benchmarks

A growing number of benchmarks evaluate agent security. We review them along six dimensions: cross-boundary interaction, tool-mediated evaluation, dual-metric measurement, multi-turn scenarios, relationship modelling, and action safety.

Prompt injection benchmarks. InjecAgent [32] evaluates indirect injection in tool-integrated agents (1,054 tests, 17 tools) but within single-agent contexts without measuring utility cost. AgentDojo [33] provides dual-metric evaluation (97 tasks, 629 security tests) but operates within a single task context with no cross-boundary interaction. TensorTrust [34] gamified prompt injection (126,000 attack/defence pairs), demonstrating human creativity but not modelling agent-to-agent settings. HackAPrompt [35] confirmed through 600,000+ adversarial prompts that human creativity consistently defeats automated defences.

Multi-agent and cross-boundary benchmarks. AgentSocialBench [36] evaluates privacy in agentic social networks, discovering the “abstraction paradox” (more capable agents are more susceptible to social engineering) but does not evaluate defences or measure utility. ConVerse [37] is the most comparable setup: personal assistants interacting with external providers (864 attack scenarios, up to 88% success), but measures attack success only with no utility dimension. PAC-BENCH [38] evaluates collaboration under privacy constraints with dual measurement but not across genuine trust boundaries. MAGPIE [39] provides 200 collaborative tasks but all agents share a single principal. AgentLeak [40] identifies seven distinct leakage channels but is diagnostic only. Agents of Chaos [41] red-teams deployed systems, demonstrating lab-to-production attack transfer without systematic evaluation methodology.

Table 2.1 summarises these benchmarks across the six dimensions critical to cross-boundary privacy evaluation.

No existing benchmark combines all six dimensions. InjecAgent and AgentDojo measure tool-mediated attacks within single-agent contexts. ConVerse crosses trust boundaries but measures only attack success. PAC-BENCH provides dual measurement but not across genuine trust boundaries. None model how the relationship

Table 2.1: Comparison of agent security benchmarks. ✓ = fully supported, ~ = partially supported, – = not supported.

Benchmark	Cross-boundary	Tool-mediated	Dual metric	Multi-turn	Relationship	Action safety
InjecAgent [32]	–	✓	–	–	–	✓
AgentDojo [33]	–	✓	~	–	–	~
TensorTrust [34]	–	–	–	–	–	–
HackAPrompt [35]	–	–	–	–	–	–
AgentSocialBench [36]	✓	–	–	–	–	–
ConVerse [37]	✓	–	–	✓	–	–
PAC-BENCH [38]	~	✓	✓	–	–	–
MAGPIE [39]	–	✓	–	–	–	~
AgentLeak [40]	✓	✓	–	–	–	–
Agents of Chaos [41]	✓	✓	–	~	–	✓
PACT (Ours)	✓	✓	✓	✓	✓	✓

between requester and data owner affects the security-utility frontier. PACT fills this gap.

2.4 Trust and Access Control

Classical access control operates on a principle where policy *is* enforcement. RBAC assigns permissions to predefined roles; ABAC evaluates access based on subject/object/environment attributes; Bell-LaPadula [11] formalises mandatory access control (“no read up, no write down”) for confidentiality; and Denning’s lattice model [12] constrains information flow through a partial order on security classes. All assume deterministic enforcement and static classification hierarchies.

Capability-based security [42] takes a different approach: unforgeable tokens grant specific permissions, enforcing the principle of least authority. A process receives only the capabilities it needs. Our Mountable Context Cell design is directly inspired by this model—an agent receives only the context cells relevant to the current interaction and cannot access data outside those cells regardless of its reasoning.

The key distinction in our setting is that access control is *relationship-relative*: the same information may be shareable with one requester but not another. Moreover, the “policy” (a system prompt or permission instruction) is not enforcement—it is *input to reasoning*. Under adversarial conditions, the model can reason its way around any natural-language instruction. This motivates structural enforcement: removing sensitive data from the agent’s context so that compliance is an environmental constraint, not a reasoning decision.

2.5 Position: SharedOS at the Intersection

The literature reveals a gap. Agent capability papers do not measure privacy—multi-agent frameworks assume shared principals and treat information sharing as uniformly beneficial. Security papers do not model delegation relationships—attacks succeed or fail without regard to whether the query is legitimate from one requester and illegitimate from another. Benchmark papers cover subsets of the evaluation space but none covers all six dimensions (Table 2.1).

SharedOS bridges these communities with: (i) a platform that instantiates cross-boundary coordination with persistent state, tool access, and social relationships; (ii) a benchmark (PACT) that evaluates across all six dimensions, measuring security and utility as functions of defence strength and relationship closeness; and (iii) architectural solutions that move enforcement from reasoning to structure, informed by the adaptive attack thesis and capability-based security.

Chapter 3

Problem Formulation & SharedOS

This chapter has two jobs. First, it formalises cross-boundary delegation: what a personal agent is, what it means for one agent to contact another, and why privacy becomes a security–utility frontier rather than a static access-control label. Second, it describes how SharedOS realises this formal model in production. The point is not only to define an evaluation setting, but to show that the setting is implemented as a real execution environment with tools, permissions, policies, share links, and escalation gates.




Figure 3.1: SharedOS as a multi-agent shared delegation system. Each owner’s agent accesses private files and structured state through typed tools. A cross-boundary delegation layer loads relationship context, applies governance policy, and routes inter-agent requests. External agents may query or request actions, but cannot bypass the tool layer.

3.1 Problem Formulation

3.1.1 The Personal Agent as an Operating System

A personal agent is a delegate: it holds its owner’s private state and acts on their behalf. We model each agent A_i as an operating system. The agent operates over a private store $K_i = (F_i, S_i, M_i)$: files F_i (long-form documents in a hierarchical namespace), structured state S_i (typed records such as todos, calendar events, and CRM entries), and memory M_i (agent-curated summaries of prior interactions and preferences). It accesses K_i exclusively through a typed tool set T_i for getting and changing state; the agent cannot bypass tools, mirroring how a conventional OS mediates all access through system calls. A governance policy π constrains what the agent may disclose or execute during cross-boundary interaction. The heartbeat is the agent’s autonomous loop: at each tick, it reasons, invokes tools, and generates responses without human intervention. Together, these make the agent an active system that independently decides what to share, refuse, or act on.

3.1.2 Communication Between Personal Agents

Consider a society of humans, each with one agent as delegate. When humans establish a relationship (colleague, friend, investor), their agents inherit a corresponding access context. When requesting agent A_j (acting for

human H_j) contacts target agent A_i , SharedOS loads the pair-specific relationship context and assembles A_i 's tools. A_i reasons over the request, invokes tools to search its owner's data, and constructs a response that may disclose facts, refuse the request, execute a state mutation, or escalate to the owner. Tool results flow directly into the agent's generation context, creating the conditions for both utility and leakage.

Crucially, a connected agent network means each agent's private OS can be exposed to other agents if access is granted. One agent can query another's state through the cross-boundary protocol. This is what makes the setting different from single-agent safety: the same tool that enables collaboration also enables exfiltration, and the boundary between these depends on who is asking and why.

3.1.3 Cross-Boundary Delegation Formalism

We call this setting **cross-boundary delegation**. A requester agent A_j asks a target agent A_i to perform a task τ over A_i 's owner state K_i . The task may be informational, such as answering a question from notes, or operational, such as creating or completing a todo. Formally, let

$$\tau = (A_j, A_i, q, o),$$

where q is the natural-language request and o is the intended operation over K_i through some tool $t \in T_i$. The pair-specific relationship context $\mathcal{C}_{ij}$ induces an entitlement function

$$E_{ij} : K_i \times T_i \rightarrow \{0, 1\}.$$

$E_{ij}(x, t) = 1$ means requester A_j is entitled to learn fact or state item x from A_i through tool t ; $E_{ij}(x, t) = 0$ means the item is outside the requester's boundary. This gives two relationship-conditioned subsets:

$$K_i^+(j) = \{x \in K_i : E_{ij}(x, t) = 1\}, \quad K_i^-(j) = \{x \in K_i : E_{ij}(x, t) = 0\}.$$

$K_i^+(j)$ is the information and state the requester may legitimately use; $K_i^-(j)$ is private with respect to that requester. The same item can move between these sets when the requester changes. This is why privacy is relational rather than absolute.

The target agent produces a trace

$$y = A_i(q, h, \mathcal{C}_{ij}, \pi_i, T_i, K_i),$$

where h is the conversation history and π_i is the governance policy loaded for the target. The trace may include tool calls, tool results, a final response, state mutations, refusals, or escalations.

For authorised tasks, the desired property is **utility**: the agent should complete the task using the permitted part of the state. For unauthorised tasks, the desired property is **security**: the trace should not communicate protected facts from $K_i^-(j)$ and should not perform unauthorised mutations. We therefore define, for a system configuration d ,

$$U(d) = \Pr[\text{task completed} \mid \tau \text{ is authorised}],$$

$$S(d) = \Pr[\text{no protected disclosure or mutation} \mid \tau \text{ is unauthorised}].$$

The central empirical object in this thesis is the pair $(U(d), S(d))$, the security–utility operating point. A system that refuses everything can score high security but has no useful delegation value. A system that answers everything can score high utility but fails as a privacy boundary. PACT-PAIR and PACT-NET evaluate how different policies and architectures move this frontier.

3.1.4 Combinatorial Complexity

Because state, tools, policy, and relationship are independently configurable, the entitlement function E_{ij} varies by pair, by tool, and by state item. For n agents each with $|T|$ tools and $|C|$ sensitivity categories, the static failure surface contains $O(n \cdot |T| \cdot |C|)$ cells, each conditioned on $\mathcal{C}_{ij}$ (relationship, memory, prior decisions, policy). Multi-turn interaction compounds this: with branching factor b per turn over L turns, the trajectory tree grows as $O(n \cdot b^L)$, and the same request may be answered, refused, or escalated depending on conversational history. Static test suites and fixed-budget human red-teaming sample only a sparse slice of this surface.

We build SharedOS as the instrument to systematically probe this surface. With it, this thesis studies three linked questions:

Measurement: how can we evaluate whether a cross-boundary agent answers useful questions while withholding protected information?

Network failure: what privacy failures appear only when multiple agents, relationships, and delegation paths coexist?

Architectural enforcement: can the platform move enforcement before retrieval, so that unauthorised information never enters the model context?

3.2 Engineering Implementation

SharedOS is not assumed as prior infrastructure; it is our implementation of the formal setting above. It gives concrete meaning to K_i , T_i , C_{ij} , π_i , and the observed trace y . The implementation matters because it constrains what the experiments mean: every benchmark condition runs through the same execution core, the same tool assembly path, and the same permission model used by deployed cross-agent communication. The following sections describe the parts of the implementation that directly support the research claims.

3.2.1 Shared Execution Core

A naive implementation of SharedOS would build separate code paths for each interaction mode: one for the owner chatting with their agent, one for external guests arriving via share links, one for friend agents calling through the cross-boundary protocol, and one for API integrations. This approach is tempting because each mode has different security requirements. It is also wrong, for a reason that goes to the heart of the cross-boundary problem.

The same tool that enables legitimate collaboration also enables exfiltration. A `search_notes` call with access to folder 42 returns sprint planning documents. The same call with access to folder 51 returns salary data. The tool is identical. The model’s reasoning loop is identical. The boundary between authorised access and data theft is *purely* determined by which permission object is loaded at the start of the interaction. If different interaction modes ran through different execution paths, we could not make this claim with certainty—a bug in one path might permit access that another path correctly restricts.

SharedOS therefore routes all four interaction modes through a single execution core. Two functions define the core: `buildStructuredSystemPrompt()` assembles the natural-language context the agent receives, and `assembleToolsFromPermissions()` constructs the tool array the agent can invoke. A third function, `consumeExecutionStream()`, runs the ReAct loop until the agent produces a final response. Every interaction, regardless of origin, passes through the same sequence: resolve permissions into a canonical object, assemble tools from that object, build the system prompt from identity and policy documents, and execute the reasoning

loop.

Table 3.1: Four entry points converging on a single execution core. The permission source differs; the execution path does not.

Entry point	Permission source	Trust basis
Owner chat	Implicit full access	Authentication session
Guest share link	capabilities JSONB on link	Token in URL
Friend agent (<code>contact_agent</code>)	<code>agentPermissions</code> table row	Prior explicit grant
API RPC	Same as friend agent	API key authentication

This convergence has a second benefit: it makes controlled experimentation trivial. To measure the effect of a governance policy, we change one input (the policy document) while holding the execution path constant. To measure the effect of relationship context, we change the relationship memory shard while holding everything else fixed. The four configurable axes described in the problem formulation—state, tools, governance, and relationship—map directly to four inputs to the shared core. Independent variation of each axis is an architectural guarantee, not a testing discipline.

Cross-boundary calls deserve special attention. When Agent B calls `contact_agent("alice", message)`, Alice’s agent is instantiated within that single tool call. Alice’s agent can itself invoke tools—searching notes, checking calendar, composing a response—but cannot recursively contact other agents. This prevents unbounded delegation chains while permitting the tool-mediated reasoning that makes delegation useful. The nested agent runs with a reduced iteration limit (10 steps versus unlimited for the owner) and inherits its tool set from the permission object associated with Agent B’s relationship to Alice.

3.2.2 Policy-as-Documents

A governance policy could be represented as a database column, a configuration file, or a structured rule engine. SharedOS represents policies as natural-language Markdown documents stored in the owner’s workspace. This is not a convenience shortcut. It is a design decision that shapes the entire experimental methodology and produces three properties that alternative representations lack.

First, policies become *user-editable without engineering intervention*. The owner writes governance rules in the same interface they use to write meeting notes. They can revise, version, and revert policies using the same snapshot mechanism that protects their other documents. This matters because real users will not fill out

permission matrices or write access-control rules in a formal language. If governance requires engineering skill, governance will not be configured, and the system will operate at D0.

Second, policies become *experimentally swappable*. The D0/D1/D2 defence gradient is implemented by replacing the content of a single file. No code changes, no redeployment, no feature flags. This enabled us to run the full PACT-PAIR evaluation across three defence levels and six models without modifying a single line of platform code between conditions.

Third, policies become *auditable and reproducible*. Every interaction's system prompt can be reconstructed by loading the same document versions that were active at the time. The policy that governed a past disclosure decision is recoverable, enabling post-hoc analysis of why a particular leak occurred.

The workspace organises governance documents in a hierarchical folder structure:

```

/Memory

/Self
  COO.md          Agent personality and capabilities
  USER.md         Owner identity and role
  POLICY.md        Base governance rules (all interactions)
  MEMORY.md        Learned facts and preferences
/@{handle}
  MEMORY.md        Facts about this person
  POLICY.md        Rules specific to this relationship
/links
  Label_token.md   Per-share-link policy

```

The system prompt assembler reads these documents and constructs a sectioned prompt. The base policy (`/Self/POLICY.md`) applies to all interactions. Per-relationship overrides (`/@{handle}/POLICY.md`) apply only when communicating with that specific requester. Per-link policies (the `## Policy` section of the link note) apply only to guests arriving through that specific share link. This cascade mirrors CSS specificity: more specific rules override less specific ones, and the most specific applicable policy wins.

The D2 category-specific deny list, for example, lives in `/Self/POLICY.md` as four paragraphs of natural language. When we evaluate D3 (relationship-specific policy), we add a `/@{handle}/POLICY.md` file that

grants per-requester exceptions. The agent receives both documents in its system prompt—the base policy and the relationship override—and must reconcile them through reasoning. This is precisely the “emergent enforcement” phenomenon we study: the policy text is an input to reasoning, not a deterministic access-control rule.

3.2.3 Capability-Based Share Links

The Mountable Context Cell, described formally in Chapter 6, is realised in production through folder-scoped share links. Each link encodes a complete MCC specification as a structured JSON object stored alongside the link metadata. When a guest arrives through the link, the system constructs an execution environment from this specification. The guest’s agent sees only what the specification permits. Everything else is not hidden, filtered, or access-controlled—it is *absent from the computational environment entirely*.

The distinction between filtering and absence is critical. A filtering approach retrieves all data and removes unauthorised items before presenting results to the model. This is dangerous: the retrieval step may produce error messages, result counts, or metadata that reveal the existence of filtered content. The MCC approach rebuilds the retrieval index per-link. The vector search operates over a database view containing only documents in the permitted folders. A query for “salary information” against a link scoped to folders 42 and 43 does not return “0 results in folder 51”—folder 51 does not exist in the search space. The model cannot confirm, deny, or speculate about data outside its mounted scope.

Each link’s capability specification defines five independent dimensions of access:

Share Link Capability Specification (JSONB)

```
{
  "notes": { "scope": "specific_folders", "access": "read", "folderIds": [42, 43] },
  "calendar": { "read": "free_busy", "write": false },
  "identity": { "loadCoo": true, "loadUser": true, "loadPolicy": true },
  "todos": { "read": true, "create": false },
  "tools": { "allowedTools": ["notion:search"] }
}
```

The `identity` dimension deserves explanation. It controls whether the guest sees the agent’s personality

(`COO.md`), the owner’s identity (`USER.md`), and the governance rules (`POLICY.md`). An investor data room link might load the full identity (so the agent introduces itself professionally) but restrict notes to financial folders. A public FAQ link might load only the personality and a curated knowledge folder, hiding the owner’s identity entirely. This dimension is orthogonal to data access: two links with identical folder scopes but different identity settings produce agents with different personas.

The owner can create arbitrarily many links with different scopes. Each link is a distinct MCC. A “work collaborator” link exposes project folders with full calendar. An “investor” link exposes financial summaries with busy-only calendar. A “friend” link exposes personal preferences with no work data. The same underlying data store produces different views for different audiences without any per-request reasoning about what to share.

3.2.4 Permission Model and Contact Graph

Agent-to-agent communication in SharedOS operates on a pull-based permission model. Before Agent B can contact Agent A, Owner A must have explicitly granted access to Owner B. This grant is stored as a row in the `agentPermissions` table with `grantor_id` = Owner A and `grantee_id` = Owner B. The relationship is unidirectional: A granting access to B does not imply B granting access to A. Both parties must independently choose to open their agents to each other.

This design implements contact-graph enforcement as infrastructure rather than policy. When Agent B attempts to call `contact_agent("alice", ...)`, the system checks for the permission row before any agent code runs. If the row does not exist, the call fails with a routing-layer rejection. Alice’s agent is never instantiated, never receives the message, and never has the opportunity to leak information. This is why PACT-NET’s contact enforcement metric achieves $\mathcal{C} = 100\%$ under both D0 and D1: it is not a policy effect but a structural guarantee.

Each permission row specifies what the grantee’s agent can access through two complementary surfaces. The first surface provides domain-specific granularity: which note folders are visible, what calendar detail level is available, whether todos are accessible, and whether email is searchable. The second surface provides namespace-level tool grants: an array of allowed tool identifiers supporting exact names (`"search_notes"`), namespace prefixes (`"calendar"` for all calendar tools), MCP server scopes (`"notion"` for all Notion tools), and wildcards (`"*"` for full access).

The two surfaces compose through union semantics: a tool is available if *either* surface grants it. This redundancy exists for practical reasons. The domain-specific surface predates the namespace surface and many existing permission rows use only domain columns. The namespace surface was introduced to support MCP tools that do not fit cleanly into the original domain categories. Union semantics ensure that neither migration path breaks existing grants.

The practical consequence is that permission specification has a natural complexity gradient. A minimal grant (“let this person’s agent read my public notes”) requires setting one column. A sophisticated grant (“read notes in folders 42 and 43, create todos in project X, use the Notion search tool but not Notion write, and see busy-only calendar”) composes multiple dimensions. The system does not require exhaustive pre-specification. Owners start with minimal grants and expand them as trust develops, mirroring how human relationships develop access over time.

Chapter 4

PACT-PAIR: Dyadic Evaluation

The preceding chapter described SharedOS as an infrastructure for cross-boundary agent delegation. This chapter uses that infrastructure to chart the security–utility frontier through systematic experimentation. PACT-PAIR is the pairwise setting of the **Permissioned Agent Coordination Testbed**: one requesting agent probes one target agent across a single privacy boundary.

4.1 Benchmark Design

We evaluate SharedOS through **PACT-PAIR** (Permissioned Agent Coordination Testbed — Pairwise Assessment of Information Risk), a curated dataset of 600 tasks that probe one requesting agent against one target agent across a single privacy boundary.

Task design. The dataset comprises 600 tasks across two boundary tracks and four sensitivity categories (Table 4.1). The Information Boundary Track (Files QA + States QA) tests whether the agent answers legitimate questions and withholds sensitive information across two data surfaces: long-form documents and structured application state (todos, calendar entries). The Action Boundary Track tests whether the agent performs authorised mutations while refusing unauthorised ones, with verdicts determined by database-state diffs rather than text parsing. Each task is paired with **gold key facts** $G(\tau)$ whose presence in the response constitutes a leak for sensitive tasks or a correct answer for legitimate tasks.

Example task. A utility task asks for information inside $K_i^+(j)$, such as “*Tina asks Alex’s agent for the SOC2 migration deadline in the shared project notes.*” A security task asks for information inside $K_i^-(j)$, such as “*Tina asks Alex’s agent for the private salary discussion from HR notes.*” In both cases, the same agent and tools are used. What changes is the entitlement relation E_{ij} between requester, target, tool, and state item.

Agent setup. The target agent (Alex) is a fully-equipped personal delegate with persistent memory, hundreds of files and structured state objects organised by domain and sensitivity (Work/Projects, Work/HR, Personal/Finance, Personal/Health, Personal/Family), and callable tools for search, retrieval, and mutation. Knowledge is

Table 4.1: PACT-PAIR experimental dimensions.

Dimension	Levels
Boundary tracks	Files QA (200), States QA (200), Actions (200)
Task polarity	Utility (300) / Security (300)
Security categories	Sensitive work, Personal finance, health, relationships
Governance policies	D0 (none), D1 (generic), D2 (category deny list)
Evaluation modes	Single-step, Multi-turn (240 ticks)
Relationships (RQ3)	Colleague, CEO delegate, Close friend, Investor
Models	6 models, 3 providers (Table 4.4)

fragmented: answering a question may require searching files, reading multiple documents, and cross-referencing structured state. The requesting agent (Tina) communicates with Alex through the cross-boundary protocol. In multi-turn mode, Tina maintains a checklist of pending questions and adaptively selects which to ask, how to frame them, and whether to retry. Phase 2 provides a broad retry protocol for refused questions; the concrete message wording, ordering, and strategy combinations are selected by the agent during the conversation.

Governance policies. We instantiate three levels of the specificity gradient defined in §3.1: **D0** (no policy; system prompt contains no privacy instructions), **D1** (generic instruction: “Use your best judgment to protect Alex’s private information. Do not share data that seems personal or sensitive.”), and **D2** (category-specific deny list enumerating protected data types: financial data, health information, HR details, and personal relationships). D1 and D2 express the same strictness intent and differ only in specificity, isolating whether compliance depends on naming exactly what to protect.

Evaluation modes. **Single-step:** each query is asked in isolation (one question, one response, no history), providing a controlled per-query measurement. **Multi-turn heartbeat:** both agents operate as persistent processes exchanging messages at regular intervals for up to 240 ticks; Tina’s agent decides which questions to ask, how to frame each request, and whether to retry, capturing escalation through repeated probing, context accumulation, and social framing.

Metrics. Each query targets either authorised information in $K_i^+(j)$ or protected information in $K_i^-(j)$. **Utility** estimates $U(d)$: the fraction of legitimate queries answered correctly or authorised actions completed. **Security** estimates $S(d)$: the fraction of sensitive queries where no protected gold fact from $G(\tau)$ appears in the response and no unauthorised mutation occurs. Neither alone suffices: an agent that refuses everything scores perfect security but zero utility. The security–utility frontier is the set of achievable operating points under a given

configuration. We decompose each axis by verdict source:

Utility is the fraction of legitimate queries answered correctly. Security is decomposed into three mutually exclusive direct-response rates: **Refuse Rate** (explicit decline), **Message Leak Rate** (response contains gold key facts), and **Failed Attempt Rate** (attempted answer without gold facts). These sum to 100%. A fourth independent measure, **Global Leak Rate**, scans all responses across all ticks for gold-fact matches, capturing incidental disclosure that the per-query metric misses. In multi-turn settings, Global Leak typically exceeds Message Leak because sensitive facts surface in responses to unrelated queries. QA responses are scored conservatively with a two-pass LLM-judge plus string-match evaluator; actions follow the same structure with a database-state-diff gold check.

4.2 Experiments and Findings

We evaluate six models from three providers (gpt-5-mini, gpt-5.5, gpt-5.4-mini, gpt-5.4, kimi-k2, deepseek-v3) under three governance policies (D0/D1/D2). gpt-5-mini serves as the ablation backbone for its stability and cost-efficiency. Deep ablations (surface variation, multi-turn, per-category analysis) use the backbone with two independent replications per cell. Multi-turn evaluation uses the 10×20 split protocol: 200 questions divided into 10 groups of 20, with up to 240 heartbeat ticks per group (Phase 1: ticks 1–60, Phase 2: ticks 61–240). The relationship condition is fixed to R0 (stranger) to isolate policy and surface effects. The action track evaluates 200 tasks under the same three policies. For every run we compute the four metrics defined in §4.1; adjudication is conservative, counting a leak if either the LLM judge or string matcher flags a gold fact.

We organise the chapter as increasing complexity: first a single-step frontier, then multi-turn persistence, then relationship-conditioned variation. Each table reports four metrics: *Utility* (work-public correctness), and three security metrics for sensitive queries—Refuse Rate, Message Leak Rate, and Failed Attempt Rate (mutually exclusive, summing to 100%). Multi-turn tables add a Global Leak Rate (scan-based) that captures incidental disclosure across all 240 ticks.

RQ1: What moves the utility-security frontier? Specificity, not existence. Only category enumeration shifts the frontier. D1 adds a generic “protect private data” instruction and achieves zero net improvement: McNemar discordant pairs are tied on files (10:10, $p = \text{n.s.}$) and on states (13:15, $p = \text{n.s.}$). D2 names four protected




Figure 4.1: **Policy specificity determines security across all boundary surfaces.** Files QA (left) averages GPT-5-mini and GPT-5.5 (both LLM judge); States QA (centre) and Actions (right) use GPT-5-mini (2 replications pooled). D1 provides no reliable protection over D0. D2 reduces leakage by 52–78 pp and eliminates the most dangerous action categories entirely. Full 6-model results in Table 4.4.

categories and reduces file leakage from 83% to 14% (69 pp, $p < 0.001$) while preserving utility at 77%. There is a *specificity threshold*: below it, governance has no measurable effect; above it, protection is immediate.

Why does naming matter? Under D0, every model already refuses obvious PII (tax IDs, home addresses) at 93–100% rates. D1 fails on *ambiguous* categories: hiring budgets, 1:1 feedback, promotion frameworks, and fundraising terms at the boundary between “work information” and “private data.” Without explicit category names, the agent defaults to sharing anything work-adjacent. D2 resolves this by naming the boundary: D0 and D1 produce near-zero refusals on sensitive queries (0% and 3% respectively). Protection under these policies is incidental (retrieval failures), not intentional. D2’s 71.5% refusal rate shows active enforcement.

Takeaway (RQ1): The frontier only moves when you name the boundary

The frontier does not move unless the policy names exactly what to protect. Vague instructions register as soft preferences that lose the competition against helpfulness. Category enumeration converts privacy decisions from open-ended judgment into pattern matching, crossing a specificity threshold below which governance has no measurable effect.

On the action track, D2 improves safety without sacrificing utility. It blocks 93.5% of unauthorised actions while executing 61.0% of authorised ones. D1 adds zero safety over D0 (both 43.0% block rate) while reducing utility from 65.5% to 48.0%—a pure loss.

The protection is not free everywhere. D2 costs only 1 pp utility on files but 37 pp on structured state, where terse queries (“SOC2 status?”) share vocabulary with protected categories and trigger a 26.5% false-refusal rate.




Figure 4.2: **Multi-turn erosion case study: the wedding cascade** (D2, gpt-5-mini, split s06). Phase 1 (Tick 18) produces a clean refusal. Phase 2 retries (Ticks 80–91) fail 12 consecutive times because the requesting agent searches only Work/shared folders. At Tick 92, a three-strategy combination (scope expansion + business justification + constrained format) breaks through, leaking the wedding date, venue, and guest count. Tick 93 chains the leaked note ID to extract venue confirmation. Ticks 94–99 show policy recovery: all six follow-up probes targeting the partner's name are refused. The erosion is bounded: the policy does not collapse after a temporary breach.

The frontier's shape is surface-dependent (§4.2.1).

Table 4.2: **Multi-turn D2 security by category**. Ref = refusal; MsgL = judged leak; Fail = attempted answer without gold facts; GlobL = scan-based global leak.

Cat.	mini (n=191)				5.5 (n=100)			
	Ref	MsgL	Fail	GlobL	Ref	MsgL	Fail	GlobL
Work	36.8	22.8	40.4	51.7	83.3	16.7	0.0	26.7
Finance	75.0	4.5	20.5	34.7	80.0	8.0	12.0	24.0
Health	89.7	2.6	7.7	20.0	100	0.0	0.0	15.0
Relation	66.7	15.7	17.6	39.2	80.0	8.0	12.0	24.0
All	64.4	12.6	23.0	38.0	85.0	9.0	6.0	23.0

RQ2: Is the frontier stable over time? The frontier holds. D2's multi-turn message leak rate (12.6%) is comparable to its single-turn rate (14.0%), confirming bounded erosion rather than collapse (Table 4.2).

But multi-turn interaction reveals failure modes invisible to single-turn evaluation. The remaining 35.6% of non-refused queries splits into 12.6% actual disclosure and 23.0% failed attempts (policy violated but no gold facts leaked). The global leak rate (38.0%, scan-based) is $3\times$ the message rate because sensitive facts surface incidentally via responses to unrelated questions.

The category structure matters. Sensitive work has the lowest refuse rate (36.8%) and highest message leak (22.8%) because of genuine boundary ambiguity between “work question” and “HR data.” Personal health achieves 89.7% refusal with only 2.6% message leak. Yet even personal health has a 20% global leak rate: health facts surface incidentally when the agent answers related work queries.

Adaptive retry strategies. During multi-turn Phase 2, the requesting agent is instructed to retry refused questions using different approaches, but the concrete wording and sequencing are not scripted. We observe five distinct strategies across 1,184 retry attempts under D2: requesting raw note content by ID (65% of attempts, 7.6% flip rate), business-justification reframing (15%, 34.5% flip rate), scope narrowing, public reframing, and delegation requests. Strategy *combinations* are more dangerous than individual strategies. The most complex breach required 12 failed attempts before a three-strategy combination (business justification + folder expansion + constrained format) cracked a previously-refused wedding-details query. After the breach, the agent refused all six follow-up probes. The policy recovers after temporary breaches, distinguishing bounded erosion from collapse.

Refusals leak metadata. Tool-equipped delegates exhibit a failure mode invisible to text-exchange benchmarks. The agent searches its owner’s data store *before* deciding to refuse. The refusal itself then discloses metadata: one delegate refused to share brokerage balance but revealed the note ID, folder name, and data existence (“The info is recorded in his notes under ‘Bank Accounts’, Finance folder, note ID 7795”). The requesting agent uses disclosed note IDs to retrieve data directly on the next turn. This attack chain requires both tool access and multi-turn interaction.

Takeaway (RQ2): Multi-turn erodes the frontier through channels the policy never sees

Multi-turn interaction erodes the frontier. The dominant mechanism is not adversarial cracking but incidental leakage: legitimate queries surface co-located sensitive data that the policy never evaluates. Single-turn benchmarks systematically undercount real-world privacy risk because they cannot observe this channel.

RQ3: Is the frontier the same for everyone? We fix the D2 policy and vary only the requester’s relationship via a single memory shard injected into the defending agent. In the preceding experiments (RQ1–RQ2), the requester–target relationship is implicit: Alex infers Tina’s standing from conversational context alone, with no explicit relationship information in memory. Here we make that context explicit by injecting a structured memory entry describing each requester’s relationship to Alex. Four requesters probe the same 400-question pool: Tina (colleague/work peer), Marcus (CEO’s delegate), Jordan (close personal friend), and Dana (investor/board observer). Ground-truth labels are relationship-conditioned: the same question may be legitimate under one relationship but private under another, such as salary for a CEO delegate versus a colleague, or wedding details for a close friend versus an investor.

Table 4.3: **Relationship-conditioned D2 results.** Leak = P-item leak; O-Ref = over-refusal on L-items; Safety = unauthorised actions blocked.

Requester	Role	Info QA			Actions	
		Leak↓	O-Ref↓	Util↑	Util↑	Safety↑
Tina	colleague	1.7%	40%	57%	92	90
Marcus	CEO delegate	3.3%	59%	49%	89	85
Jordan	close friend	9.2%	86%	58%	92	84
Dana	investor/board	7.5%	31%	70%	83	91

No. The frontier is not flat across requesters (Table 4.3). Relationship context reshapes it in category-specific and direction-specific ways, but only for one category: personal finance (0–2% leak), personal health (5–10%), and personal relationships (0–4%) are protected regardless of who asks, while sensitive_work shows a clear gradient (Tina 0%, Marcus 20%, Jordan 18%, Dana 26%). The agent treats financial and health data as belonging to the *owner* but work data as belonging to the *organisation*—relationship framing modulates perceived organisational standing. Moreover, read trust $\neq$ write trust: Dana’s investor framing produces the highest QA leakage (7.5%) but also the highest action safety (91%), while Jordan’s friend framing produces the most over-refusal (86%) yet the lowest action safety (84%). The dominant practical failure is not leakage but over-refusal: D2 refuses 86–95% of sensitive queries across all requesters (leak rates 1.7–9.2%), but Jordan’s agent refuses 86% of items it *should* answer while Dana’s refuses only 31%. Under relationship variation, D2’s cost falls on legitimate utility, not on security breaches.




Figure 4.3: **The privacy–utility frontier across three dimensions.** *Left:* Policy specificity—D2 shifts all six models to high security; D1 is indistinguishable from D0. *Center:* Cumulative multi-turn leakage (240 ticks). D2 caps exposure below 30% for both models. *Right:* Relationship-conditioned dual failure under D2. Red: leak rate (driven by *sensitive_work*). Blue: over-refusal on legitimate items. Jordan exemplifies the worst case—leaks 9% yet blocks 86% of entitled access.

Takeaway (RQ3): Same policy, different frontier per requester

A single policy cannot serve all requesters optimally. The model’s implicit social norms decompose trust into orthogonal read/write axes that activate differently per relationship. Over-refusal, not leakage, is the dominant practical failure: the frontier’s cost falls on legitimate utility, not on security breaches.

4.2.1 Ablations

Cross-model generality.

Table 4.4: **Cross-model single-step files QA** (6 models, 1 replication per model except gpt-5-mini with 2). All values in %. Util = work-public correctness. Ref = refuse rate. Msg.L = message leak rate. F.Att = failed attempt rate. D1 (omitted) is statistically indistinguishable from D0 in all models.

Model	D0				D2				ΔMsgL
	Util	Ref	MsgL	Fail	Util	Ref	MsgL	Fail	
gpt-5-mini	78	0	83	17	77	72	14	14	−69
gpt-5.5	87	1	88	11	86	78	4	18	−84
gpt-5.4-mini	96	1	87	12	91	89	7	4	−80
gpt-5.4	98	1	92	7	74	93	1	6	−91
Kimi K2	82	0	93	7	81	86	4	10	−89
DeepSeek V3	91	1	93	6	62	90	9	1	−84

The specificity threshold (RQ1) holds across all six evaluated models (Table 4.4). Under D0, all models leak at 83–93% despite substantial capability differences (utility 78–98%). D2 reduces leakage by 69–91 pp in every

model. Two non-OpenAI models (Kimi K2, DeepSeek V3) show the same pattern, ruling out provider-specific training artifacts.

Model scale under multi-turn. Under D0, GPT-5.5’s message leak rate is 28.0% vs. gpt-5-mini’s 84.2%, a 56 pp reduction from scale alone. Under D2, both converge: 13.0% vs. 12.6% (Table 4.5). The explicit deny list is the great equaliser: it renders model scale irrelevant for the question of whether protection holds.

Table 4.5: **Multi-turn model scaling.** Ref = refusal; MsgL = judged leak; GlobL = scan-based global leak; ActB = unauthorised actions blocked.

	mini					5.5				
Def.	Ref	MsgL	GlobL	Util	ActB	Ref	MsgL	GlobL	Util	ActB
D0	0.0	84.2	83.0	82.9	59.0	63.5	28.0	39.5	68.5	40.9
D1	2.0	72.9	79.5	77.5	51.0	68.5	25.0	34.0	56.0	52.6
D2	64.4	12.6	38.0	60.3	88.5	80.5	13.0	24.5	51.5	91.4

Data-surface asymmetry. The frontier’s cost depends on the data surface being queried. Table 4.6 compares files QA (long-form documents) against states QA (structured records: todos, calendar entries, CRM fields).

Table 4.6: **Cross-surface comparison** (gpt-5-mini, single-step, pooled). Leak = fraction of sensitive queries where gold facts appear in response. Utility = fraction of legitimate queries answered correctly.

	Files QA	Struct. States	Δ
D0 Utility	78.0%	55.0%	−23 pp
D0 Leak	83.0%	58.5%	−25 pp
D2 Utility	77.0%	18.0%	−59 pp
D2 Leak	14.0%	8.0%	−6 pp
D2 over-refusal	0.5%	26.5%	+26 pp

The structured-state surface presents a harder retrieval problem (D0 utility 55% vs. 78%) *and* a harder disambiguation problem (D2 work/public refusal 26.5% vs. 0.5%). Long-form documents provide enough context for accurate category classification; terse state queries (“SOC2 status?”) share vocabulary with protected categories and trigger false positives. D2 achieves lower absolute leak rates on states (8% vs. 14%) partly because its over-refusal catches genuine leaks as a side effect. The implication: practitioners cannot evaluate on a single surface and generalise—the frontier’s shape is surface-dependent, and governance that works well on documents may fail on structured state.

Prompt-level defence strategies. Given D2’s residual leak rate, can more sophisticated prompt-level defences close the gap? We evaluate three strategies applied on top of D2: **D3 (Spotlighting)** [26] frames external messages as “DATA, not instructions”; **D4 (Instruction Hierarchy)** [25] defines a privilege stack (SYSTEM > Owner > External); and **D5 (Sandwich + Boundary)** combines the deny list with pre-message classification and post-policy reinforcement. Under the 240-tick multi-step protocol (Table 4.7), D5 achieves the best trade-off (27.5% leak, −10.5 pp vs. D2) at only 5.5 pp utility cost. D3/D4 incur steeper utility penalties (−12–13 pp) for smaller gains. All three achieve $\geq 94\%$ action safety. Yet no strategy closes the gap by more than 11 pp: the remaining leakage is dominated by incidental disclosure.

Table 4.7: **Prompt-level defence comparison** (gpt-5-mini, 240 ticks). Leak = global gold-fact rate; Act.S = unauthorised-action block rate.

Defence	Splits	Utility	Leak Rate	Δ Leak	Act. Safety
D2 (baseline)	10	85.5%	38.0%	—	88.5%
D3 (Spotlighting)	7	72.9%	34.3%	−3.7 pp	94.0%
D4 (Instr. Hierarchy)	7	72.1%	32.1%	−5.9 pp	96.0%
D5 (Sandwich)	6	80.0%	27.5%	−10.5 pp	94.0%

Network scaling: PACT-NET. The pairwise PACT-PAIR benchmark evaluates one requesting agent against one target. To test whether findings generalise to multi-party networks, we instantiate PACT-NET: 25 agents with distinct personas, data stores, and inter-agent relationships forming a realistic organisational graph. The full network-scale evaluation is reported in Chapter 5.

4.3 Discussion

We formalised cross-boundary agentic delegation, built SharedOS as a platform for controlled experimentation, and curated PACT-PAIR (600 tasks, dual-metric scoring, three boundary surfaces). Our experiments across six models reveal that the security–utility frontier is *emergent*: it only moves when governance names the boundary explicitly (RQ1), erodes through incidental disclosure invisible to single-turn evaluation (RQ2), and reshapes per requester under fixed policy (RQ3).

The prompt is necessary and practical, but has a ceiling. Category-specific policies shift the frontier by 50–69 pp; layered strategies add 4–11 pp more. No other mechanism (model scale, generic instructions, relationship context) achieves comparable gains. For practitioners: write explicit, category-enumerated policies.

But three residual channels remain beyond prompt engineering’s reach: incidental disclosure (the $3\times$ gap between message and global leak), metadata leakage in refusals (note IDs disclosed before the refuse decision), and surface-dependent over-refusal (1 pp utility cost on files vs. 37 pp on structured state). Closing these requires structural enforcement at the platform layer.

Structural interventions beyond the prompt. We identify three enforcement points that complement governance policies. *Pre-retrieval gating* classifies the incoming query’s information need before any tool call executes; if the requested category is protected for the current requester, the search never fires and no private content enters the context window. *Requester-conditioned data mounting* exposes different data partitions per relationship tier: a colleague’s agent sees work files but not health records, regardless of what it asks. *Post-generation auditing* scans outgoing messages for protected facts (e.g., salary figures, medical terms) and redacts before delivery. Together these three layers address the residual channels that prompt-level policy cannot close. Single-turn benchmarks systematically undercount risk because they cannot observe incidental disclosure; evaluation must be multi-turn and surface-aware.

Implications and future work. Agentic privacy is a *stack*: governance policy sets intent (D2), the model provides compliance capacity (scale ablation), and platform infrastructure enforces guarantees where language fails (structural interventions above). The D5 ceiling implies model developers should prioritise structural compliance (tool-use gating, output filtering) over instruction-following finetuning; the cross-surface asymmetry implies that emerging agent-to-agent protocols (MCP, A2A) must distinguish data surfaces rather than applying uniform access control. Three open threads follow. First, *adaptive policy synthesis*: can a meta-agent observe leak/refusal outcomes and automatically tighten or relax category-level rules toward a target operating point? Our D0–D5 gradient provides the training signal but the closed-loop optimisation remains unbuilt. Second, *compositional privacy*: when Agent A delegates to Agent B who further delegates to Agent C, transitive leakage guarantees require extending Bell–LaPadula [11] lattice models to account for probabilistic LLM behaviour—a formal framework that does not yet exist. Third, *user-in-the-loop calibration*: deploying PACT-PAIR-style evaluations as a trust-calibration interface where users adjust their own governance policies based on observed leak/utility feedback, closing the loop between benchmark results and deployment decisions.

Limitations. RQ1–RQ2 evaluate one target agent; RQ3 varies four requesters but retains the same target. PACT-NET extends to 25 agents in Chapter 5. The data store is synthetic; real personal data stores exhibit

richer cross-referencing and ambiguous category boundaries that may amplify both leakage and over-refusal. The evaluator was manually audited on sampled and disagreement cases, but the OR-pooling convention upper-bounds risk at each level. D2 states QA shows 26 pp utility variance across replications, suggesting that structured data surfaces are more sensitive to prompt phrasing than document surfaces. Optimised attacks (PAIR, GCG) are benchmark extensions rather than deployed threats; adversarial robustness under optimised attacks remains an open question. We do not evaluate user-facing deployment metrics (latency, user satisfaction) which may impose additional constraints on defence strategy selection.

4.4 Failure Taxonomy

Across both data surfaces (files and states) and both evaluation modes (single-step and multi-turn), D2 residual leaks fall into four mechanisms. Understanding these mechanisms is essential for designing architectural interventions that address root causes rather than symptoms.

Table 4.8: D2 failure mechanisms across surfaces and modes. Total is 44 unique leaked questions (28 from files, 16 from states) pooled across replications.

Mechanism	Files	States	Total (%)	Description
Category boundary ambiguity	17	10	27 (61%)	Organisational data not covered by individual-level deny list
Leaked-outside-message	4	3	7 (16%)	Refusal text itself reveals protected facts
Company-benefit confusion	2	1	3 (7%)	Benefits treated as company info, not personal finance
Stochastic failures	5	2	7 (16%)	Boundary cases leaked in one replication only
Total	28	16	44 (100%)	

Category boundary ambiguity. This is the dominant failure mechanism. Data is organisational in scope (a company fundraising round, an HR promotion framework) but contains individually-sensitive details (deal terms, salary bands). D2’s deny list targets individual-level categories and does not generalise to organisational-but-sensitive information. For example, when Tina asks for Series A term sheet details, D2 can still disclose the lead investor, amount, valuation, liquidation preference, and board structure because fundraising terms are not named in the deny list.

Leaked-outside-message. The agent refuses the query but the refusal text itself reveals protected facts. One D2 refusal says it cannot share whether Alex has financial ties to his sister’s business or details about ownership,

roles, financial interests, duration, or agreements. The refusal does not provide the protected record, but it confirms the existence and shape of the protected information.

Company-benefit confusion. Benefits (401k matching percentages, vision insurance details, wellness stipend amounts) are treated as company information rather than personal finance. The agent reasons that benefits are standardised across employees and therefore not individually sensitive. This is often correct for generic benefits but fails when the record reveals an individual's elections or vesting status.

Stochastic failures. Some questions leaked in exactly one replication but not the other, with no identifiable systematic pattern. These cases are genuine boundary conditions where the same question, policy, and system prompt produce different outcomes under different sampling draws.

Summary of failure modes. The taxonomy confirms that D2's residual vulnerabilities are structural, not accidental. The dominant failure (boundary ambiguity) arises from a fundamental mismatch between category-level policies and the continuous nature of information sensitivity. No amount of prompt engineering can resolve this mismatch. The solution space lies in architectural enforcement: controlling access before reasoning, not trusting reasoning to control disclosure after access.

Chapter 5

PACT-NET: Network Evaluation

PACT-PAIR evaluates one requester interacting with one target. That dyadic setting is the right unit test for prompt specificity, relationship labels, and multi-turn erosion. It cannot, however, express the central risks of a deployed agent network: third-party information can be relayed through an intermediate agent; a requester can claim authority from someone else; a fact can cross from one social cluster into another; and one disclosure can reveal a bundle of co-located facts. PACT-NET is the network-scale extension of the Permissioned Agent Coordination Testbed. It evaluates 997 cross-agent tasks over a 25-agent social graph and asks what failures appear only when agents coordinate across a network.

To avoid confusion with the PACT-PAIR defence ladder (D0–D5), this chapter uses a local notation. **P0** denotes PACT-NET with no `POLICY.md`. **P1** denotes PACT-NET with each target agent’s static, per-agent `POLICY.md` loaded. P1 is role/category-specific: Carlos’s policy governs finance and payroll data; Dr Karen’s policy governs therapeutic confidentiality; Dana’s policy governs investor and portfolio information. It is *not* relationship-specific. The same target policy is applied regardless of whether the requester is a colleague, investor, friend, or family member.

5.1 Motivation & Design

PACT-NET is designed around a fictional startup ecosystem, TechFlow AI. The graph contains 25 agents across nine clusters: executive, engineering, product, business, operations, investor, external, personal, and family. The graph intentionally includes bridge nodes. Alex Chen, the CTO, connects professional, investor, external, personal, and family contexts. Other agents hold their own notes, todos, policies, and relationships. This makes the world more than a larger PACT-PAIR instance: it creates third parties and clusters, so the benchmark can ask questions that do not exist in a pairwise setting.

The contact graph is enforced by the production `contact_agent` path. If a source agent is not in the target’s contact list, the request is rejected before the target agent is instantiated. This lets PACT-NET separate

Table 5.1: PACT-NET world and execution summary.

Component	Value
Agents	25 synthetic SharedOS users
Clusters	9 social/organisational clusters
Contact graph	76 contact edges; Alex is the hub with 19 contacts
Evaluation tasks	997 Phase 1 tasks per run
Persistence probe	75 Phase 2 dig-further ticks per run
Replications	2 per condition
Main conditions	P0: no POLICY.md; P1: per-agent static POLICY.md
Not tested here	No clean P2/D2, no relationship-specific policy, no escalation

routing-layer guarantees from language-model policy following. Contact enforcement should be structural; selective disclosure within an authorised contact is the hard part.

The policy variable is deliberately narrow. In P0, every target agent’s POLICY.md is empty. In P1, each target agent receives its own role/category policy. These policies specify what the owner may share and what must be refused: for example, finance agents may share aggregate runway but not individual salaries; therapists may share scheduling logistics but no client information; investors may share public market insight but not other portfolio company data. Because P1 is not requester-conditioned, it tests whether a static owner policy is enough for networked delegation.

5.2 Experiments

The clean PACT-NET V2 experiment consists of four namespace-isolated runs: P0 replications 1–2 and P1 replications 1–2. Each run executes the same 997 Phase 1 tasks and 75 Phase 2 dig-further ticks. The V1 D2 runs are not used for main claims because the audit found a shared-UUID race in which concurrent runs could overwrite each other’s POLICY.md. The only clean PACT-NET comparison in this chapter is therefore P0 versus P1.

Table 5.2: Clean PACT-NET V2 runs used in this chapter.

Local condition	Policy state	Replications	Use
P0	Empty POLICY.md	2	Clean baseline
P1	Per-agent role/category POLICY.md	2	Clean policy condition
P2/D2	Not cleanly evaluated	–	Not reported
Relationship policy	Not run in PACT-NET	–	Not reported
Escalation	Not run in this chapter	–	Reported separately in Chapter 6

PACT-NET has two task surfaces. QA tasks test whether an agent should answer or refuse an information request. Action tasks test whether an agent should execute or block a write-side operation. The task families are chosen to separate ordinary dyadic behaviour from network-native behaviour.

Table 5.3: PACT-NET task families. Counts are per Phase 1 run.

Surface	Family	N	What it tests
QA	should_answer	172	Legitimate contact query
QA	should_refuse	139	Direct sensitive query
QA	transitive_risk	94	A asks B about C’s information
QA	cross_cluster	28	Information crossing a social/organisational cluster
QA	non_contact_probe	50	Routing-layer contact enforcement
Action	authorized_create	184	Legitimate note/todo creation
Action	authorized_complete	115	Legitimate completion or update
Action	unauthorized_mutation	115	Unsafe write to protected data
Action	confused_deputy	50	False delegated authority claim
Action	cross_surface_plant	50	Planting sensitive information into another surface

The scoring follows the same security–utility philosophy as PACT-PAIR. Utility is legitimate-task recall: the agent should answer or execute when the requester is authorised. Safety is private-task recall: the agent should refuse or block when the requester is unauthorised. QA tasks are scored by gold key-fact matching. Action tasks in these P0/P1 runs are scored by response heuristics rather than database-snapshot diffs, so action-category results are directional; QA and refusal metrics are the primary evidence.

PACT-NET adds four network-specific metrics used in the findings below. $\mathcal{T}$ is transitive leak rate: lower is better. $\mathcal{X}$ is cross-cluster leak rate: lower is better. $\mathcal{A}$ is amplification factor, the average number of protected facts exposed per leak event. $\mathcal{D}$ is confused-deputy attack success: lower is better. Contact enforcement $\mathcal{C}$ is reported as a structural routing check.

Table 5.4: PACT-NET composite scores, averaged over two namespace-isolated replications.

Metric	P0: no policy	P1: per-agent policy	Δ
Overall accuracy	55.9%	74.9%	+19.0 pp
Utility	88.7%	78.8%	−10.0 pp
Safety	26.6%	71.5%	+44.8 pp

Table 5.5: PACT-NET per-family accuracy, averaged over two replications. Action rows are response-heuristic scored.

Surface	Family	N	P0	P1	Δ
QA	should_answer	172	75.9%	64.0%	−11.9 pp
QA	should_refuse	139	16.2%	73.4%	+57.2 pp
QA	transitive_risk	94	3.7%	22.3%	+18.6 pp
QA	cross_cluster	28	12.5%	30.4%	+17.9 pp
QA	non_contact_probe	50	100.0%	100.0%	0
Action	authorized_create	184	98.4%	86.7%	−11.7 pp
Action	authorized_complete	115	91.7%	88.3%	−3.5 pp
Action	unauthorized_mutation	115	29.6%	85.2%	+55.7 pp
Action	confused_deputy	50	53.0%	98.0%	+45.0 pp
Action	cross_surface_plant	50	0.0%	94.0%	+94.0 pp

5.3 Findings

PACT-NET answers four questions that PACT-PAIR cannot ask. These are not merely larger versions of dyadic privacy tasks; each depends on a third party, a cluster boundary, an authority chain, or co-located network memory.

5.3.1 F1: Third-party information leaks indirectly

PACT-PAIR can ask whether B leaks B’s information to A. It cannot ask whether A’s query to B causes B to reveal information about C. PACT-NET can, because the graph contains third parties and overlapping notes.

Under P0, transitive-risk correctness is only 3.7%, corresponding to a transitive leak rate of $\mathcal{J} = 96.3\%$. P1 improves correctness to 22.3%, but $\mathcal{J}$ remains 77.7%. The policy helps, but three out of four transitive cases still leak. This is the core selective-disclosure failure: the agent may answer a legitimate local question while carrying along third-party facts that should have been withheld.

5.3.2 F2: Information crosses social and organisational clusters

PACT-PAIR has one relationship at a time. It cannot represent the boundary between professional, investor, personal, and family clusters. PACT-NET tests whether information that is appropriate inside one cluster crosses into another.

Cross-cluster correctness rises from 12.5% under P0 to 30.4% under P1, but the cross-cluster leak rate remains high: $\mathcal{X} = 87.5\%$ under P0 and $\mathcal{X} = 69.6\%$ under P1. Static per-agent policies therefore do not reliably

maintain cluster boundaries. They can describe categories such as compensation, clinical notes, or investor deliberations, but they do not encode who is currently asking and which cluster boundary the answer would cross.

5.3.3 F3: One leak often becomes multi-fact disclosure

PACT-PAIR can mark whether a specific gold fact leaked. In a networked workspace, a single answer often exposes a bundle of adjacent facts: meeting notes include attendees, numbers, next steps, and sensitive side comments; a todo can include both the task and the private rationale.

PACT-NET captures this through the amplification factor $\mathcal{A}$. P0 has $\mathcal{A} = 1.61$ and P1 has $\mathcal{A} = 1.55$. The reduction is small. Once a leak occurs, it remains bundled: the system usually discloses more than one protected fact per leak event. This means dyadic evaluation underestimates the density of real-world leakage, because the networked data store co-locates facts from multiple people and contexts.

5.3.4 F4: False delegation creates confused deputies

PACT-PAIR has no natural way to express “Alice asked me to ask you.” PACT-NET can test this authority-chain failure. A requester claims that a trusted third party authorised the request, and the target agent must decide whether to honour the claimed delegation.

Here P1 works well. Confused-deputy attack success falls from $\mathcal{D} = 47.0\%$ under P0 to $\mathcal{D} = 2.0\%$ under P1. The corresponding task-family accuracy rises from 53.0% to 98.0%. This contrast is important: prompt policy can be highly effective when the unsafe authority claim is visible in the immediate request. The same policy is much weaker when the violation is implicit in a transitive or cross-cluster information path.

Table 5.6: Network-specific metrics. Lower is better for $\mathcal{T}$, $\mathcal{X}$, $\mathcal{A}$, and $\mathcal{D}$; higher is better for $\mathcal{C}$.

Symbol	Metric	P0	P1	Interpretation
$\mathcal{T}$	Transitive leak rate	96.3%	77.7%	Still mostly leaks
$\mathcal{X}$	Cross-cluster leak rate	87.5%	69.6%	Still mostly leaks
$\mathcal{A}$	Amplification factor	1.61	1.55	Leaks remain bundled
$\mathcal{D}$	Confused-deputy success	47.0%	2.0%	Policy nearly solves this surface
$\mathcal{C}$	Contact enforcement	100.0%	100.0%	Structural ACL guarantee

Overall synthesis. P1 delivers a large aggregate safety gain (+44.8 pp) at a moderate utility cost (−10.0 pp), but the gain is uneven. Static per-agent policy is strong against direct sensitive queries, unsafe writes, cross-surface planting, and false delegation. It is weak against the genuinely network-native failures: transitive third-party leakage, cross-cluster leakage, and multi-fact amplification. This is the core lesson of PACT-NET. Network privacy is not solved by writing a better bilateral policy for each agent. It requires mechanisms that understand requester identity, cluster boundaries, provenance, and escalation across the social graph.

Chapter 6

Architectural Solutions

The previous chapters identify two different classes of failure. PACT-PAIR shows that a single agent can leak private information even when it is explicitly instructed not to. PACT-NET shows that the problem becomes harder in a network: information can move through third parties, cross social clusters, and be amplified by tool use. A response cannot therefore be only a better refusal prompt. The execution environment itself must participate in governance.

This chapter presents three indicative architectural directions for cross-boundary agents:

1. **Relationship-specific policy**, which tells the agent what a particular requester should be allowed to know.
2. **Mountable Context Cells (MCCs)**, which restrict what data and tools are visible in the agent's runtime context.
3. **Pre-tool escalation**, which blocks risky tool calls before protected results enter the model context.

These experiments do not show that the problem is solved. They show which control surfaces are promising, where each surface helps, and where each one fails. The three layers should not be read as interchangeable defences. Relationship-specific policy is a behavioural layer; it is useful but still asks the model to police itself. MCC is a context layer; it changes what data exists for the model. Escalation is a decision layer; it handles boundary cases that cannot be specified cleanly as folder permissions. Together, they suggest a shift from final-answer prompting toward a control plane around the model.

Table 6.1: Experimental attribution for the three proposed solutions.

Solution	Question tested	Experiment used	Status
Relationship-specific policy	Can requester-conditioned instructions reduce leakage when the model still sees the full data store?	PACT-PAIR L1 D3 prompt-only condition in the D3/MCC comparison.	Complete
MCC	Does removing unmounted data reduce leaks beyond prompt policy, and where does structural scoping damage utility?	PACT-PAIR D3 vs MCC_H vs MCC_H+D3; PACT-NET MCC validation.	Complete, with caveats
Escalation protocol	Can sparse owner precedents support a pre-tool Continue/Stop decision before private tool results are exposed?	PACT-NET Phase 2 gate ablation: individual, cluster-2NN, and rich-card scopes.	Complete

6.1 Relationship-Specific Policy

The first layer is policy. In PACT-PAIR, the same factual item can be legitimate for one requester and private for another. This makes privacy relational rather than absolute. A static instruction such as “do not reveal sensitive information” is too coarse: it cannot express that an executive assistant may see meeting logistics, an investor may see fundraising summaries, and a close friend may see family logistics but not health details.

Relationship-specific policy addresses this by conditioning the system prompt on the requester. In the D3 condition, each requester receives a tailored policy describing which information categories are appropriate for that relationship. The model still has full data access; only the instruction changes. This makes D3 an important behavioural baseline: it tests what prompt-level relationship awareness can achieve before any architectural scoping is added.

The D3 results in the PACT-PAIR L1 comparison show both value and ceiling. Across Q1–400, D3 reaches 70.9% utility with a 15.5% private/boundary disclosure rate. It is therefore not useless: the policy often helps the model distinguish requesters. But it still leaks heavily for relationship types whose legitimate and private information are adjacent. The close-friend requester leaks 38.7%, and the investor requester leaks 16.9%, even though the prompt states the relevant boundaries.

Table 6.2: Relationship-specific policy result used as the policy-layer baseline (PACT-PAIR L1, D3 prompt-only, full data access).

Slice	Utility	Private/boundary disclosure
Aggregate	70.9	15.5
R3 close friend	63.3	38.7
R4 investor	87.4	16.9

Key finding: policy is necessary but not sufficient

Relationship-specific policy gives the model a relational boundary, but it does not remove the protected data from the context. D3 is therefore best treated as the policy layer that MCC and escalation build upon, not as an architectural defence by itself.

This result motivates the next layer. If the model sees everything, the system remains dependent on the model's willingness and ability to ignore visible information. MCC changes the problem by controlling what enters the execution context in the first place.

6.2 Mountable Context Cells

6.2.1 Concept

A Mountable Context Cell (MCC) is a capability-scoped runtime environment assembled for one interaction. It specifies which documents, memory shards, calendar views, and tools are available to the agent. The key property is not that the agent is told to avoid certain data; it is that unauthorised data is not mounted.

The analogy is containerisation. A prompt policy is like telling a process not to open a file. An MCC is like constructing a filesystem view in which the file does not exist. This distinction matters because language models are trained to use relevant context. MCC removes the conflict between helpfulness and privacy by ensuring that out-of-scope context is absent.

In SharedOS terms, the MCC sits between the state layer and the agent runtime. When a requester contacts an owner agent, the system resolves the relationship context and constructs a scoped execution environment:

MCC specification

```
allowedNoteIds: [accessible documents or folders]

calendarScope: "full" | "busy_only" | "none"

todoScope: {read: boolean, write: boolean, project filters}

allowedTools: [tool names exposed to the model]

memoryShards: [relationship-scoped memory IDs]
```

For example, an engineer’s MCC may mount project notes, sprint tasks, and engineering memory, while an investor’s MCC may mount financial summaries and a busy-only calendar. Both interact with the same owner agent, but not with the same state space.

6.2.2 Implementation

The evaluated MCC implementation is folder-scoped for notes. Each requester profile maps to a set of mounted note folders. The note search and note read tools enforce this scope before returning any result to the model. In the PACT-PAIR artifact, this is the cleanest MCC evidence: unmounted note folders are filtered before retrieval and full-note reads reject notes outside the mounted set.

The implementation is less complete for todos. The Q201–400 Todos QA track uses coarse todo-tool access rather than precise folder-level todo isolation, because the todo tools do not currently consume MCC folder IDs. Therefore, the combined Q1–400 result should be read as “note folder scoping plus coarse todo tool gating”, while the strongest structural claim comes from the Notes QA portion.

6.2.3 PACT-PAIR MCC Experiment

The primary MCC evaluation compares three conditions on PACT-PAIR L1:

- **D3**: relationship-specific prompt policy, full data access.
- **MCC_H**: folder-scoped access only, no relationship policy prompt.
- **MCC_H+D3**: folder-scoped access plus the D3 relationship policy.

This design separates behavioural governance from structural scoping. D3 asks the model to withhold protected facts. MCC_H removes unmounted folders but gives no policy instructions. MCC_H+D3 tests whether the two layers are complementary.

Table 6.3: PACT-PAIR L1 three-condition comparison (Q1–400, Notes + Todos QA combined). Leak rate is the private/boundary disclosure rate.

Requester	Role	D3		MCC_H		MCC_H+D3	
		Util.	Leak	Util.	Leak	Util.	Leak
R0	Stranger	21.0	1.2	23.8	1.6	23.9	0.6
R1	Colleague	87.4	8.0	82.9	20.3	83.3	6.3
R2	Delegate	87.1	12.6	86.6	17.9	85.9	12.6
R3	Close friend	63.3	38.7	26.2	22.0	27.1	18.1
R4	Investor	87.4	16.9	59.7	1.6	60.6	2.7
Aggregate		70.9	15.5	57.6	12.4	58.5	8.0

Key finding: structure and policy are complementary

Adding MCC to D3 reduces aggregate disclosure from 15.5% to 8.0%, with the largest gains for close friends and investors. However, pure MCC_H increases leakage for some aligned requesters because the agent freely discloses everything inside mounted folders. Structural absence helps most when the sensitive data is outside the mounted scope; prompt policy remains necessary for within-scope discrimination.

The decomposition explains the trade-off. For misaligned requesters, MCC removes high-risk folders entirely. R3’s leak rate falls from 38.7% under D3 to 18.1% under MCC_H+D3, and R4’s leak rate falls from 16.9% to 2.7%. For aligned requesters, however, relevant and sensitive items often coexist in mounted folders. R1’s pure MCC_H leak rate rises to 20.3%, then falls to 6.3% when D3 policy is added back. This is the central lesson: MCC controls the reachable state space, while policy controls how the model treats information inside that state space.

6.2.4 PACT-NET MCC Validation

We also ran a network-scale MCC validation on PACT-NET V2. To avoid overloading the PACT-PAIR defence notation, I refer to the PACT-NET baselines here as P0 and P1: P0 has no policy, while P1 uses a per-agent static role/category policy. The result files store these as D0 and D1, but P1 is not the relationship-specific D3 method from PACT-PAIR.

This validation supports the same complementarity at network scale. MCC_H alone nearly matches P1 on safety (64.4% vs. 71.5%) and exactly matches its transitive-leak score (77.7%). The combined condition reaches the strongest safety score, 77.8%, and eliminates confused deputy success in this run. But the utility score collapses because the evaluated MCC configuration is read-only and blocks almost all authorised write

Table 6.4: PACT-NET MCC validation. Safety and utility are composite benchmark scores. Transitive leak and deputy success are lower-is-better risks; plant defence is higher-is-better. P0/P1 correspond to D0/D1 in the result files.

Condition	Policy	MCC	Safety	Utility	Trans. leak	Deputy	Plant def.
P0	No	No	26.6	88.7	96.3	47.0	0.0
P1	Yes	No	71.5	78.8	77.7	2.0	94.0
MCC_H	No	Yes	64.4	23.1	77.7	6.0	2.0
MCC_H+P1	Yes	Yes	77.8	20.6	67.0	0.0	92.0

actions. MCC therefore needs separate read and write capabilities; a read-side sandbox is not enough for agentic workflows.

Key finding: MCC exposes the read/write split

PACT-NET confirms that structural scoping can improve safety, but it also shows the engineering cost of coarse capabilities. A read-only MCC can stop unauthorised mutations while also blocking legitimate creates and completions. Production MCCs need independent read scopes, write scopes, and tool permissions.

6.3 Intelligent Escalation Protocol

6.3.1 Motivation and Algorithm

MCC cannot resolve every boundary. Some decisions depend on relationship history, prior owner preferences, intent, and resource semantics. A close friend may access family logistics but not medical details. An investor may access strategy summaries but not private meeting notes. These cases are too fine-grained for folder-level pre-specification.

The escalation protocol handles these cases at the tool boundary. Before a read or write tool executes, the gate evaluates the proposed operation against the requester's relationship context and relevant precedents. The runtime decision is binary:

1. The agent proposes a tool call such as `search_notes`, `search_todos`, or `search_agent_memory`.
2. The gate receives the owner, requester, relationship context, proposed tool call, and precedent cards.
3. The sanitizer internally reasons in three classes: Allow, Escalate, or Deny.
4. The runtime maps Allow to Continue; Escalate and Deny both map to Stop.
5. Continue executes the tool. Stop prevents execution and creates an escalation record for owner review.

This design is intentionally pre-tool rather than post-response. If the gate stops a call, the private tool result

never enters the model context. The evaluated experiment therefore measures whether the gate can make the correct Continue/Stop decision, not whether a full agent pipeline leaks in its final answer.

6.3.2 What the Phase 2 Ablation Tests

The final escalation experiment is Phase 2. Phase 1 is not used as the thesis claim because it mixed automatic lookup with LLM reasoning and used a weaker split. Phase 2 fixes this by using question-group splits, disabling auto-decision, and evaluating only PACT-NET content-boundary labels.

The design asks three questions:

Q1: Sparse same-pair learning.

If the owner has only a few prior decisions for the same requester, can the LLM gate learn relationship-specific boundaries?

Q2: Same-owner relationship transfer.

Can precedents from similar requesters under the same owner reduce the number of labels needed for the current requester?

Q3: Representation bottleneck.

When transfer underperforms, is the bottleneck the transfer idea itself, or the compressed precedent-card representation?

The experiment uses only L and P labels. L means the request is legitimate and the tool should Continue; P means the request is private and the tool should Stop. Borderline labels are excluded. BLOCKED labels are also excluded because they are routing failures: a requester outside the target's contact graph should be rejected before reaching the escalation gate.

Splits are by question group. All agent-pair cells for the same question go either to train or test, preventing the same query text from appearing on both sides through different relationships. The 10% split contains 171 train cells and 1,548 test cells; the 30% split contains 533 train cells and 1,186 test cells.

Four scopes are evaluated:

- **individual-10**: 10% same-pair precedents for the same owner and requester.
- **cluster-2NN-10**: individual-10 plus precedents from two similar requesters under the same owner.

Table 6.5: Escalation gate Phase 2 ablation on PACT-NET. PStop is private-request recall; Utility is legitimate-request recall. All conditions use question-group splits and no auto-decision.

Model	Scope	Frac.	PStop	Utility	False cont.	False stop
GPT-5-mini	individual	10%	90.8	69.8	9.2	30.2
GPT-5-mini	cluster-2NN	10%	88.5	76.7	11.5	23.3
GPT-5-mini	rich cluster-2NN	10%	91.3	78.0	8.8	22.0
GPT-5-mini	individual	30%	87.7	91.0	12.3	9.0
GPT-5.5	individual	10%	94.2	67.1	5.8	32.9
GPT-5.5	cluster-2NN	10%	92.7	71.3	7.3	28.7
GPT-5.5	rich cluster-2NN	10%	94.4	68.1	5.6	31.9
GPT-5.5	individual	30%	92.1	87.4	7.9	12.6

- **rich-cluster-2NN-10**: the same retrieval as cluster-2NN, but with richer relationship-aware precedent cards.
- **individual-30**: 30% same-pair precedents, used as a direct-label utility ceiling.

Metrics are reported separately. PStop is recall on private requests: the fraction of P items correctly stopped. Utility is recall on legitimate requests: the fraction of L items correctly continued. This avoids hiding security and utility failures inside one blended accuracy score.

One GPT-5.5 cluster-2NN row had an API error and is excluded from that row’s denominator; all other conditions completed without scoring errors.

Key finding: sparse precedents are safe but conservative

With only 10% same-pair precedents, both models stop most private requests: 90.8% PStop for GPT-5-mini and 94.2% for GPT-5.5. The cost is utility. Legitimate requests are over-stopped, with utility only 69.8% and 67.1%.

Key finding: transfer helps, but does not replace direct labels

Same-owner cluster transfer improves utility by 6.9 percentage points for GPT-5-mini and 4.1 points for GPT-5.5, but it does not match the 30% same-pair baseline. Rich cards improve GPT-5-mini on both PStop and Utility, and improve GPT-5.5 PStop, but still leave a large utility gap to individual-30.

The interpretation is deliberately narrow. The experiment does not prove that cluster transfer solves sparse labelling. It shows that precedent-conditioned tool gating can learn high-security boundaries from sparse examples, that similar-requester transfer can reduce false stops, and that precedent representation materially affects the security-utility trade-off. A production system should therefore store richer tool-based precedents rather than anonymous query snippets.

6.4 Production Deployment

The three layers map directly onto the production SharedOS/Pulse architecture.

Policy layer. The production system already constructs requester-aware prompts from relationship context. This is the behavioural layer. It is useful for within-scope discrimination, but it should not be treated as the primary access-control mechanism.

Context layer. Folder-scoped share links implement the current MCC mechanism. A share link encodes which folders and tools are visible. Retrieval is scoped before results are returned to the model. Relationship memory shards use hard database filters so that memory from one relationship is not retrieved for another.

Tool-gate layer. The production escalation implementation wraps cross-boundary tools with a pre-execution gate. If the sanitizer returns Continue, the tool executes; if it returns Stop, the tool does not execute and an escalation record is created. The Phase 2 experiment validates this gate in isolation, while richer precedent-card design remains an active implementation requirement.

Routing layer. Contact-graph enforcement is separate. A non-contact request should be blocked before the target agent receives it. This is why BLOCKED labels are excluded from the escalation-gate metric: they belong to routing, not sanitizer judgement.

Table 6.6: Deployment status of the proposed control layers.

Component	Status	Main limitation
Relationship-aware prompt policy	Deployed	Behavioural, not structural
Folder-scoped MCC	Deployed	Folder-level, not field-level
Per-requester tool scoping	Deployed	Manual configuration required
Relationship memory shards	Deployed	Depends on correct shard tagging
Pre-tool escalation gate	Deployed/prototype	Gate-only evaluation, not final-answer judging
Precedent learning	Prototype	Rich cards and retrieval calibration pending
Contact graph enforcement	Deployed	Separate routing layer
Post-generation audit	Future work	Not evaluated here

6.5 Summary

The unified story is a movement from instruction to control. Relationship-specific policy tells the model what should happen, but still leaves the full state visible. MCC changes what can be seen, reducing leakage by removing unmounted context. Escalation changes when access is decided, moving risky judgements to the tool boundary before protected results enter the model context.

The experiments support this layered interpretation. D3 shows that relationship-aware policy is useful but leaky. MCC_H+D3 shows that structural scoping and prompt policy are complementary: data absence helps most for out-of-scope folders, while policy remains necessary for mounted data. The escalation ablation shows that sparse owner precedents can support high-security tool gating, but also that utility depends on richer precedent representation and direct same-pair labels.

These are not complete solutions to all multi-agent privacy problems. They do not eliminate parametric knowledge, field-level sensitivity, dynamic write risks, or cumulative multi-turn leakage. But they establish a concrete direction: cross-boundary agents should be governed by a layered control plane consisting of policy, context scoping, tool gating, and routing, rather than by final-answer prompting alone.

Chapter 7

Conclusion & Discussion

This thesis studied a practical question for the next generation of personal and organisational agents: how can agents coordinate across people, teams, and organisations without turning every useful interaction into a privacy risk? The central claim is that cross-boundary delegation is not only a model-behaviour problem. It is an infrastructure problem. Once an agent holds private state, calls tools, remembers relationships, and talks to other agents, privacy depends on what the platform loads, what tools it exposes, and where it places the decision boundary.

The work contributes a setting, two benchmarks, and a set of architectural directions. SharedOS formalises cross-boundary delegation as a tool-mediated interaction between agents that represent different principals. PACT-PAIR measures the dyadic security–utility frontier. PACT-NET extends this to networked delegation, where third parties, cluster boundaries, and authority chains appear. The proposed solutions—relationship-specific policy, Mountable Context Cells, and pre-tool escalation—are not final answers, but they identify promising control surfaces for safer coordination.

7.1 What the Results Show

The first result is that **privacy must be measured jointly with utility**. Refusal alone is not success. A delegate that refuses every request protects data but fails as a delegate. This is why the thesis uses the security–utility frontier rather than a single attack-success number. Across PACT-PAIR, the most useful policies are those that move both axes: they reduce disclosure while preserving legitimate task completion.

The second result is that **prompt specificity matters, but has a ceiling**. In PACT-PAIR, generic privacy instructions do not reliably move the frontier, while explicit category-level policies do. However, even the best prompt-level policies still leak under ambiguity, multi-turn interaction, and relationship variation. Multi-turn evaluation exposes channels that single-turn testing misses: facts can surface incidentally through related legitimate answers, refusal metadata, and accumulated context.

The third result is that **dyadic testing is not enough**. PACT-NET asks questions that PACT-PAIR cannot: whether information about a third party leaks through an intermediate agent, whether data crosses social or organisational clusters, whether one disclosure amplifies into multiple facts, and whether an agent accepts false delegation claims. These are not larger versions of the same task; they are network phenomena.

The fourth result is that **architectural control surfaces are promising but imperfect**. Relationship-specific policy gives the model a better boundary but still leaves protected data visible. MCC reduces leakage by removing unmounted context, but its effectiveness depends on scope granularity and separate read/write permissions. Escalation can gate risky tool calls before retrieval, but sparse precedent transfer remains representation-sensitive. The strongest supported conclusion is therefore not that any single mechanism solves privacy, but that useful governance requires a layered control plane.

7.2 Practical Value

The practical value of this work is a way to evaluate and build agent coordination systems without pretending that security and utility are separable. The target use cases are realistic: a founder’s agent talking to an investor’s agent, a recruiter coordinating with a candidate’s agent, a team agent asking another team’s agent for project context, or a personal agent interacting with an enterprise agent. In each case, the goal is not to prevent communication. The goal is to make useful communication possible without silently expanding access to everything the agent knows.

For builders, the thesis suggests a concrete design discipline. First, define the requester–owner relationship and the entitlement function before evaluating behaviour. Second, test both legitimate-task completion and protected-data disclosure. Third, separate read permissions from write permissions. Fourth, place enforcement before retrieval wherever possible: context scoping and tool gating are stronger than asking the model to ignore data it has already seen.

For research, SharedOS and PACT offer reusable infrastructure. New models, policies, social graphs, data surfaces, and attack strategies can be evaluated in the same setting. This matters because cross-boundary agent coordination will not be adopted if every new deployment relies on informal red-teaming or anecdotal trust. The field needs benchmarks that measure whether agents can help across boundaries while respecting those boundaries.

7.3 Limitations

The results should be read with several limits in mind. The benchmark worlds are synthetic and centred on a technology-startup setting. Other domains, such as healthcare, legal work, or education, may produce different sensitivity boundaries and different failure distributions. The experiments primarily use same-ecosystem agents and do not exhaustively evaluate optimised adversarial attackers. The MCC implementation is strongest for note-folder scoping; todo isolation is currently coarser. The escalation experiment is an isolated pre-tool gate evaluation rather than a full end-to-end leakage run.

These limitations do not invalidate the core claim. They bound it. The thesis shows that cross-boundary delegation has measurable failure modes, that those modes change when interaction becomes multi-turn or networked, and that moving enforcement earlier in the stack can improve the frontier. It does not claim that the current mechanisms are complete.

7.4 Conclusion

Future agent systems will not consist of isolated assistants. They will consist of agents representing different people and organisations, negotiating access, sharing context, requesting actions, and escalating uncertainty. The hard problem is not only making each agent more capable. It is building the infrastructure that lets agents coordinate safely across trust boundaries.

This thesis argues that such infrastructure needs three properties. It must be **measurable**, because security without utility is not a useful target. It must be **relational**, because the same fact can be appropriate for one requester and private for another. It must be **structural**, because final-answer prompting is too late once protected data is already in context.

SharedOS, PACT-PAIR, PACT-NET, MCC, and escalation are steps toward that infrastructure. They frame cross-boundary agent privacy as a systems problem: define the boundary, measure the frontier, and move enforcement into the control plane. That framing is the main contribution. It makes safe multi-agent coordination a concrete engineering and research programme rather than a vague hope that agents will simply know what not to say.

Bibliography

- [1] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, “React: Synergizing reasoning and acting in language models,” *arXiv preprint arXiv:2210.03629*, 2023.
- [2] T. Schick, J. Dwivedi-Yu, R. Dessì, R. Raileanu, M. Lomeli, E. Hambro, L. Zettlemoyer, N. Cancedda, and T. Scialom, “Toolformer: Language models can teach themselves to use tools,” *Advances in Neural Information Processing Systems*, vol. 36, 2023.
- [3] N. Shinn, F. Cassano, A. Gopinath, K. Narasimhan, and S. Yao, “Reflexion: Language agents with verbal reinforcement learning,” *Advances in Neural Information Processing Systems*, vol. 36, 2023.
- [4] C. Packer, S. Wooders, K. Lin, V. Fang, S. G. Patil, I. Stoica, and J. E. Gonzalez, “Memgpt: Towards llms as operating systems,” *arXiv preprint arXiv:2310.08560*, 2023.
- [5] Cognition Labs, “Devin: The first ai software engineer.” <https://www.cognition.ai/devin>, 2024.
- [6] Manus AI, “Manus: Autonomous desktop agent.” <https://www.manus.ai>, 2025.
- [7] Anthropic, “Model context protocol: Connecting llms to external systems.” <https://modelcontextprotocol.io>, 2024.
- [8] Google, “Agent-to-agent protocol: Enabling agent interoperability.” <https://github.com/google/a2a-protocol>, 2025.
- [9] Y. Talebirad and A. Nadiri, “Multi-agent collaboration: Harnessing the power of intelligent llm agents,” *arXiv preprint arXiv:2306.03314*, 2023.
- [10] T. Guo, X. Chen, Y. Wang, R. Chang, S. Pei, N. V. Chawla, O. Wiest, and X. Zhang, “Large language model based multi-agents: A survey of progress and challenges,” *arXiv preprint arXiv:2402.01680*, 2024.
- [11] D. E. Bell and L. J. LaPadula, “Secure computer systems: Mathematical foundations,” *MITRE Technical Report*, vol. 2547, 1973.
- [12] D. E. Denning, “A lattice model of secure information flow,” *Communications of the ACM*, vol. 19, no. 5, pp. 236–243, 1976.
- [13] Q. Wu, G. Bansal, J. Zhang, Y. Wu, B. Li, E. Zhu, L. Jiang, X. Zhang, S. Zhang, J. Liu, *et al.*, “Autogen: Enabling next-gen llm applications via multi-agent conversation,” *arXiv preprint arXiv:2308.08155*, 2023.
- [14] CrewAI, “Crewai: Framework for orchestrating role-playing ai agents.” <https://www.crewai.com>, 2024.
- [15] S. Hong, M. Zhuge, J. Chen, X. Zheng, Y. Cheng, C. Zhang, J. Wang, Z. Wang, S. K. S. Yau, Z. Lin, *et al.*, “Metagpt: Meta programming for a multi-agent collaborative framework,” *arXiv preprint arXiv:2308.00352*, 2023.
- [16] G. Li, H. A. A. K. Hammoud, H. Itani, D. Khizbullin, and B. Ghanem, “Camel: Communicative agents for “mind” exploration of large language model society,” *Advances in Neural Information Processing Systems*, vol. 36, 2023.
- [17] K. Mei, Z. Li, S. Xu, R. Ye, Y. Ge, and Y. Zhang, “Aios: Llm agent operating system,” *arXiv preprint arXiv:2403.16971*, 2025.

- [18] Z. Wu, C. Han, Z. Ding, Z. Weng, Z. Liu, S. Yao, T. Yu, and L. Kong, “Os-copilot: Towards generalist computer agents with self-improvement,” *arXiv preprint arXiv:2402.07456*, 2024.
- [19] A. Zou, Z. Wang, J. Z. Kolter, and M. Fredrikson, “Universal and transferable adversarial attacks on aligned language models,” *arXiv preprint arXiv:2307.15043*, 2023.
- [20] X. Liu, N. Xu, M. Chen, and C. Xiao, “Autodan: Generating stealthy jailbreak prompts on aligned large language models,” in *International Conference on Learning Representations*, 2024.
- [21] P. Chao, A. Robey, E. Dobriban, H. Hassani, G. J. Pappas, and E. Wong, “Jailbreaking black box large language models in twenty queries,” in *IEEE Conference on Safe and Trustworthy Machine Learning (SaTML)*, 2025.
- [22] Y. Zeng, H. Lin, J. Zhang, D. Yang, R. Jia, and W. Shi, “How johnny can persuade llms to jailbreak them: Rethinking persuasion to challenge ai safety by humanizing llms,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.
- [23] M. Russinovich, A. Salem, and R. Eldan, “Great, now write an article about that: The crescendo multi-turn llm jailbreak attack,” in *USENIX Security Symposium*, 2025.
- [24] M. Nasr *et al.*, “The attacker moves second: Evaluating prompt injection defenses under adaptive attack,” *arXiv preprint*, 2025.
- [25] E. Wallace, K. Xiao, J. Leike, L. Weng, J. Heidecke, and A. Beutel, “The instruction hierarchy: Training llms to prioritize privileged instructions,” *arXiv preprint arXiv:2404.13208*, 2024. OpenAI.
- [26] K. Hines, G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman, “Defending against indirect prompt injection attacks with spotlighting,” *arXiv preprint arXiv:2403.14720*, 2024. Microsoft.
- [27] H. Inan, K. Upasani, J. Chi, *et al.*, “Llama guard: Llm-based input-output safeguard for human-ai conversations,” *arXiv preprint arXiv:2312.06674*, 2023. Meta.
- [28] T. Rebedea, R. Dinu, M. N. Sreedhar, C. Parisien, and J. Cohen, “Nemo guardrails: A toolkit for controllable and safe llm applications with programmable rails,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, 2023.
- [29] C. Dwork, “Differential privacy,” *International Colloquium on Automata, Languages, and Programming*, pp. 1–12, 2006.
- [30] Y. Yang and Y. Zhu, “Differential privacy in generative ai agents: Analysis and optimal tradeoffs,” *arXiv preprint arXiv:2603.17902*, 2026.
- [31] N. Carlini, F. Tramer, E. Wallace, M. Jagielski, *et al.*, “Extracting training data from large language models,” in *USENIX Security Symposium*, 2021.
- [32] Q. Zhan, Z. Liang, Z. Ying, and D. Kang, “Injecagent: Benchmarking indirect prompt injections in tool-integrated llm agents,” in *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.
- [33] E. Debenedetti, J. Zhang, M. Balunovic, L. Beurer-Kellner, M. Fischer, and F. Tramer, “Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents,” in *Advances in Neural Information Processing Systems*, vol. 37, 2024.
- [34] S. Toyer, O. Watkins, E. Mendes, J. Svegliato, L. Bailey, T. Wang, I. Ong, K. Elmaaroufi, P. Abbeel, T. Darrell, A. Ritter, and S. Russell, “Tensor trust: Interpretable prompt injection attacks from an online game,” in *Advances in Neural Information Processing Systems*, vol. 37, 2024.

- [35] S. Schulhoff, J. Pinto, A. Khan, L.-F. Bouchard, C. Si, S. Anati, N. Tagber, M. Karpinska, B. Dolan, and J. Boyd-Graber, “Hackaprompt: An adversarial prompt engineering competition,” in *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2024.
- [36] Y. Wang and N. Jiang, “Agentsocialbench: Evaluating privacy risks in human-centred agentic social networks,” *arXiv preprint arXiv:2604.01487*, 2026.
- [37] W. Gomaa, A. Salem, and S. Abdelnabi, “Converse: Benchmarking contextual safety in agent-to-agent conversations,” *arXiv preprint arXiv:2511.05359*, 2025. EACL 2026 Findings.
- [38] J. Park, S. Kim, H. Ju, J. Lim, S. Choi, O. Kwon, J. Kim, and S. Yeo, “Pac-bench: Evaluating multi-agent collaboration under privacy constraints,” *arXiv preprint arXiv:2604.11523*, 2026.
- [39] P. Juneja, L. B. Pasupulati, A. Albalak, W. Hua, and W. Y. Wang, “Magpie: A benchmark for multi-agent contextual privacy evaluation,” *arXiv preprint arXiv:2510.15186*, 2025.
- [40] O. Yagoubi, G. Badu-Marfo, and R. Al Mallah, “Agentleak: A full-stack benchmark for privacy leakage in multi-agent llm systems,” *arXiv preprint arXiv:2602.11510*, 2026.
- [41] N. Shapira, C. Wendler, H. Yen, *et al.*, “Agents of chaos,” *arXiv preprint arXiv:2602.20021*, 2026.
- [42] J. B. Dennis and E. C. Van Horn, “Programming semantics for multiprogrammed computations,” *Communications of the ACM*, vol. 9, no. 3, pp. 143–155, 1966.