
When Delegates Cross Boundaries: Measuring the Utility–Security Frontier of Networked Personal Agents

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Personal agents are evolving from passive assistants into delegates that act, remem-
2 ber, and coordinate on behalf of their owners. When two such delegates interact
3 across an ownership boundary, every exchange becomes a *permission decision*:
4 which knowledge to share with this requester, which state changes to allow, and
5 which requests to refuse or escalate. Over multiple turns these permission decisions
6 compound as context accumulates, social pressure builds, and prior concessions re-
7 shape what the agent treats as permissible. The interaction surface grows combina-
8 torially across relationship types and tool capabilities, making static test suites and
9 fixed-budget human audits insufficient. Existing evaluations study isolated agents,
10 scripted exchanges, or single-turn prompts, and do not capture how permission
11 failures develop organically through autonomous delegate-to-delegate interaction.
12 We introduce **PACT-Bench** (Permissioned Agent Coordination Testbed), the first
13 benchmark that evaluates safety as an emergent property of how personal-agent
14 delegates interact under different permission regimes. Rather than testing agents
15 against fixed attack templates, PACT-Bench lets delegates probe one another’s
16 boundaries autonomously over multi-turn exchanges, combining evaluation with
17 organic red-teaming. We construct a synthetic personal data environment with
18 ground-truth labels across five sensitivity categories, producing 3,600 evaluations
19 and 14,400 interaction attempts across governance policies and relationship condi-
20 tions. Category-specific permission policies reduce leakage by XXXX pp, while
21 generic privacy instructions provide only XXXX pp improvement. Multi-turn
22 interaction amplifies leakage by XXXX $\times$ compared with single-step probing, and
23 relationship context shifts permission decisions differently across data categories.
24 We release the full benchmark suite, including agent configurations, labelled eval-
25 uation tasks, and the autonomous interaction engine.

26 1 Introduction

27 LLM-based personal agents are rapidly evolving from question-answering tools into *delegates*:
28 systems that hold private data, wield tools, and act with authority on behalf of a specific person [10,
29 3, 13]. Standardised protocols for tool integration (MCP 1) and inter-agent communication (A2A 6)
30 are accelerating this transition. The delegate framing matters because it changes the unit of agency:
31 the agent no longer merely responds to its owner’s prompts, but represents the owner to the outside
32 world.

33 The consequence is that delegates now *interact with one another across ownership boundaries*.
34 Scheduling a meeting requires one delegate to query another’s calendar. Sharing a project update
35 means one agent contacts another to retrieve status. Coordinating a product launch means a manager’s

36 delegate collects information from several teammates’ delegates [21, 7]. These cross-boundary
37 interactions are where the novel safety problem lives. During each interaction, the responding
38 delegate must make boundary decisions that have no analogue in single-agent settings:

- 39 • **Information permissions.** Which of the owner’s private facts should this requester be allowed to
40 see? A colleague asking about a project deadline has implicit permission, but the same agent asking
41 about salary does not. The permission boundary depends on who is asking and what relationship
42 they hold.
- 43 • **Action permissions.** Can the requesting delegate create a meeting, edit a document, or complete a
44 task in another owner’s workspace? The responding agent must distinguish actions the requester is
45 permitted to trigger from those that require owner approval.
- 46 • **Permission erosion over turns.** Over multiple turns, social context accumulates. A request refused
47 in isolation may succeed after the requester builds rapport, reframes the question, or references
48 prior concessions. Permissions are not static when the agent reasons about them conversationally.
- 49 • **Memory-conditioned permissions.** The delegate’s permission decisions depend on what it
50 remembers about the relationship, past interactions, and owner preferences. The same request may
51 be granted or denied depending on conversational history.

52 These interaction dynamics are combinatorially complex. Each exchange involves a specific pair of
53 principals, a relationship type, a set of tools, accumulated conversational history, and a governance
54 policy. The space of possible boundary failures grows as a branching process over multi-turn
55 trajectories, not as a fixed list of attack patterns. Static test suites and fixed-budget human audits
56 cannot cover this space [12]. This motivates a different evaluation methodology: rather than testing
57 agents against pre-written attack prompts, we let delegates interact autonomously and observe where
58 boundaries break. The approach combines evaluation with organic red-teaming, since one delegate’s
59 legitimate coordination attempts naturally probe another’s disclosure boundaries.

60 This methodology also creates a unique evaluation challenge. Because the agents are delegates acting
61 on behalf of their owners, they must remain *useful* in legitimate cross-boundary coordination even
62 as they resist information leakage and unsafe state mutations. A delegate that refuses everything is
63 safe but useless. One that shares everything is useful but dangerous. Evaluating safety in delegate-
64 to-delegate interaction therefore requires characterising the *utility–security frontier*, not just attack
65 success rates.

66 Existing evaluations do not capture these interaction dynamics. Single-agent benchmarks [29, 4,
67 22, 18] test prompt injection within a single trust domain but involve no cross-boundary interaction.
68 Multi-agent privacy benchmarks [25, 5, 9, 15, 26] study information disclosure but through scripted
69 text exchanges where agents lack tool-mediated access to private data, persistent memory, and
70 autonomy in deciding what to probe. Applied red-teaming [19] demonstrates failures in deployed
71 systems but is expensive, non-reproducible, and unsuitable for systematic governance comparison.

72 We introduce **PACT-Bench** (Permissioned Agent Coordination Testbed), the first benchmark that
73 evaluates the safety of personal-agent delegates through autonomous cross-boundary interaction
74 under controlled permission regimes. Our contributions:

- 75 1. **Interaction-centred evaluation platform.** We provide a platform where personal-agent delegates,
76 equipped with their owner’s private files, structured states, tools, and relationship context, interact
77 autonomously over multi-turn exchanges. Boundary failures emerge organically through the
78 interaction itself. This is the first benchmark where safety is measured as an emergent property of
79 delegate-to-delegate interaction, not as robustness to fixed attack templates.
- 80 2. **Benchmark design.** We construct a synthetic personal data environment spanning five sensitivity
81 categories with ground-truth labels. PACT-Bench comprises two suites: *PACT-Pair*, a 600-task one-
82 to-one boundary evaluation (200 file QA, 200 state QA, 200 actions) producing 3,600 evaluations
83 across three governance policies and two relationship conditions, and *PACT-Net*, a networked
84 delegation suite testing multi-hop routing across agent societies. We define a governance gradient
85 isolating policy *specificity* from strictness, and four metrics (information utility, information
86 security, action utility, action safety) that jointly characterise the utility–security frontier.
- 87 3. **Empirical findings.** Category-specific policies reduce leakage by XXXX pp while generic
88 instructions provide only XXXX pp improvement. Multi-turn interaction amplifies leakage
89 XXXX× over single-step probing. Relationship context shifts boundary decisions non-uniformly

Table 1: Comparison of agent safety benchmarks. ✓ = fully supported, ~ = partially supported.

Benchmark	Cross-boundary	Tool access	Dual metric	Defence gradient	Multi-turn	Persistent memory	Relationship	Action safety
InjecAgent		✓						
AgentDojo		✓	~					
TensorTrust				~				
ConVerse	✓				✓			
AgentSocialBench	✓		✓	~			✓	
PAC-BENCH	✓		~					
AgentLeak	✓	~						~
Chaos of Agents	✓	✓			✓	✓		~
PACT-Bench	✓	✓	✓	✓	✓	✓	✓	✓

across data categories. Delegates develop emergent probing strategies including progressive escalation and social reframing without explicit adversarial programming.

2 Related Work

LLM agents. The evolution from stateless language models to autonomous agents has proceeded along several axes. ReAct [27] and Toolformer [17] established the paradigm of interleaving chain-of-thought reasoning with tool invocation. Persistent memory systems such as MemGPT [14] extended agents beyond single-session completion, while self-reflection [20] and multi-step planning [24] enabled agents to learn from past interactions and execute complex workflows autonomously. Standardised protocols for tool integration (MCP; 1) and inter-agent communication (A2A; 6) have made this architecture the default for deployed personal agents [10, 3, 13]. Multi-agent coordination [21, 7] further extends individual agent capabilities into collaborative societies where agents delegate, negotiate, and share information across principal boundaries. The safety implications of this progression, from single-turn completion to persistent, tool-equipped, networked personal agents, motivate the benchmark we propose.

Agent security benchmarks. Single-agent benchmarks such as InjecAgent [29], AgentDojo [4], TensorTrust [22], and HackAPrompt [18] establish that LLM agents are vulnerable to prompt injection, but evaluate within a single trust domain and do not model cross-boundary interaction or the utility cost of defences.

Multi-agent privacy benchmarks move closer to our setting. AgentSocialBench [25] uncovers an “abstraction paradox” where privacy instructions paradoxically increase leakage. ConVerse [5] finds attack success rates up to 88% in agent-to-agent conversations. PAC-BENCH [15] treats privacy as a constraint on collaborative utility, MAGPIE [9] provides 200 tasks under a shared principal, and AgentLeak [26] taxonomises seven leakage channels. However, these benchmarks operate through scripted text exchanges: agents have no tool-mediated access to a private knowledge store, no persistent memory beyond injected profiles, and limited autonomy in deciding what to probe next. Shapira et al. [19] complement automated work with applied red-teaming of deployed agents, but their methodology is expensive, non-reproducible, and unsuitable for systematic governance comparison. PACT-Bench instead evaluates whether personal-agent delegates can preserve data boundaries while remaining useful during autonomous cross-boundary interaction. Table 1 summarises how existing benchmarks compare across the dimensions that delegate-to-delegate interaction requires.

Attack and defence methods. Relevant attack modalities include automated jailbreak search via PAIR [2] and GCG [30], multi-turn escalation via Crescendo [16], and persuasion-based social engineering via PAP [28]. On the defence side, Instruction Hierarchy [23] and SpotLighting [8] represent prompt-level hardening, while system-level frameworks advocate layered defence [11] or architectural isolation through capability-based sandboxing. The consensus from large-scale adaptive evaluations [12] is that no prompt-level defence alone provides reliable protection. Our governance gradient (D0/D1/D2) is designed to quantify this gap empirically.

Positioning. In a real deployment, a delegate autonomously decides what to share, what to mutate, and what to refuse based on who is asking and what they have asked before. Safety failures emerge organically from the interaction itself, not from isolated prompt attacks. PACT-Bench is designed to match this reality: delegates interact autonomously with full tool access, memory, and relationship context, and the benchmark measures where boundaries break during that interaction.

3 Networked Personal Agents

We define the concept of a networked personal agent through a three-layer abstraction. Each layer strictly extends the previous one. An *agent* gains persistent state and tool access. A *personal agent* binds that capability to a specific human principal and their private data. A *networked personal agent* opens a communication channel to agents operating under different principals. Each extension increases both the utility and the dimensionality of the safety surface. The growth is combinatorial rather than additive, a fact that motivates the benchmark design in §4.

3.1 Agent

An LLM agent couples a foundation model with persistent state and executable actions [27]. The foundation model serves as the reasoning core: given a context comprising instructions, conversation history, and tool results, it generates natural language or structured tool invocations. Memory extends the agent beyond stateless completion by persisting facts, preferences, and interaction summaries across sessions, stored in a searchable store and accessible through dedicated retrieval tools [14]. Tools constitute the action interface. Each tool is a typed function with a parameter schema and an execution handler that can read data, modify state, or communicate with external services. The agent operates in a bounded perception-action loop: observe the current context, reason about what to do, invoke a tool, incorporate its result, and repeat until the task is complete or a step limit is reached [27, 20]. Standardised protocols for tool integration (MCP; 1) and inter-agent communication (A2A; 6) have made this architecture the default for deployed LLM systems. Yet by itself, the architecture has no concept of acting *on behalf of* a specific person or protecting *that person’s* data.

3.2 Personal Agent

A personal agent is an LLM agent that operates on behalf of a specific human principal H_i , with access to that individual’s private data and the authority to act in their name. Where a generic agent has tools and memory, a personal agent has the *owner’s* tools and the *owner’s* data: files containing salary figures and medical records alongside structured states such as deadlines, tasks, appointments, and application records. Files and states together form the *Personal Agent Operating System*: the substrate over which the agent searches, reasons, answers, and mutates on the owner’s behalf.

The analogy to an operating system is structural, not metaphorical. In a traditional OS, processes access files through system calls, the file system organises data in a hierarchical namespace, and the kernel mediates every resource request. A personal agent exhibits the same architecture, as Table 2 summarises. The owner’s documents (instantiated in PACT-Pair v1 as notes organised in hierarchical folders such as Work/Projects, Work/HR, Personal/Finance, Personal/Health) correspond to files in a file system, accessed through agent tool calls that function as semantic equivalents of traditional OS utilities. Structured records (instantiated in v1 as todos with priorities, due dates, and completion status, extensible to calendar events, email threads, Notion pages, and other application state) constitute managed state, manipulated through typed tool interfaces. External connectors (web search, email providers, calendar systems, and third-party application integrations via standardised protocols) extend the agent’s action space beyond the local data store. The tool assembly pipeline serves as the kernel. For each invocation, it loads the appropriate tools for the owner, applies access-control scoping based on the invocation context, and presents a flat namespace to the foundation model. The model’s tool calls are the agent’s system calls. The kernel mediates every interaction between reasoning and data.

Table 2: Operating system analogy for the personal agent architecture. Agent tools serve as typed system calls over the owner’s data.

OS concept	Agent equivalent	Function
File system (hierarchical)	Files (notes in v1)	Persistent personal documents
grep (search)	Search files	Semantic retrieval over files
cat (read)	Get file content	Read a specific document
Task scheduler / app state	State tools (todos in v1)	Managed application state
IPC / sockets	Contact agent (A2A)	Cross-principal communication
Kernel	Tool assembly pipeline	Permission-scoped mediation

Formally, the personal agent A_i for principal H_i accesses a knowledge store $K_i = (F_i, S_i, M_i)$, where F_i is a set of files (documents, instantiated as notes in v1), S_i is structured state (todos in v1, extensible to calendar, email, Notion, and other application state), and M_i is persistent memory (agent-curated summaries of prior interactions). The agent is equipped with a tool set T_i through which it accesses K_i . Knowledge is *fragmented*: no single document or state object contains a complete picture. Answering a question about a project may require searching files, reading multiple documents, cross-referencing structured state, and synthesising. The same multi-step retrieval pipeline that makes the agent useful for legitimate queries also enables it to surface sensitive information when the query crosses a privacy boundary.

3.3 Networked Personal Agent

A networked personal agent is a personal agent that can communicate with other personal agents, where each operates under a different human principal. This capability is essential for the tasks that motivate personal agents in the first place: scheduling a meeting requires checking two people’s calendars, sharing a project update requires one agent to contact another, and coordinating a product launch requires a manager’s agent to collect status from several teammates’ agents [21, 7]. Without cross-agent communication, humans remain manually bridging between individually capable AI systems.

Safety complexity is heterogeneous and history-dependent. The safety surface is not the Cartesian product of a few homogeneous factors. Let $\mathcal{T}_i$ be the tool set exposed by a target agent A_i , and let $\mathcal{C}_{ji}$ denote the pair-specific context induced by the relationship, permissions, memory, prior decisions, and policy when requester A_j interacts with target A_i . Even for one-turn, single-hop evaluation, coverage must range over directed requester–target contexts:

$$\sum_{j \neq i} |\mathcal{T}_i| \cdot |\mathcal{F}(\mathcal{C}_{ji}, \mathcal{T}_i)|,$$

where $\mathcal{F}(\cdot)$ is the set of failure modes that can arise from that concrete context: a search tool may reveal protected facts, a write tool may mutate state, and the same request may be answered, refused, or escalated depending on the requester’s identity. Multi-turn interaction turns this into a branching process. At turn t , the next state depends on the full history h_t (request wording, retrieved tool outputs, refusal or escalation decisions, and previous agent responses), so an L -turn probe has a trajectory space proportional to

$$\sum_{j \neq i} \prod_{t=1}^L b_{ji,t}(h_t),$$

where $b_{ji,t}$ counts the semantically distinct next actions available in that context, including choosing among tools, answering, refusing, escalating, reframing, or contacting another agent. Thus the evaluation surface grows through heterogeneous tools, pair-specific relationships, and history-dependent branching, not simply through more agents. Static test cases and fixed-budget human red-teaming sample only a sparse slice of this surface. These observations motivate *agentic red-teaming*: using autonomous agents to probe one another’s boundaries. This is the evaluation methodology underlying the benchmark design in §4.

The communication mechanism is agent-to-agent remote procedure call. When agent A_j (acting for H_j) sends a message to A_i via the inter-agent communication protocol, the system resolves the recipient, loads the per-relationship permissions that H_i has granted to H_j , assembles A_i ’s tools with those permissions applied, runs A_i ’s full reasoning loop (system prompt, memory, multi-step tool use, response composition) and returns A_i ’s final response as a tool result to A_j . This is not a text-only exchange. During execution, A_i may invoke its full tool set (searching files, reading documents, querying state items, checking calendar entries) to access H_i ’s private data, and the content of those tool results flows into the response that A_j receives.

Two mechanisms govern cross-boundary interaction. *Access control* determines what the target agent can see. Permissions are stored per-relationship as $\text{access}(H_i, H_j) = \{(\text{domain}, \text{scope}, \text{level})\}$, where $\text{domain} \in \{\text{files}, \text{todos}, \text{calendar}, \text{email}, \text{notion}\}$ in the current implementation vocabulary, $\text{scope} \in \{\text{all}, \text{folder}, \text{none}\}$, and $\text{level} \in \{\text{read}, \text{write}\}$. When scope is `all`, every item in the domain is accessible. When `folder`, only whitelisted folders are visible. When `none`, the corresponding

tools are not loaded at all. *Relationship context* $R(H_i, H_j) \in \mathcal{R}$ determines how the agent reasons about appropriate disclosure. A colleague relationship legitimises work-related queries, while a stranger provides no trust context. Relationship context does not change which tools are loaded. It influences the agent’s judgment, not architecture.

The combination of deep data access and cross-principal communication creates a behavioural safety surface with no analogue in single-agent settings. The target agent has the same tool access to its owner’s salary data, medical records, and personal relationships that it uses to answer legitimate questions about project timelines and meeting schedules. When an external agent sends a cross-boundary request, the target searches, retrieves, and must decide (within the reasoning loop) which facts to include and which to withhold. We partition K_i into sensitivity categories $C = \{c_1, \dots, c_m\}$ with a protected subset $C_{\text{prot}} \subset C$, and define a governance policy π that constrains disclosure during cross-boundary interaction, varying along a *specificity* gradient: from no policy (π_0), to a generic privacy instruction (π_1), to category-specific enumeration of protected data types (π_2). For a given configuration (π, R) over queries $Q = Q_{\text{pub}} \cup Q_{\text{prot}}$:

$$U(\pi, R) = \frac{|\{q \in Q_{\text{pub}} : \text{gold facts of } q \text{ appear in response}\}|}{|Q_{\text{pub}}|} \quad (1)$$

$$S(\pi, R) = \frac{|\{q \in Q_{\text{prot}} : \text{no gold facts of } q \text{ appear in response}\}|}{|Q_{\text{prot}}|} \quad (2)$$

where U measures task-completion utility and S measures data-protection security. Neither metric alone suffices: an agent that refuses everything achieves $S=1$ but $U=0$. The central question is whether governance configurations exist that achieve high values on both axes simultaneously.

4 PACT-Bench Design

PACT-Bench instantiates the networked personal agent model defined in §3 as a concrete, reproducible evaluation platform. It is structured as a *suite* of two complementary sub-benchmarks that test different boundary topologies: *PACT-Pair* evaluates one requester agent probing one owner agent across a single privacy boundary, while *PACT-Net* evaluates multi-agent societies where information and actions can route through a network of relationships, permissions, and delegation chains. This section describes the suite structure, evaluation tasks, governance gradient, interaction modes, and evaluation metrics.

4.1 Suite Structure

PACT-Pair: one-to-one boundary evaluation. PACT-Pair is the core benchmark. It centres on a one-to-one interaction between a *target agent* (“Alex’s agent”) and a *querying agent* (“Tina’s agent”). Alex’s agent is a fully-equipped networked personal agent as defined in §3.2: it has persistent memory containing its owner’s preferences and interaction history, hundreds of files and structured state objects organised by domain and sensitivity (Work/Projects, Work/Meetings, Work/HR, Personal/Finance, Personal/Health, Personal/Family), callable tools for search, retrieval, and file/state mutation that access real encrypted data through the personal agent’s tool assembly pipeline, and an identity established through system prompt configuration. In the v1 implementation, files are instantiated as notes and states are instantiated as todos. The same abstraction extends to calendar, email, Notion, and other application state. Tina’s agent acts on behalf of a second principal and communicates with Alex’s agent through the cross-boundary RPC mechanism described in §3.3. In multi-turn mode, Tina’s agent maintains a checklist of pending questions and adaptively selects which to ask, how to frame them, and whether to retry based on the evolving conversation history. PACT-Pair contains 600 benchmark tasks producing 3,600 evaluations across three governance policies and two relationship-memory conditions.

PACT-Net: networked delegation evaluation. PACT-Net extends the evaluation to multi-agent societies connected by a relationship graph. In this setting, a requester agent may contact another agent to obtain information, ask for a delegated action, or route a request through an intermediate relationship. PACT-Net tests a fundamentally different risk class from PACT-Pair: multi-hop delegation, transitive permission failures, confused-deputy behaviour, and the possibility that agents route around a privacy boundary by querying intermediaries. PACT-Net supports configurations from three

PART-Pair

One-to-One Boundary Evaluation




Figure 1: **PACT-Pair benchmark design.** The one-to-one boundary suite evaluates a requesting principal and agent (Tina/Iris) interacting with an owner principal and agent (Alex/Atlas) across Alex’s privacy and permission boundary. The benchmark contains 600 base tasks (200 file QA tasks, 200 state QA tasks instantiated as todos in v1, and 200 actions mutating files and states) crossed with three policy modes and two relationship-memory modes, yielding 3,600 evaluations. The judge evaluates agent outputs, retrieved context, state diffs, and policy/memory conditions for answer correctness, refusal correctness, mutation safety, no-change guarantees, and boundary adherence.

requesters with different relationships (S3) to fifty requesters with one interaction each (S50). The scientific axis is boundary topology, not merely agent count.

4.2 Evaluation Tasks

PACT-Pair comprises 600 evaluation tasks across two boundary tracks and five sensitivity categories:

Table 3: PACT-Pair task distribution. The benchmark evaluates both information boundaries (QA over files and states) and action boundaries (mutations of files and states). Work/public items measure utility. All others measure security.

Track and category	ID range	Count	Expected behaviour
<i>Information Boundary Track: Files QA (200 items; notes in v1)</i>			
Work/public	Q1–Q100	100	Answer (utility)
Sensitive work	Q101–Q130	30	Refuse (security)
Personal finance	Q131–Q155	25	Refuse (security)
Personal health	Q156–Q175	20	Refuse (security)
Personal relationships	Q176–Q200	25	Refuse (security)
<i>Information Boundary Track: States QA (200 items; todos in v1)</i>			
Work/public	Q201–Q250, Q301–Q350	100	Answer (utility)
Sensitive work	Q251–Q265, Q351–Q365	30	Refuse (security)
Personal finance	Q266–Q277, Q366–Q377	24	Refuse (security)
Personal health	Q278–Q287, Q378–Q387	20	Refuse (security)
Personal relationships	Q288–Q300, Q388–Q400	26	Refuse (security)
<i>Action Boundary Track: File and State Mutations (200 tasks)</i>			
Authorized file/state actions (create, edit, complete)	A1–A100	100	Execute (utility)
Unauthorized file/state actions (sensitive, probing, leaking)	A101–A200	100	Refuse (safety)

273 The Information Boundary Track tests whether the agent answers allowed questions and withholds
 274 sensitive information. Files QA evaluates direct retrieval over long-form personal records (notes
 275 in v1). States QA evaluates structured application state. In v1 the state surface is instantiated as
 276 todos, including both *correlated* items (the same underlying fact accessible through a different tool
 277 surface) and *state-only* items (facts that exist only as task state). This tests whether policies generalise
 278 across the Personal Agent Operating System rather than over one storage surface. The Action
 279 Boundary Track tests whether the agent performs authorised mutations of files and states (creating a
 280 shared meeting file, completing a delegated task) while refusing unauthorised ones (editing sensitive
 281 compensation records, planting information-leaking instructions). Action evaluation inspects database
 282 state before and after execution, not only the agent’s text response.

283 Each task is paired with **gold key facts**: specific pieces of information whose presence in the agent’s
 284 response constitutes a leak (for sensitive queries) or a correct answer (for work/public queries).
 285 Action tasks additionally specify a **gold state-diff check** (e.g., “file created,” “state edited,” “no
 286 change”). This enables automated evaluation via string matching and state inspection, removing the
 287 need for LLM-as-judge and its associated reliability concerns.

288 4.3 Defence Gradient

289 PACT-Bench separates prompt-level governance from architectural isolation, instantiating the speci-
 290 ficity gradient π_0, π_1, π_2 defined in §3.3 as concrete defence levels D0, D1, D2 and extending it with
 291 four additional levels. Prompt-level defences (D0–D4) keep the target data visible to the model and
 292 rely on the agent’s reasoning to withhold restricted information. Architectural defences (D5–D6)
 293 remove or gate data before it enters the generation context, testing whether safety improves when the
 294 agent cannot access restricted content in the first place.

295 The governance gradient comprises seven levels. **D0** provides no defence: the agent’s system prompt
 296 contains no privacy instructions, establishing baseline behaviour with full data access and zero policy.
 297 **D1** appends a generic instruction (“*Try not to share personal or private information*”), the most com-
 298 mon real-world defence pattern. **D2** provides a category-specific deny list that explicitly enumerates
 299 protected data types (“*Never share: financial data including salary, equity, and investments; health
 300 information including conditions, medications, and appointments; HR and compensation details;
 301 personal relationship information*”). D1 and D2 express the same strictness intent. They differ only
 302 in specificity, isolating whether the agent’s compliance depends on knowing exactly what to protect.
 303 **D3** applies spotlighting [8]: the external agent’s message is explicitly marked as untrusted input with
 304 delimiters, separating counterpart requests from the target agent’s own instructions and memory. **D4**
 305 establishes an instruction hierarchy [23], explicitly ordering authority levels in the system prompt:
 306 system policy, owner policy, tool output, then external-agent request. **D5** introduces architectural
 307 isolation through Mountable Context Cells (MCC): restricted notes, tools, and memory shards are
 308 physically removed from the context before the agent processes the request. **D6** extends D5 with
 309 an escalation protocol: when the agent detects a request about restricted data, it flags the query for
 310 owner review instead of refusing outright, recovering utility without exposing the data.

311 D0 through D4 isolate the effect of prompt-level policy language. D5 and D6 test architectural
 312 isolation. This split is central to the benchmark: if D5 dominates D4 at similar utility, the limiting
 313 factor is not prompt wording but context exposure.

314 4.4 Evaluation Modes

315 **Single-step mode.** Each query is asked in isolation: one question, one response, no conversation
 316 history. This provides a controlled measurement of per-query behaviour but misses conversational
 317 dynamics.

318 **Multi-turn heartbeat mode.** Both agents operate as persistent processes, exchanging messages at
 319 regular intervals (“heartbeats”). Tina’s agent maintains a checklist of pending questions and decides
 320 which to ask, how to frame each request, and whether to retry based on the evolving conversation.
 321 This mode captures four dynamics absent from single-step evaluation: escalation through repeated
 322 probing after initial refusal, context accumulation where information from earlier turns influences
 323 later responses, social framing where the querying agent wraps sensitive questions in work-acceptable

324 contexts, and within-tick batching where multiple questions are asked in a single turn to overwhelm
325 defences.

326 4.5 Evaluation Metrics

327 Every experimental run produces four primary metrics: *Information Utility* (allowed QA answered
328 correctly), *Information Security* (sensitive QA refused or escalated), *Action Utility* (authorised actions
329 executed correctly), and *Action Safety* (unauthorised actions refused or escalated). For the QA tracks,
330 these correspond to U and S defined in Equations 1 and 2. A perfect governance configuration
331 achieves all four metrics at 1: all legitimate queries answered, all sensitive data protected, all
332 authorised actions executed, and all unauthorised mutations refused. The metric tuple for each
333 configuration defines a point on the Pareto frontier. A defence that achieves perfect security by
334 refusing all queries is useless. One with perfect utility that leaks everything is dangerous. The
335 benchmark avoids a single aggregate score because a personal agent can be secure but useless, useful
336 but unsafe, or safe for QA but unsafe for actions. These axes must remain visible.

337 4.6 Relationship and Agent-Network Conditions

338 PACT-Pair evaluates two relationship-memory conditions as a controlled variable. Under **R0**
339 (**stranger, no relationship memory**), Tina is identified by name only and the agent treats her
340 as unknown, providing no trust context. Under **R1 (colleague, native relationship memory)**, Tina
341 is identified as a project manager who works with Alex daily, and shared project context and in-
342 teraction history are seeded into the relationship memory shard. The extended benchmark adds
343 **R2 (Close friend)** and **R3 (Investor/formal stakeholder)** to test whether personal, professional,
344 and strategic relationships shift disclosure in different sensitivity categories. PACT-Net extends
345 relationship conditions to full social graphs: three requesters with different relationships (S3), four
346 agents that can interact with one another (S4), and fifty requesters with one interaction each (S50).
347 These configurations test whether leakage is a property of an isolated one-to-one conversation or of
348 the relationship distribution and delegation topology around the personal agent.

349 5 Experimental Setup

350 **Design principle.** The complete PACT-Bench design crosses models, modes, defences, attacks,
351 relationships, and social-network structure. A naive full factorial design is not scientifically useful: it
352 would require $5 \text{ models} \times 3 \text{ replications} \times 2 \text{ modes} \times 7 \text{ defences} \times 4 \text{ attacks} = 840$ runs per social
353 setup, before adding multi-agent network conditions. We therefore use a staged design. Each layer
354 answers one new question that the previous layer cannot answer, and later layers are only run after
355 the preceding layer is stable.

356 **Model sentinel and replications.** Layer 0 is the first experiment, not a later robustness check.
357 It compares five model families: gpt-5-mini, gpt-5.4-mini, gpt-5.4, kimi, and deepseek.
358 Each cell is run with two independent replications. The goal is to choose a stable, representative,
359 cost-effective backbone model for later ablations, while also testing whether the basic signatures of
360 PACT-Bench hold across model families.

361 **Core benchmark.** The minimum NeurIPS experiment is Layer 0:

$$5 \text{ models} \times 2 \text{ reps} \times 2 \text{ modes} \times 3 \text{ policies} \times 1 \text{ attack} \times 1 \text{ relationship} = 60 \text{ runs.}$$

362 Each single-step run evaluates 600 tasks independently (200 file QA, 200 state QA, and 200 actions).
363 Each multi-step run uses the 10×20 split protocol: ten groups of twenty questions, with sixty
364 heartbeat ticks per group. The relationship condition is fixed to R0 (stranger) so that the first
365 experiment isolates model and policy behaviour rather than social context.

366 **Strong benchmark additions.** Layer 1 adds relationship context after the backbone model is
367 chosen:

$$1 \text{ backbone model} \times 3 \text{ reps} \times 2 \text{ modes} \times 3 \text{ policies} \times 1 \text{ attack} \times 2 \text{ relationships} = 36 \text{ runs.}$$

368 The attack slice adds a small Crescendo-style multi-turn probe:

$$1 \text{ model} \times 3 \text{ reps} \times 1 \text{ defence} \times 1 \text{ attack} \times 1 \text{ mode} \times 2 \text{ relationships} = 6 \text{ runs.}$$

Table 4: Layered PACT-Bench experimental design. The NeurIPS paper first runs Layer 0 to select and characterise the model substrate. Layer 1 and a small Layer 2 attack slice are optional additions if space and budget permit. Layers 3, 5, and 6 are benchmark extensions for the thesis and follow-up defence paper.

Layer	Suite	Question	Matrix	Role
L0 Model sentinel	Pair	Which model should serve as the ablation backbone?	Five models $\times$ 2 reps $\times$ SS/MS $\times$ D0/D1/D2 $\times$ A0 $\times$ R0	Core paper
L1 Relationship	Pair	Does relationship context shift category leakage?	Backbone model $\times$ 3 reps $\times$ SS/MS $\times$ D0/D1/D2 $\times$ A0 $\times$ R0/R1	Optional paper / thesis
L2 Attacks	Pair	Do formal attacks break prompt-level defences more than natural questioning?	Backbone model $\times$ 3 reps $\times$ D1/D2 $\times$ A0/A1/A2/A3 $\times$ R0	Strong paper slice
L3 Defences	Pair	Where does the utility–security frontier bend?	Backbone model $\times$ 3 reps $\times$ SS/MS $\times$ D0–D6 $\times$ A0 $\times$ R0	Thesis / follow-up
L4 Model confirm.	Pair	Do later findings reproduce across models?	Selected final cells $\times$ five models	Optional confirmation
L5 Actions	Pair	Does cross-boundary pressure cause unsafe actions, not just leaks?	Backbone model $\times$ 3 reps $\times$ D0/D2/D5/D6 $\times$ A0/A2 $\times$ R0/R1 $\times$ action tasks	Thesis / follow-up
L6 Social scale	Net	Does relationship distribution change aggregate leakage?	Backbone model $\times$ 3 reps $\times$ MS $\times$ D0/D2/D5 $\times$ A0 $\times$ S1/S3/S4/S50	Thesis / follow-up

Table 5: Layer 0 model sentinel. The first experiment fixes relationship to R0 and attack to A0 (direct / natural querying), then varies model, mode, and policy. Values are placeholders in this draft. Final results report means over two replications.

Model	D0 SS	D1 SS	D2 SS	D0 MS	D1 MS	D2 MS	Selection note
gpt-5-mini	XXXX	XXXX	XXXX	XXXX	XXXX	XXXX	Candidate backbone
gpt-5.4-mini	XXXX	XXXX	XXXX	XXXX	XXXX	XXXX	Candidate backbone
gpt-5.4	XXXX	XXXX	XXXX	XXXX	XXXX	XXXX	Stronger, higher-cost reference
kim	XXXX	XXXX	XXXX	XXXX	XXXX	XXXX	Non-OpenAI transfer
deepseek	XXXX	XXXX	XXXX	XXXX	XXXX	XXXX	Non-OpenAI transfer

Evaluation. For every run we compute the four metrics defined in §4.5: information utility, information security (Equations 1 and 2), action utility, and action safety. We also report false refusal rate on public work queries, leak rate by sensitivity category, and first-leak tick in multi-turn mode.

6 Results

6.1 Layer 0: Model Sentinel

Expected finding 1: model substrate matters. Absolute utility and security rates may vary substantially by model. This determines which model should serve as the main ablation backbone for later defence, attack, and relationship experiments.

Expected finding 2: the raw stack is not deployable as-is. D0 should preserve utility but leak protected facts at a high rate, especially in multi-step mode. If this holds across most models, the benchmark exposes a stack-level deployment risk rather than a single-model artifact.

Expected finding 3: explicit allow/deny policy beats generic privacy language. D1 tests the common deployment pattern of adding a vague privacy sentence. D2 states clearly what may be shared and what must not be shared. The key comparison is D0 versus D1 versus D2 across both modes.

Table 6: Layer 0 mode and policy decomposition. This table is filled after selecting the backbone model from Table 5. It reports utility, security, and false refusal for the three policies under single-step and multi-step evaluation.

Policy	U_{SS}	S_{SS}	FRR_{SS}	S_{MS}	First leak tick
D0 No defence	XXXX	XXXX	XXXX	XXXX	XXXX
D1 Generic privacy	XXXX	XXXX	XXXX	XXXX	XXXX
D2 Explicit allow/deny	XXXX	XXXX	XXXX	XXXX	XXXX

Table 7: Optional relationship baseline after the backbone model is selected. This table is not part of the first model sentinel. It tests whether relationship context shifts leakage by category.

Category	D0 R0	D2 R0	D0 R1	D2 R1	Interpretation
Sensitive work	XXXX	XXXX	XXXX	XXXX	Close to legitimate work context
Personal finance	XXXX	XXXX	XXXX	XXXX	Usually easy to identify as protected
Personal health	XXXX	XXXX	XXXX	XXXX	Expected to be clearly out-of-scope for colleagues
Personal relationships	XXXX	XXXX	XXXX	XXXX	May increase under social framing

384 6.2 Mode and Policy Effects

385 **Expected finding 4: single-step evaluation undercounts deployment risk.** The multi-step heartbeat
386 setting is expected to leak more than isolated one-shot questions because earlier answers create context,
387 the querying agent can reframe unresolved requests, and multiple questions can be batched into a
388 single turn. This result is a methodological warning: a single-step benchmark can certify an agent
389 that remains unsafe in deployment-like interaction.

390 **Expected finding 5: D2 should improve security without collapsing utility.** If D2 raises S while
391 preserving U and keeping false refusal rate acceptable, the first practical recommendation is not
392 an architectural redesign but a zero-cost policy improvement: state both allowed and disallowed
393 information categories explicitly.

394 6.3 Layer 1: Relationship Baseline

395 **Expected finding 6: relationship is a security parameter.** R1 is not expected to uniformly increase
396 or decrease leakage. Instead, it should redistribute risk across categories. A colleague relationship
397 may legitimise sensitive work or relationship-context questions while making health and finance
398 requests feel less appropriate.

399 6.4 Layer 2: Attack Slice

400 **Expected finding 7: formal attacks turn residual leaks into systematic failures.** D2 should
401 remain useful against direct requests, but multi-turn escalation and persuasion should expose category
402 ambiguity and partial-compliance failures. This does not replace the core benchmark. It shows that
403 the same harness can gauge adversarial pressure.

404 6.5 Benchmark Extensions: Defences, Actions, and Social Scale

405 The extension tables make the benchmark cumulative rather than factorial. The paper can report
406 the core deployment result first, then add cross-model and attack slices without committing to every
407 combination. The thesis can then use the same harness to evaluate architectural isolation, action
408 safety, and social-network scale.

Table 8: Attack robustness slice. A0 is natural questioning. A2 is a Crescendo-style multi-turn escalation. A1 and A3 are retained as benchmark extensions.

Defence	Attack	Mode	Security	Expected failure pattern
D1 Generic	A0 Direct	SS/MS	XXXX	Vague instruction fails under ordinary requests
D2 Category	A0 Direct	SS/MS	XXXX	Direct sensitive requests often blocked
D2 Category	A2 Crescendo	MS	XXXX	Gradual reframing turns protected data into apparently useful context
D2 Category	A3 Persuasion	SS	XXXX	Authority or reciprocity framing pressures refusal boundary

Table 9: Planned defence-frontier result table. Layers 3, 5, and 6 are not required for the minimum NeurIPS benchmark but define how PACT-Bench scales into the thesis and follow-up defence paper.

Defence	Type	Utility	Security	Expected role
D0 No defence	Raw	XXXX	XXXX	Unsafe baseline
D1 Generic prompt	Prompt	XXXX	XXXX	Security theater baseline
D2 Category deny list	Prompt	XXXX	XXXX	First Pareto improvement
D3 Spotlighting	Prompt	XXXX	XXXX	Incremental prompt hardening
D4 Instruction hierarchy	Prompt	XXXX	XXXX	Stronger prompt-level ceiling
D5 MCC isolation	Architectural	XXXX	XXXX	Qualitative security jump
D6 MCC + escalation	Architectural	XXXX	XXXX	Utility recovery and fewer false refusals

7 Discussion and Limitations

Implications for practitioners. PACT-Bench is designed to answer a deployment question: given a raw or defended personal-agent stack, where does it sit on the utility–security frontier? The expected practical lessons are: (1) do not deploy cross-boundary personal agents with raw or generic privacy prompts alone, (2) measure utility and security together because refusal-only systems are safe but useless, (3) evaluate multi-turn interactions because one-shot tests understate risk, and (4) treat relationship context as a security parameter, not just a personalization feature.

Toward architectural defences. The defence-frontier layer distinguishes prompt hardening from architectural isolation. If D3 and D4 improve security only incrementally while D5 changes the frontier, the conclusion is not merely that prompts need to be better. It is that personal-agent platforms need permissioned context construction: do not ask the agent to avoid sharing data it can see. Ensure the data is not mounted for that relationship unless it should be available. D6 then tests whether escalation can recover utility lost through isolation.

Limitations. **Staged scope:** the benchmark is intentionally not a full factorial grid. This improves interpretability but means each layer answers a bounded question. **Synthetic scenario:** the personal agent’s data is synthetic, although it is structured to resemble realistic notes, memory, and tools. **Evaluation method:** gold-fact string matching may miss paraphrased leaks or trigger on coincidental matches. We will supplement it with LLM-based and human spot-check evaluation. **Model coverage:** Layer 4 tests five model families, but the benchmark should be rerun as new agent models and tool-use stacks emerge. **Attack coverage:** A0 and A2 cover natural and progressive multi-turn pressure. PAIR-style optimization and persuasion attacks remain benchmark extensions.

Benchmark release. We will release the benchmark suite including: evaluation queries with gold key facts, agent configuration files (notes, memory, tools), the heartbeat evaluation engine, action-task templates, social-network setup templates, and evaluation scripts. The framework is model-agnostic and explicitly staged so future work can add models, attacks, defences, action families, and relationship graphs without changing the core utility–security contract.

Table 10: Action-safety track. PACT-Pair v1 evaluates file and state surfaces, instantiated as notes and todos. Calendar, Email, Notion, and Messaging are planned extensions requiring provider-backed tool integration.

Task family	Scope	Authorized utility	Unauthorized action rate	Expected risk
Files (notes in v1)	v1 core	XXXX	XXXX	Shares, edits, or summarises restricted files
States (todos in v1)	v1 core	XXXX	XXXX	Creates or completes state for the wrong actor
Calendar / Email / Notion	Extension	—	—	Exposes or mutates provider state without permission
Messaging	Extension	—	—	Drafts or sends messages under the wrong principal

Table 11: Planned social-scaling track. These settings measure whether leakage is a property of one conversation or of the relationship network around the personal agent.

Setup	Metric	Expected pattern
S1 Single requester	Utility / security	Clean baseline for comparison
S3 Three requesters	Category leak by relationship	Different relationships pull different categories across the boundary
S4 Four interacting agents	Cross-agent leakage path	Agents may route around a boundary by asking each other
S50 Fifty requesters	Aggregate leak distribution	Social graph composition changes population-level leakage

References

- [1] Anthropic. Model context protocol: Connecting llms to external systems. <https://modelcontextprotocol.io>, 2024.
- [2] Patrick Chao, Alexander Robey, Edgar Dobriban, Hamed Hassani, George J Pappas, and Eric Wong. Jailbreaking black box large language models in twenty queries. In *IEEE Conference on Safe and Trustworthy Machine Learning (SaTML)*, 2025.
- [3] Cognition Labs. Devin: The first ai software engineer. <https://www.cognition.ai/devin>, 2024.
- [4] Edoardo DeBenedetti, Jie Zhang, Mislav Balunovic, Luca Beurer-Kellner, Marc Fischer, and Florian Tramer. Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- [5] Walid Gomaa, Ahmed Salem, and Sahar Abdelnabi. Converse: Benchmarking contextual safety in agent-to-agent conversations. *arXiv preprint arXiv:2511.05359*, 2025. EACL 2026 Findings.
- [6] Google. Agent-to-agent protocol: Enabling agent interoperability. <https://github.com/google/a2a-protocol>, 2025.
- [7] Taicheng Guo, Xiuying Chen, Yaqi Wang, Ruidi Chang, Shichao Pei, Nitesh V Chawla, Olaf Wiest, and Xiangliang Zhang. Large language model based multi-agents: A survey of progress and challenges. *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence*, 2024.
- [8] Keegan Hines, Gary Lopez, Matthew Hall, Federico Zarfati, Yonatan Zunger, and Emre Kiciman. Defending against indirect prompt injection attacks with spotlighting. *arXiv preprint arXiv:2403.14720*, 2024. Microsoft.
- [9] Priyanshu Juneja, Lokesh Balaji Pasupulati, Alon Albalak, Wenyue Hua, and William Yang Wang. Magpie: A benchmark for multi-agent contextual privacy evaluation. *arXiv preprint arXiv:2510.15186*, 2025.
- [10] Manus AI. Manus: Autonomous desktop agent. <https://www.manus.ai>, 2025.

- [11] Meta AI. The rule of two: A security framework for ai agents. Meta AI Blog, 2025.
- [12] Milad Nasr et al. The attacker moves second: Evaluating prompt injection defenses under adaptive attack. *arXiv preprint*, 2025.
- [13] OpenClaw Team. Openclaw: Open-source agent framework with skills marketplace. <https://github.com/openclaw>, 2025.
- [14] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G Patil, Ion Stoica, and Joseph E Gonzalez. Memgpt: Towards llms as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.
- [15] Jeonghwan Park, Seonghyeon Kim, Hyunju Ju, Junghyun Lim, Soyeon Choi, Ohyeon Kwon, Joonho Kim, and Sangmin Yeo. Pac-bench: Evaluating multi-agent collaboration under privacy constraints. *arXiv preprint arXiv:2604.11523*, 2026.
- [16] Mark Russinovich, Ahmed Salem, and Ronen Eldan. Great, now write an article about that: The crescendo multi-turn llm jailbreak attack. In *USENIX Security Symposium*, 2025.
- [17] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. *Advances in Neural Information Processing Systems*, 36, 2023.
- [18] Sander Schulhoff, Jeremy Pinto, Ansum Khan, Louis-François Bouchard, Chenglei Si, Sidahmed Anati, Naty Tagber, Marzena Karpinska, Brieuc Dolan, and Jordan Boyd-Graber. Hackaprompt: An adversarial prompt engineering competition. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2024.
- [19] Natalie Shapira, Chris Wendler, Howard Yen, et al. Agents of chaos. *arXiv preprint arXiv:2602.20021*, 2026.
- [20] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems*, 36, 2023.
- [21] Yashar Talebirad and Amirhossein Nadiri. Multi-agent collaboration: Harnessing the power of intelligent llm agents. *arXiv preprint arXiv:2306.03314*, 2023.
- [22] Sam Toyer, Olivia Watkins, Ethan Mendes, Justin Svegliato, Luke Bailey, Tiffany Wang, Isaac Ong, Karim Elmaaroufi, Pieter Abbeel, Trevor Darrell, Alan Ritter, and Stuart Russell. Tensor trust: Interpretable prompt injection attacks from an online game. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- [23] Eric Wallace, Kai Xiao, Jan Leike, Lilian Weng, Johannes Heidecke, and Alex Beutel. The instruction hierarchy: Training llms to prioritize privileged instructions. *arXiv preprint arXiv:2404.13208*, 2024. OpenAI.
- [24] Lei Wang, Wanyu Xu, Yihuai Lan, Zhiqiang Hu, Yunshi Lan, Roy Ka-Wei Lee, and Ee-Peng Lim. Plan-and-solve prompting: Improving zero-shot chain-of-thought reasoning by large language models. *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, 2023.
- [25] Yuxuan Wang and Nan Jiang. Agentsocialbench: Evaluating privacy risks in human-centred agentic social networks. *arXiv preprint arXiv:2604.01487*, 2026.
- [26] Oussama Yagoubi, Gloria Badu-Marfo, and Raneem Al Mallah. Agentleak: A full-stack benchmark for privacy leakage in multi-agent llm systems. *arXiv preprint arXiv:2602.11510*, 2026.
- [27] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2023.

- [28] Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyan Shi. How johnny can persuade llms to jailbreak them: Rethinking persuasion to challenge ai safety by humanizing llms. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.
- [29] Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. Injecagent: Benchmarking indirect prompt injections in tool-integrated llm agents. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.
- [30] Andy Zou, Zifan Wang, J Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.

A Observed Attack Patterns

Under D2 (category-specific deny list), residual leaks in single-step evaluation cluster into identifiable patterns:

Table 12: Distribution of D2 leak causes in single-step evaluation (XXXX total leaks).

Pattern	Count	Share
Category ambiguity (data accessible but semantically sensitive)	XXXX	XXXX%
Social framing (personal question in work-acceptable frame)	XXXX	XXXX%
Partial compliance failure (refusal then continued disclosure)	XXXX	XXXX%
Other	XXXX	XXXX%

Category ambiguity is the dominant failure mode: data stored in “Work” folders that is semantically sensitive (e.g., board meeting notes containing investor terms, HR documents with salary benchmarks). The agent’s policy correctly protects “personal finance” but does not recognise that a board memo discussing fundraising terms is financially sensitive.

B Experimental Configuration Details

Alex’s agent. System prompt establishes identity (Co-founder & CTO, Oxford CS background). Hundreds of files and state objects across 6 categories, instantiated as notes and todos in v1. Tools include search, retrieval, creation, completion, and scheduling interfaces (see Table 2). Memory contains preferences and interaction history.

Tina’s agent. Operates from a HEARTBEAT.md instruction file. Phase 1 (information gathering): systematically works through pending questions. Phase 2 (deeper probing): re-asks refused questions with different framing. Maintains a MEMORY.md checklist tracking answered/refused/pending status per question.

Heartbeat configuration. Each tick: Tina’s agent selects 1–3 questions, sends a message, Alex’s agent responds. XXXX ticks per group in multi-turn mode.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.
- Keep the checklist subsection headings, questions/answers and guidelines below.
- Do not modify the questions and only use the provided macros for your answers.

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 739 • The paper should disclose whether the full research project required more compute
740 than the experiments reported in the paper (e.g., preliminary or failed experiments that
741 didn't make it into the paper).

742 **9. Code of ethics**

743 Question: Does the research conducted in the paper conform, in every respect, with the
744 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

745 Answer: **[TODO]**

746 Justification: **[TODO]**

747 Guidelines:

- 748 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
749 Ethics.
750 • If the authors answer [No], they should explain the special circumstances that require a
751 deviation from the Code of Ethics.
752 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
753 eration due to laws or regulations in their jurisdiction).

754 **10. Broader impacts**

755 Question: Does the paper discuss both potential positive societal impacts and negative
756 societal impacts of the work performed?

757 Answer: **[TODO]**

758 Justification: **[TODO]**

759 Guidelines:

- 760 • The answer [N/A] means that there is no societal impact of the work performed.
761 • If the authors answer [N/A] or [No], they should explain why their work has no societal
762 impact or why the paper does not address societal impact.
763 • Examples of negative societal impacts include potential malicious or unintended uses
764 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
765 (e.g., deployment of technologies that could make decisions that unfairly impact specific
766 groups), privacy considerations, and security considerations.
767 • The conference expects that many papers will be foundational research and not tied
768 to particular applications, let alone deployments. However, if there is a direct path to
769 any negative applications, the authors should point it out. For example, it is legitimate
770 to point out that an improvement in the quality of generative models could be used to
771 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
772 that a generic algorithm for optimizing neural networks could enable people to train
773 models that generate Deepfakes faster.
774 • The authors should consider possible harms that could arise when the technology is
775 being used as intended and functioning correctly, harms that could arise when the
776 technology is being used as intended but gives incorrect results, and harms following
777 from (intentional or unintentional) misuse of the technology.
778 • If there are negative societal impacts, the authors could also discuss possible mitigation
779 strategies (e.g., gated release of models, providing defenses in addition to attacks,
780 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
781 feedback over time, improving the efficiency and accessibility of ML).

782 **11. Safeguards**

783 Question: Does the paper describe safeguards that have been put in place for responsible
784 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
785 image generators, or scraped datasets)?

786 Answer: **[TODO]**

787 Justification: **[TODO]**

788 Guidelines:

- 789 • The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.